

HOMEWORK 6

Due Wed, Nov 11 at 11:30 pm

Transaction Processing

Objectives: Introduction to Transaction Processing Concepts

20.16. Add the operation commit at the end of each of the transactions T_1 and T_2 in Figure 20.2, and then list all possible schedules for the modified transactions. Determine which of the schedules are recoverable, which are cascadeless, and which are strict.

$$(5+3)! / (5! * 3!) = 8*7*6*5*4*3*2*1 / 5*4*3*2*1*3*2*1 = 56.$$

You don't need to list all 56 possible schedules; only list 2 strict, 2 recoverable, 2 non-recoverable, and 2 cascadeless schedules.

Below are the 56 possible schedules, and the type of each schedule:

S 1 : r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); C 1 ; r 2 (X); w 2 (X); C 2 ; strict

S 21 : r 1 (X); r 2 (X); w 1 (X); r 1 (Y); w 1 (Y); C 1 ; w 2 (X); C 2 ; strict

S 2 : r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); r 2 (X); C 1 ; w 2 (X); C 2 ; recoverable

S 3 : r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); r 2 (X); w 2 (X); C 1 ; C 2 ; recoverable

S 9 : r 1 (X); w 1 (X); r 1 (Y); r 2 (X); w 2 (X); w 1 (Y); C 2 ; C 1 ; non-recoverable

S 10 : r 1 (X); w 1 (X); r 1 (Y); r 2 (X); w 2 (X); C 2 ; w 1 (Y); C 1 ; non-recoverable

S 22 : r 1 (X); r 2 (X); w 1 (X); r 1 (Y); w 1 (Y); w 2 (X); C 1 ; C 2 ; cascadeless

S 23 : r 1 (X); r 2 (X); w 1 (X); r 1 (Y); w 1 (Y); w 2 (X); C 2 ; C 1 ; cascadeless

20.17. List all possible schedules for transactions T_1 and T_2 in Figure 20.2, and determine which are conflict serializable (correct) and which are not.

Below are the 15 possible schedules, and the type of each schedule:

S 1 : r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); r 2 (X); w 2 (X); serial (and hence also serializable)

S 2 : r 1 (X); w 1 (X); r 1 (Y); r 2 (X); w 1 (Y); w 2 (X); (conflict) serializable

S 3 : r 1 (X); w 1 (X); r 1 (Y); r 2 (X); w 2 (X); w 1 (Y); (conflict) serializable

S 4 : r 1 (X); w 1 (X); r 2 (X); r 1 (Y); w 1 (Y); w 2 (X); (conflict) serializable

S 5 : r 1 (X); w 1 (X); r 2 (X); r 1 (Y); w 2 (X); w 1 (Y); (conflict) serializable

S 6 : r 1 (X); w 1 (X); r 2 (X); w 2 (X); r 1 (Y); w 1 (Y); (conflict) serializable

S 7 : r 1 (X); r 2 (X); w 1 (X); r 1 (Y); w 1 (Y); w 2 (X); not (conflict) serializable

S 8 : r 1 (X); r 2 (X); w 1 (X); r 1 (Y); w 2 (X); w 1 (Y); not (conflict) serializable

S 9 : r 1 (X); r 2 (X); w 1 (X); w 2 (X); r 1 (Y); w 1 (Y); not (conflict) serializable

S 10 : r 1 (X); r 2 (X); w 2 (X); w 1 (X); r 1 (Y); w 1 (Y); not (conflict) serializable

S 11 : r 2 (X); r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); w 2 (X); not (conflict) serializable

S 12 : r 2 (X); r 1 (X); w 1 (X); r 1 (Y); w 2 (X); w 1 (Y); not (conflict) serializable

S 13 : r 2 (X); r 1 (X); w 1 (X); w 2 (X); r 1 (Y); w 1 (Y); not (conflict) serializable

S 14 : r 2 (X); r 1 (X); w 2 (X); w 1 (X); r 1 (Y); w 1 (Y); not (conflict) serializable

S 15 : r 2 (X); w 2 (X); r 1 (X); w 1 (X); r 1 (Y); w 1 (Y); serial (and hence also serializable)

20.23. Consider the three transactions T_1 , T_2 , and T_3 , and the schedules S_1 and S_2 given below.

Draw the serializability (precedence) graphs for S_1 and S_2 , and state whether each schedule is serializable or not. If a schedule is serializable, write down the equivalent serial schedule(s).

T_1 : $r_1(X)$; $r_1(Z)$; $w_1(X)$;

T_2 : $r_2(Z)$; $r_2(Y)$; $w_2(Z)$; $w_2(Y)$;

T_3 : $r_3(X)$; $r_3(Y)$; $w_3(Y)$;

S_1 : $r_1(X)$; $r_2(Z)$; $r_1(Z)$; $r_3(X)$; $r_3(Y)$; $w_1(X)$; $w_3(Y)$; $r_2(Y)$; $w_2(Z)$; $w_2(Y)$;

S_2 : $r_1(X)$; $r_2(Z)$; $r_3(X)$; $r_1(Z)$; $r_2(Y)$; $r_3(Y)$; $w_1(X)$; $w_2(Z)$; $w_3(Y)$; $w_2(Y)$;

20.23

S_1

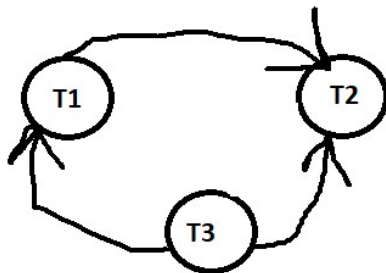
S_2

T1	T2	T3
R(x)	R(z)	R(x)
R(z)	R(y)	R(y)
W(x)	W(z)	W(y)
	W(y)	

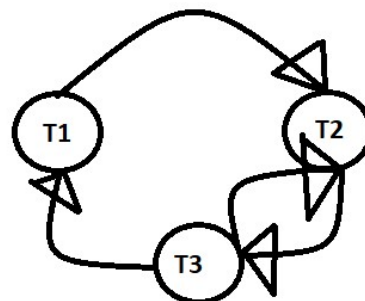
T1	T2	T3
R(x)		
R(z)	R(z)	
		R(x)
W(x)		R(y)
	R(y)	W(y)
	W(z)	
	W(y)	

T1	T2	T3
R(x)		
	R(z)	
R(z)		R(x)
	R(y)	
W(x)		R(y)
	W(z)	
	W(y)	W(y)

S_1



S_2



S_1 is conflict serializable because there is NOT a cycle.

S_2 is NOT conflict serializable because there is a cycle.

20.24. Consider schedules S_3 , S_4 , and S_5 below. Determine whether each schedule is strict, cascadeless, recoverable, or nonrecoverable. (Determine the strictest recoverability condition that each schedule satisfies.)

S_3 : $r_1(X); r_2(Z); r_1(Z); r_3(X); r_3(Y); w_1(X); c_1; w_3(Y); c_3; r_2(Y); w_2(Z); w_2(Y); c_2$;

S_4 : $r_1(X); r_2(Z); r_1(Z); r_3(X); r_3(Y); w_1(X); w_3(Y); r_2(Y); w_2(Z); w_2(Y); c_1; c_2; c_3$;

S_5 : $r_1(X); r_2(Z); r_3(X); r_1(Z); r_2(Y); r_3(Y); w_1(X); c_1; w_2(Z); w_3(Y); w_2(Y); c_3; c_2$;

Strict schedule: A schedule is strict if it satisfies the following conditions:

1. T_j reads a data item X **after** T_i has written to X and T_i is terminated (aborted or committed)
2. T_j writes a data item X **after** T_i has written to X and T_i is terminated (aborted or committed)

S_3

- **is not strict** because T_3 reads X ($r_3(X)$) **before** T_1 has written to X ($w_1(X)$) but T_3 commits **after** T_1

S_4

- **is not strict** because T_3 reads X ($r_3(X)$) **before** T_1 has written to X ($w_1(X)$) but T_3 commits **after** T_1 .

S_5

- **is not strict** because T_3 reads X ($r_3(X)$) **before** T_1 has written to X ($w_1(X)$) but T_3 commits **after** T_1 .

Cascadeless schedule: A schedule is cascadeless if the following condition is satisfied:

1. T_j reads X only **after** T_i has written to X and terminated (aborted or committed).

S3

- is **not cascadeless** because T_3 reads X ($r_3(X)$) before T_1 commits.

S4

- is **not cascadeless** because T_3 reads X ($r_3(X)$) before T_1 commits.

S5

- is **not cascadeless** because T_3 reads X ($r_3(X)$) **before** T_1 commits or T_2 reads Y ($r_2(Y)$) **before** T_3 commits.

Recoverable schedule: A schedule is recoverable if the following condition is satisfied:

1. T_j commits after T_i if T_j has read any data item written by T_i .
2. Abort operations will be used in place of commits, one at a time
3. Notations Used:
 - a. Commit Notation: $C_i > C_j$ means C_i happens **before** C_j .
 - b. Abort Notation: A_i denotes abort T_i .
 - c. Strictness:
 - i. a transaction neither reads nor writes to a data item, which was written to by a transaction that has not committed yet.

S3.

- If $A1 > C3 > C2$, then S3 is **recoverable** because rolling back of $T1$ does not affect $T2$ and $T3$.
- If $C1 > A3 > C2$, S3 is **not recoverable** because $T2$ read the value of Y ($r2(Y)$) **after** $T3$ wrote X ($w3(Y)$) and $T2$ committed but $T3$ rolled back. Thus, $T2$ used non-existent value of Y .
- If $C1 > C3 > A3$, then S3 is **recoverable** because roll back of $T2$ does not affect $T1$ and $T3$.
- Strictest condition of S3 is $C3 > C2$.

S4

- If $A1 > C2 > C3$, then S4 is **recoverable** because roll back of $T1$ does not affect $T2$ and $T3$.
- If $C1 > A2 > C3$, then S4 is **recoverable** because the roll back of $T2$ will restore the value of Y that was read and written to by $T3$ ($w3(Y)$). It will not affect $T1$.
- If $C1 > C2 > A3$, then S4 is **not recoverable** because $T3$ will restore the value of Y which was not read by $T2$.
- Strictest condition of S4 is $C3 > C2$, but it is not satisfied by S4.

S5

- If $A1 > C3 > C2$, then S5 is **recoverable** because neither $T2$ nor $T3$ writes to X , which is written by $T1$. If $C1 > A3 > C2$, then S5 is **not recoverable** because $T3$ will restore the value of Y , which was not read by $T2$. Thus, $T2$ committed with a non-existent value of Y .
- If $C1 > C3 > A2$, then S5 is **recoverable** because it will restore the value of Y to the value, which was read by $T3$. Thus, $T3$ committed with the right value of Y .
- Strictest condition of S3 is $C3 > C2$, but it is not satisfied by S5.