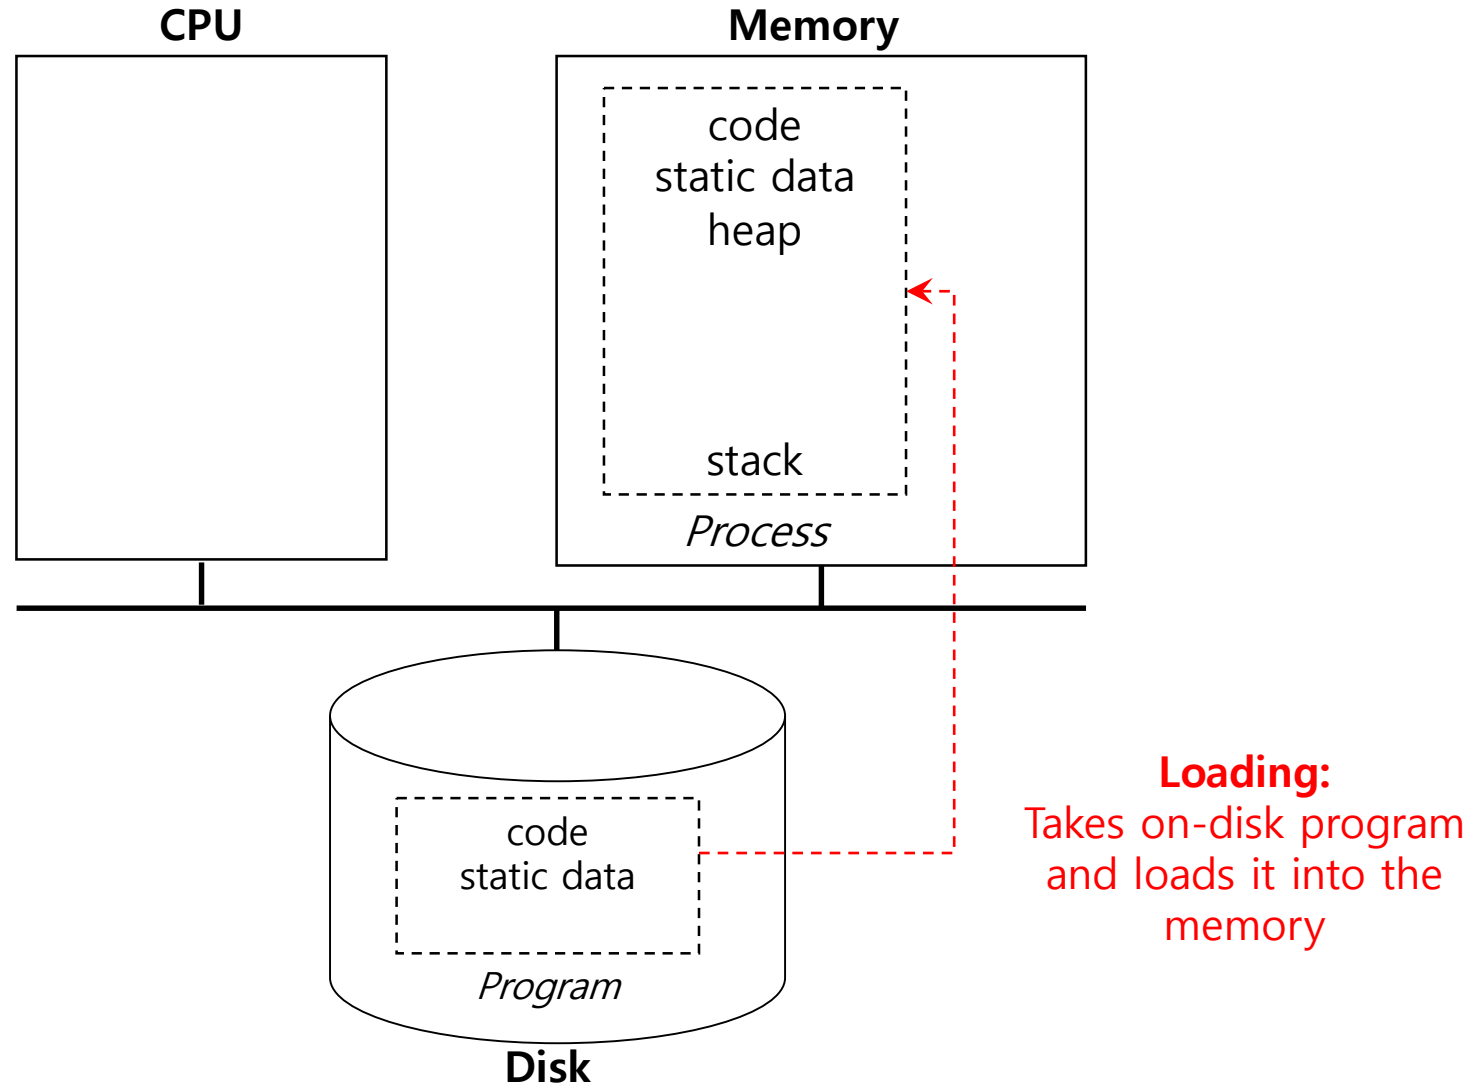


Virtualization: The CPU

Sridhar Alagar

Executing a Program



What is a Process?

- A program in execution
- It is an abstraction
- What constitutes a process?
 - Whatever it can affect
 - Memory: Address space - code, heap, stack
 - Registers, IP
 - Open files

Process API

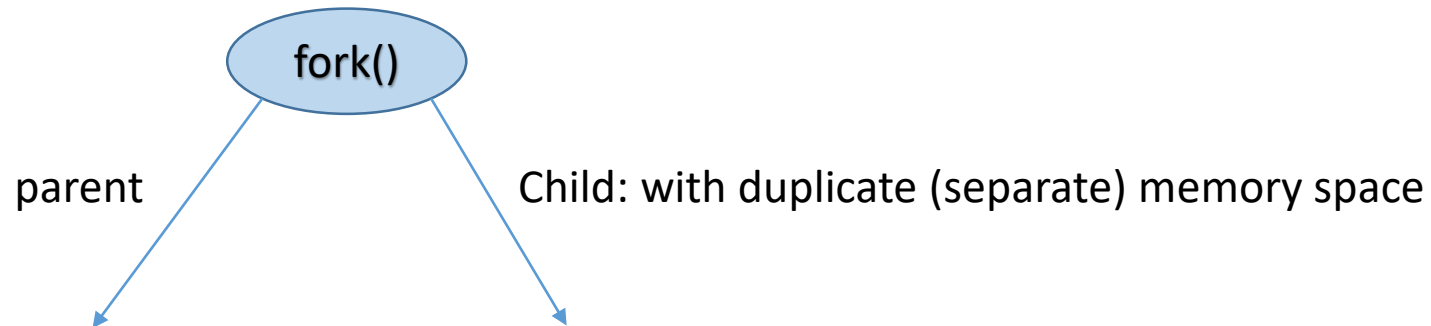
- Creation
- Destroy
- Wait
 - wait for a process to stop running
- Control
 - suspend and resume
- Status
 - get some status info about the process

Process creation

UNIX examples

fork() system call creates new child process

Parent and child processes have separate memory spaces and execute independently

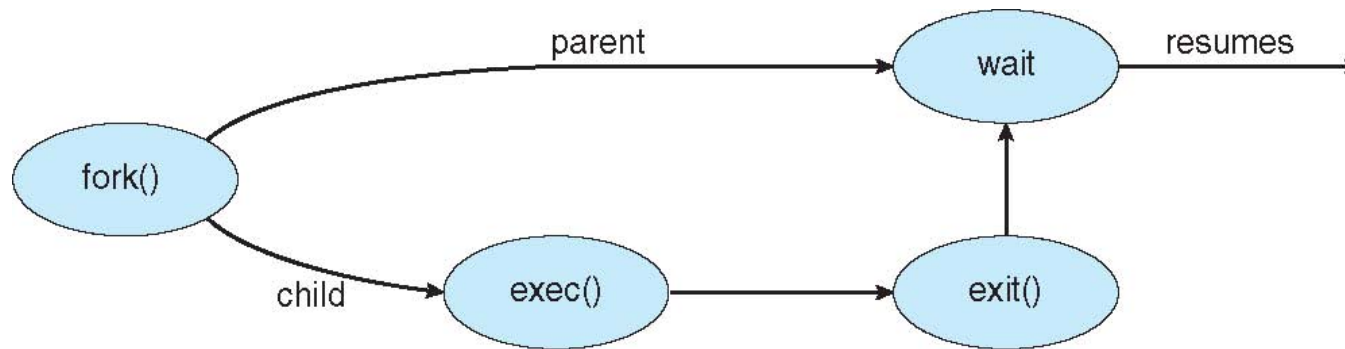


Process creation

UNIX examples

fork() system call creates new child process

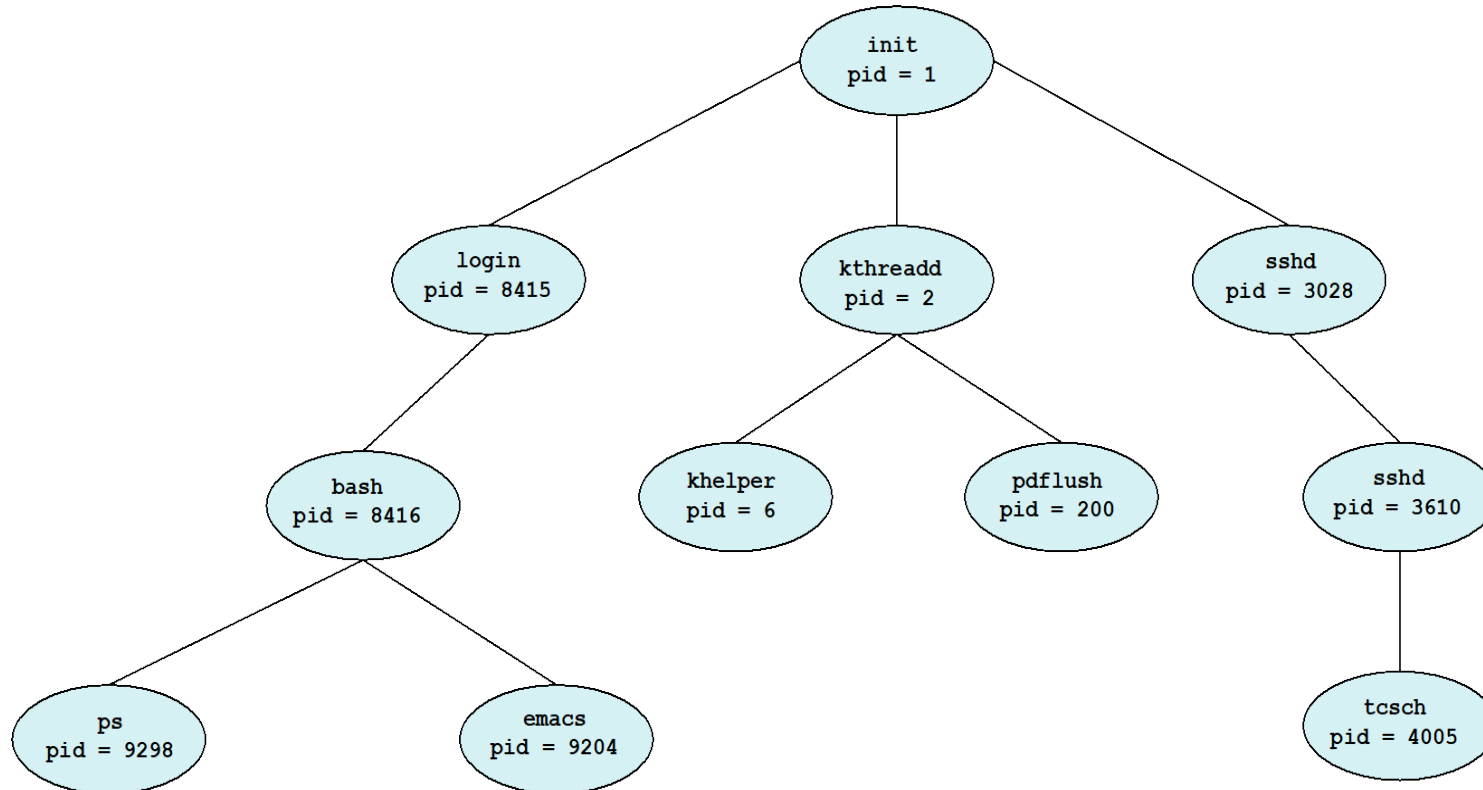
exec() system call used after a **fork()** to replace the process' memory space with a new program



Virtualize the CPU

- CPU needs to be (time) shared by many processes
- Transparent to the process (app)
- Give an illusion that a process own the CPU
 - Each process is allocated a virtual CPU

Processes tree



How to provide good CPU performance?

- **Direct Execution**
 - Runs directly on CPU
 - Full control of CPU - OS creates process and transfers control to it
- Who should be in control?

Problems with direct execution?

1. Process could do something restricted
Could read/write other process data (disk or memory)
2. Process could run forever (slow, buggy, or malicious)
OS needs to be able to switch between processes
3. Process could do something slow (like I/O)
OS wants to use resources efficiently and switch CPU to other process

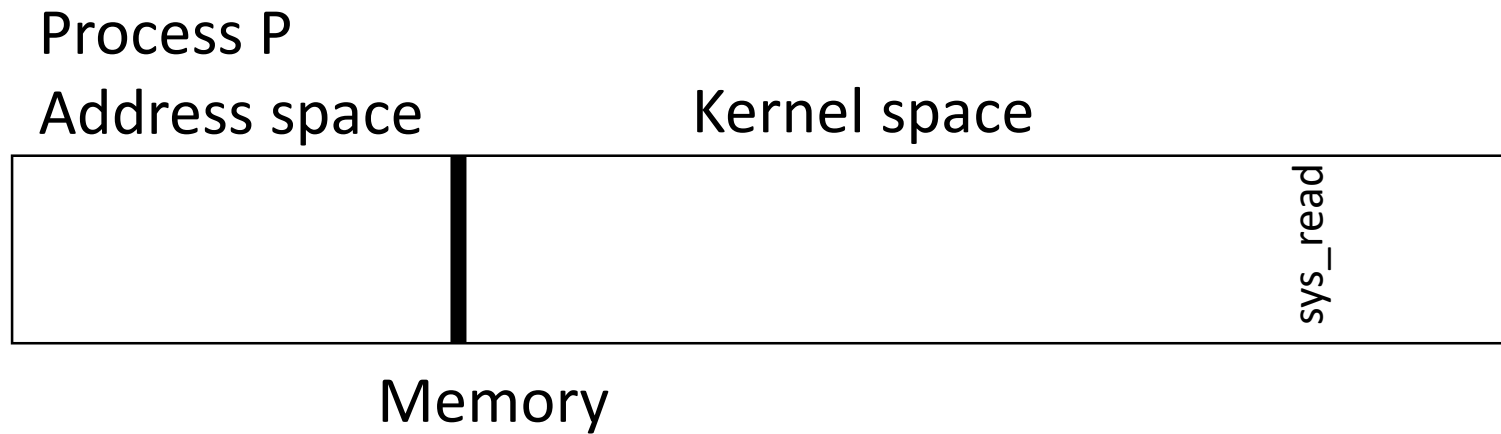
Solution:

Limited direct execution – OS and hardware maintain some control

Problem 1: How to restrict process?

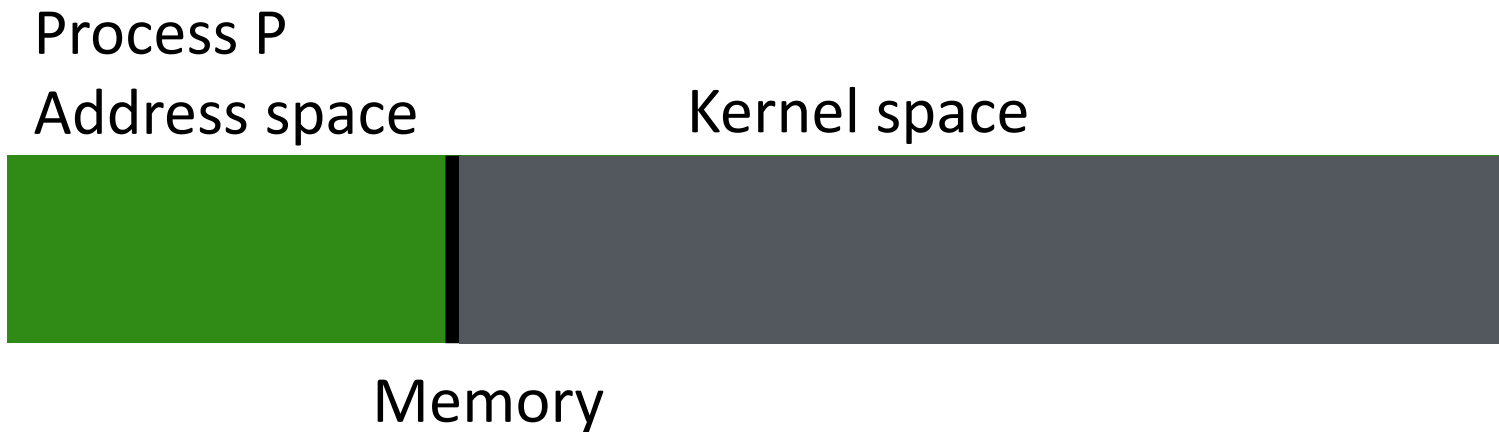
- Use privilege levels supported by the hardware
 - Process runs in user level (restricted)
 - OS runs in kernel level (unrestricted)
- How can processes access devices?
 - Need to request OS to do the job through System calls
 - Change privilege level through trap (int)

System Call



P needs to call `sys_read()`

System Call

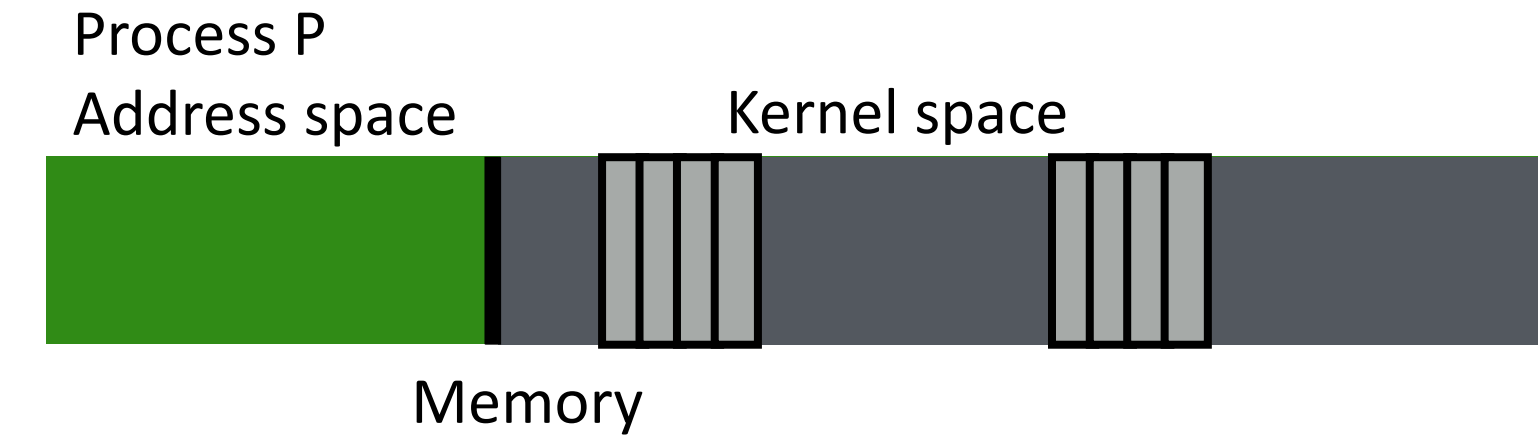


P is in user mode (restricted). It cannot see kernel space or any other space

P wants to call `sys_read`. But no way to call directly

Need hardware instruction (int) to indirectly call the routine

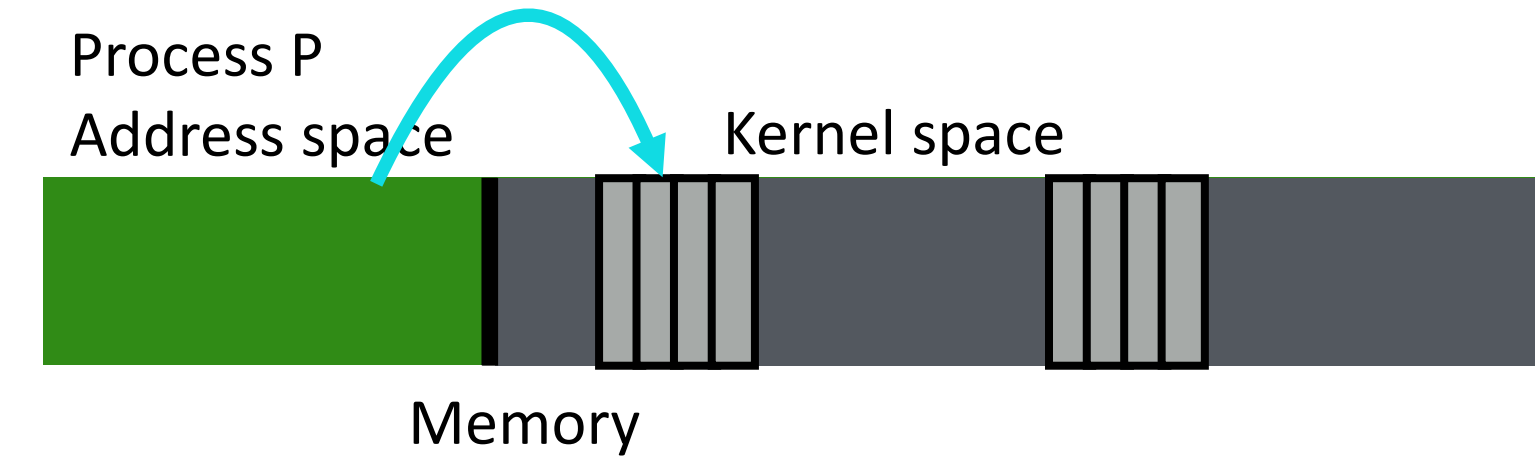
System Call



read():

```
movl $6, %eax;    int $64
```

System Call



read():

`movl $6, %eax;`

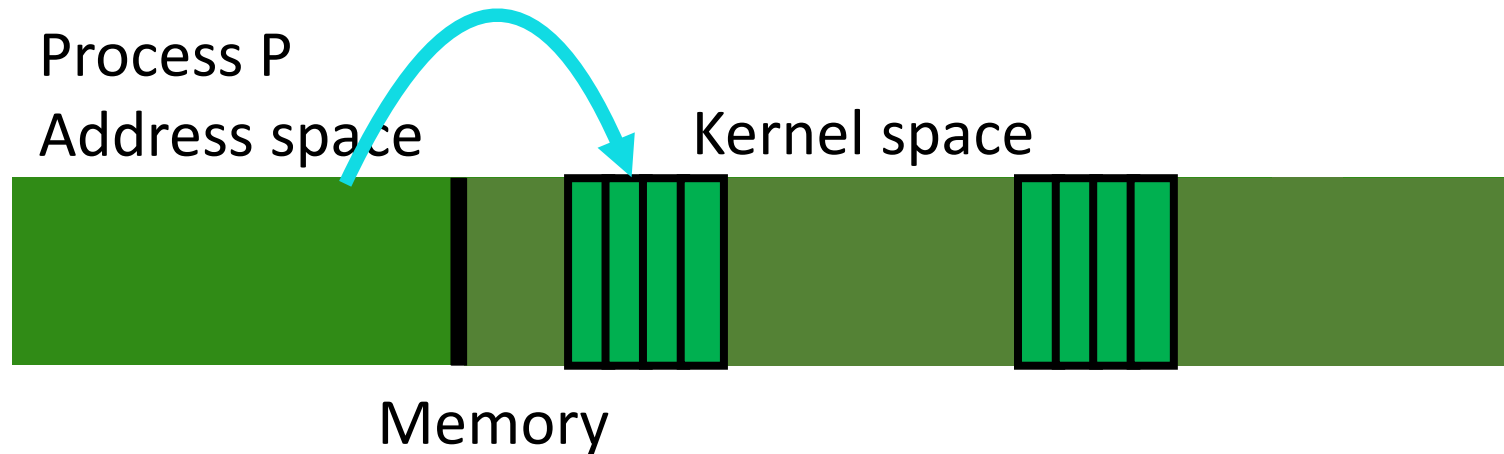
syscall-table index

`int $64`

trap-table index

System Call

Kernel mode: all are visible
and we can do anything



read():

`movl $6, %eax;`


syscall-table index

`int $64`

trap-table index

System Call

Kernel can access user memory to fill in user buffer;
return-from-trap at end to return to Process P



Limited direct execution protocol

OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from -trap

restore regs from kernel stack
move to user mode
jump to main

Run main()

...
Call system
trap into OS

Limited direct execution protocol

OS @ run
(kernel mode)

Hardware

Program
(user mode)

(Cont.)

Handle trap
Do work of syscall
return-from-trap

save regs to kernel stack
move to kernel mode
jump to trap handler

restore regs from kernel stack
move to user mode
jump to PC after trap

...
return from main
trap (via `exit()`)

Free memory of process
Remove from process list

What to limit?

- User process are not allowed to perform
 - General memory access
 - Disk I/O
 - Special x86 instructions like lidt
- What if process tries to do something restricted?



?

•

•

•

•

•

•