

# CPU Virtualization: Scheduling

Sridhar Alagar

# Sharing the CPU

- Mechanism - Dispatcher
  - How to switch to another process?
- Policy - Scheduler
  - Which process to switch to?

# Workload assumptions

1. Each job runs for the **same amount of time**
2. All jobs **arrive** at the same time
3. All jobs only use the **CPU** (i.e., they perform no I/O)
4. The **run-time** of each job is known

# Performance metric

- Turnaround time

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Other metrics?
  - Fairness
- Metrics can be conflicting with each other

# FIFO (or FCFS)

- Run jobs the order in which they arrived
  - Easy to implement
- Example

jobs	arrival time (s)	run time (s)
A	~0	10
B	~0	10
C	~0	10

# FIFO - Event trace

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	10
---	----	----

B	~0	10
---	----	----

C	~0	10
---	----	----

Time	Event
------	-------

0	A arrives
---	-----------

0	B arrives
---	-----------

0	C arrives
---	-----------

0	run A
---	-------

10	complete A
----	------------

10	run B
----	-------

20	complete B
----	------------

20	run C
----	-------

30	complete C
----	------------


# FIFO - Gantt Chart

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	10
---	----	----

B	~0	10
---	----	----

C	~0	10
---	----	----



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ secs}$$

# Workload assumptions

- ~~1. Each job runs for the **same amount of time**~~
2. All jobs **arrive** at the same time
3. All jobs only use the **CPU** (i.e., they perform no I/O)
4. The **run-time** of each job is known




# FIFO - Big job first

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	100
---	----	-----

B	~0	10
---	----	----

C	~0	10
---	----	----



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

# Convoy effect



# Passing the tractor

- Problems with FIFO
  - Short jobs have to wait for long jobs to finish
- New Scheduler - Shortest job first
  - choose the job with the smallest run time as the next job


# Shortest job first (SJF)

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	100
---	----	-----

B	~0	10
---	----	----

C	~0	10
---	----	----



$$\text{Average turnaround time} = \frac{120 + 10 + 20}{3} = 50 \text{ sec}$$

# Shortest job first (SJF)

- Moving shorter jobs before longer jobs improves the turnaround time for shorter jobs
- Moving longer jobs later does not affect the overall completion time
- SJF is optimal (provable)

# Workload assumptions

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
3. All jobs only use the **CPU** (i.e., they perform no I/O)
4. The **run-time** of each job is known


# SJF - different arrival times

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	100
---	----	-----

B	~10	10
---	-----	----

C	~10	10
---	-----	----



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.3 \text{ secs}$$

Stuck behind the tractor again!!

# Shortest time to completion first (STCF)

- When a new job arrives, check if it will complete sooner than the job running in the CPU
  - if so, switch to the new job
- Preemption is required
  - we already know how to switch context
- After completing a job, execute the job with shortest time




# STCF - different arrival times

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	100
---	----	-----

B	~10	10
---	-----	----

C	~10	10
---	-----	----




$$\text{Average turnaround time} = \frac{120 + (20 - 10) + (30 - 10)}{3} = 50 \text{ secs}$$

# Workload assumptions


- ~~1. Each job runs for the **same amount of time**~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (i.e., they perform no I/O)~~
4. The **run-time** of each job is known

# Not I/O aware



Poor use of resources

# I/O aware



- Treat job A as several separate CPU bursts
- When job A completes I/O, another job A is ready
- Each CPU burst is shorter than job B, so with SCTF, job A preempts job B

# New metric: Response time

- Important when a job is started than when it is finished
  - interactive jobs

$$T_{response} = T_{firstrun} - T_{arrival}$$

- FCFS, SJF, STCF are not good at minimizing response time


# Response time vs turnaround time

jobs	arrival time (s)	run time (s)
------	------------------	--------------

A	~0	100
---	----	-----

B	~10	10
---	-----	----

C	~10	10
---	-----	----




*response time for C = (20 - 10) = 10 secs*

How to minimize the response time?

# Round robin scheduler

- Alternate ready processes every fixed-length time-slice
- time-slice also called as time-quantum
- Length of quantum could be a multiple of timer interrupt period


# RR vs SJF scheduling



**SJF (Bad for Response Time)**

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5\ secs$$

$$T_{average\ turnaround} = \frac{5 + 10 + 15}{3} = 10\ secs$$



**RR with a time-slice of 1sec (Good for Response Time)**

$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1\ secs$$

$$T_{average\ turnaround} = \frac{13 + 14 + 15}{3} = 14\ secs$$



# Length of time-slice is critical

- Shorter time-slice
  - better response time
  - too many context switching. Overhead becomes high
- Longer time-slice
  - few context switching
  - poorer response time
- A trade-off

# Another benefit of round robin

- Run time need not be known
- In fact, RR does not care whether a process terminates or not

# Workload assumptions

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (i.e., they perform no I/O)~~
- ~~4. The run-time of each job is known~~

# Disclaimer

- Some of the materials in this lecture slides are from the materials prepared by Prof. Arpaci, and Prof. Youjip. Thanks to all of them.