

Virtualizing Memory

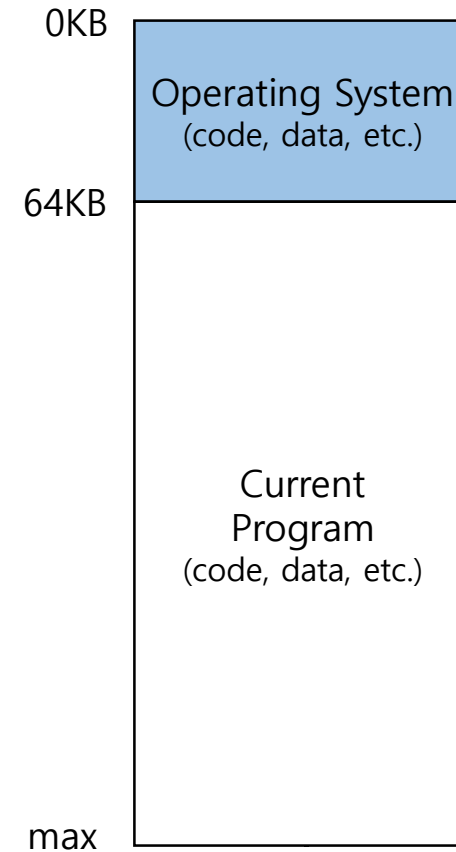
Sridhar Alagar

Virtualization


- CPU
 - Abstraction: Process
 - Sharing:
 - Policy
 - Mechanism
 - Isolation
 - Limited direct execution
- Memory
 - Abstraction
 - Sharing
 - Isolation

Early OS

- Uniprogram: load only one program
- Drawback?
 - poor utilization
 - process can destroy OS



Multiprogramming




Physical Memory

Multiprogramming Goals

- Transparency
 - processes are not aware of sharing
 - works regardless of where it is loaded in memory
- Protection
 - cannot access OS or other process' memory
- Efficient
- Share memory among cooperating process

Abstraction: Address Space

- Each process has a set of addresses that map to bytes
- What is in an address space?
- Static: code and global variables
- Dynamic: heap and stack



Quiz: Find that address location

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

- Possible segments: static data, code, stack, heap

Address	Location
x	Static data
main	Code
y	Stack
z	Stack
*z	Heap

Example: Memory Access

```
void func()  
    int x;  
    ...  
    x = x + 3; //this is the line of code we are interested
```

Load a value from memory

Increment it by three

Store the value back into memory

Example: Memory access


```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax
132 : addl $0x03, %eax          ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)      ; store eax back to mem
```

Load a value from memory

Increment it by three

Store the value back into memory

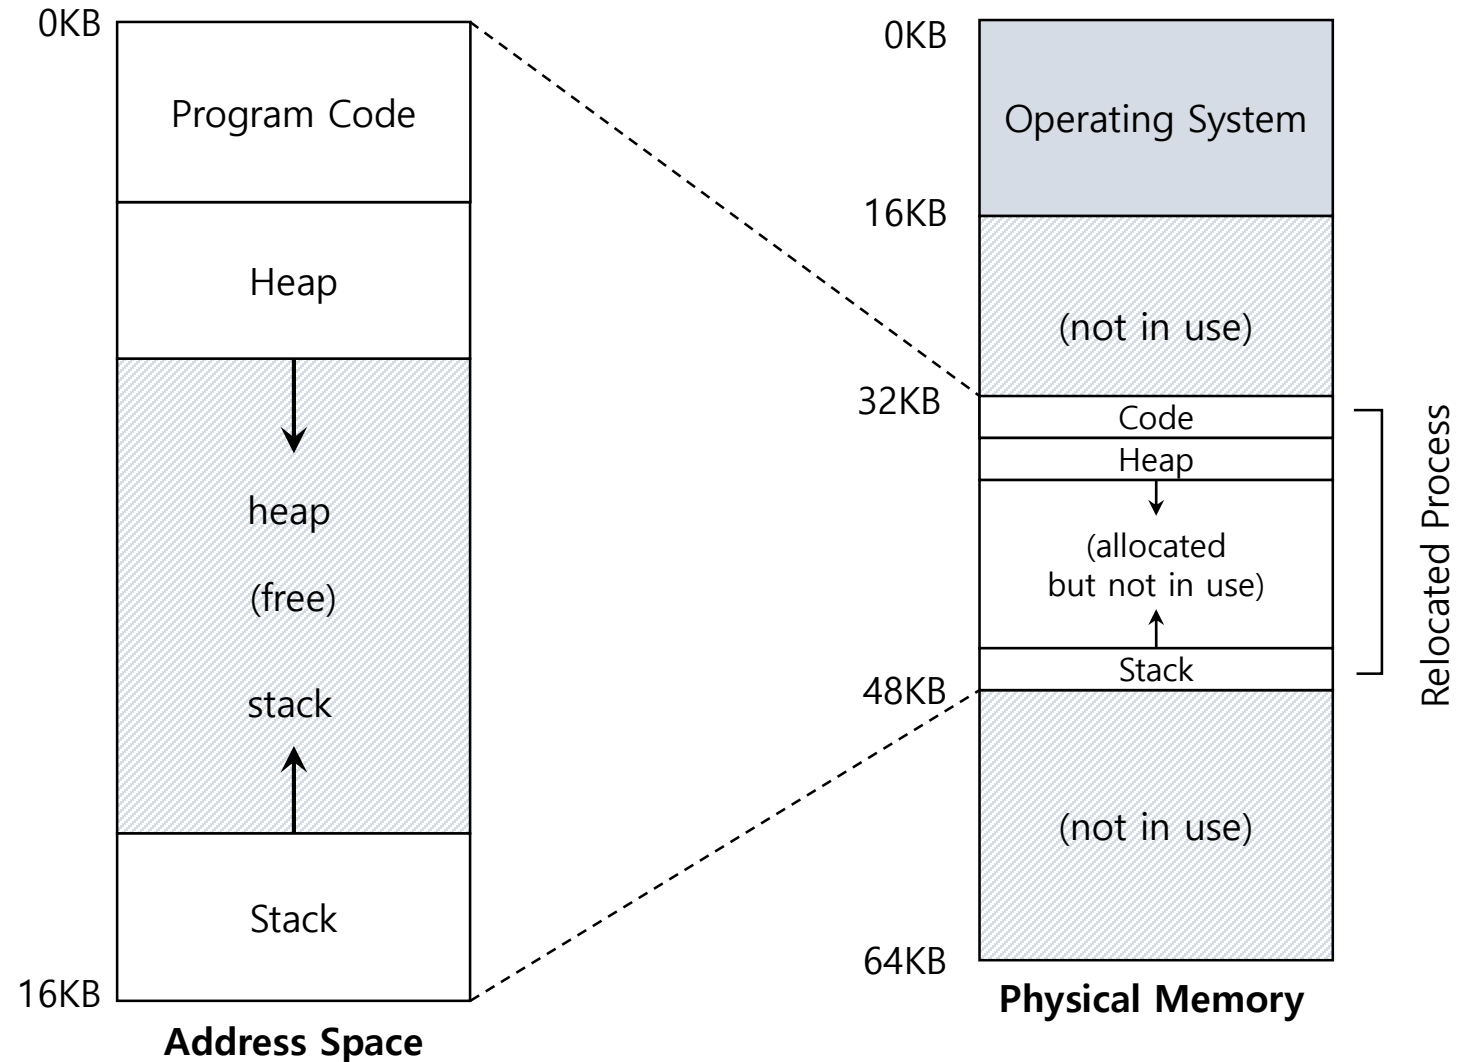
Example: Memory access



- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

Relocation

- OS wants to relocate the process

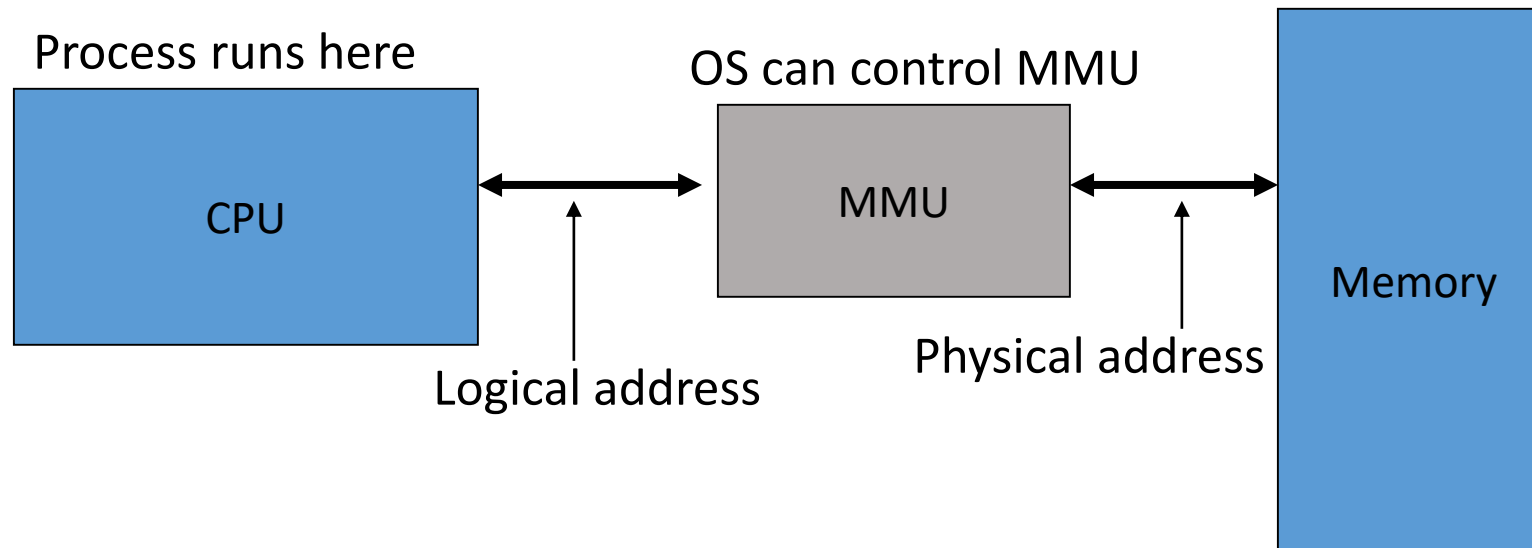


Dynamic relocation

- Process generates virtual address
- Memory access needs real address
- Every virtual address referenced by a process must be dynamically translated to a real address


Dynamic relocation

- H/W does the translation: Memory Management Unit (MMU)




Dynamic relocation: Base register


- MMU contains base register and adds its value to virtual address




Dynamic relocation: Example



Dynamic relocation: Example



Dynamic relocation: Example

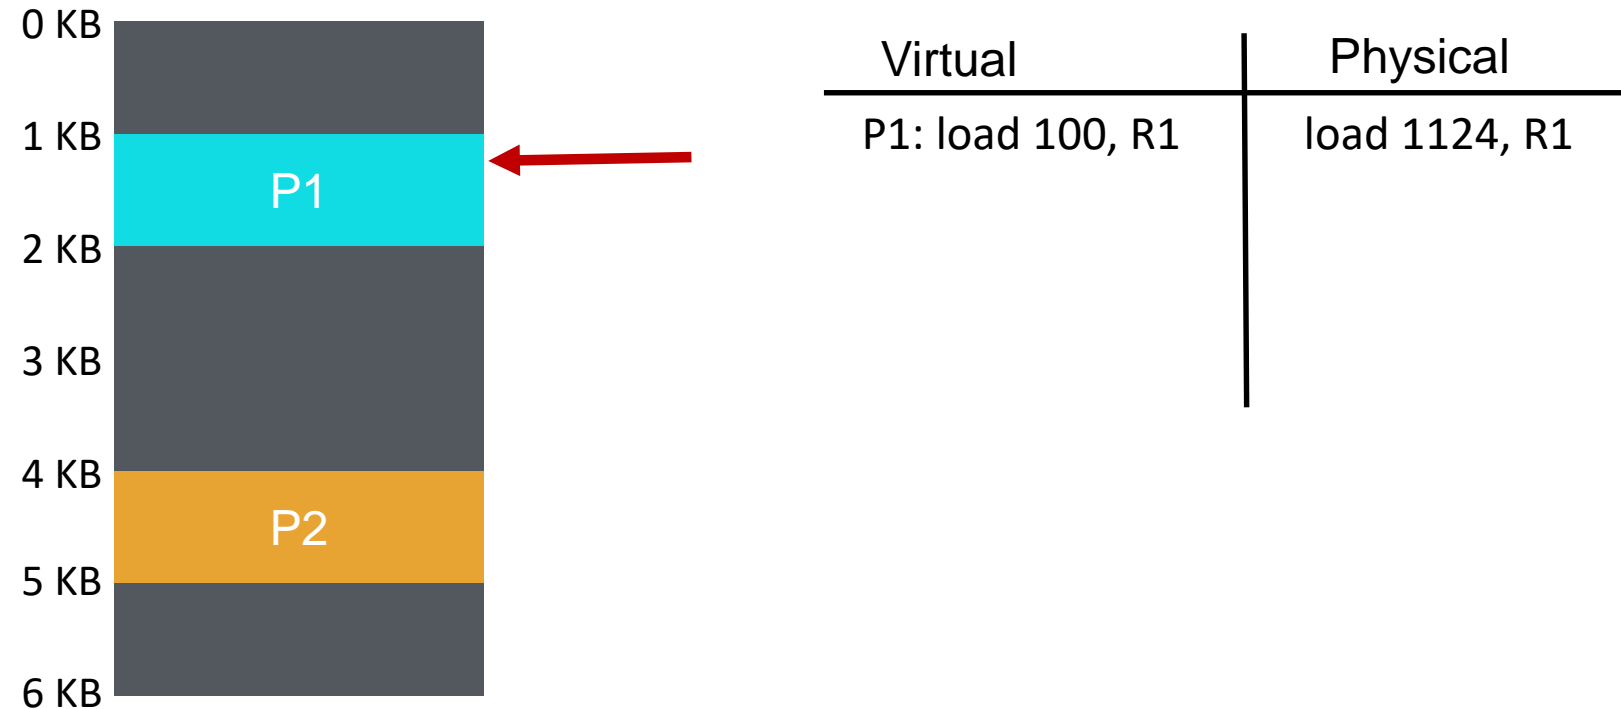


Dynamic relocation: Example



Virtual	Physical
P1: load 100, R1	

Dynamic relocation: Example




Dynamic relocation: Example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	

Dynamic relocation: Example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1

Dynamic relocation: Example



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1

Dynamic relocation: Example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1

Who controls base register?

- Who should translate the address with base register?
 - (a) process
 - (b) OS
 - (c) h/w
- Who should modify the base register?
 - (a) process
 - (b) OS
 - (c) h/w

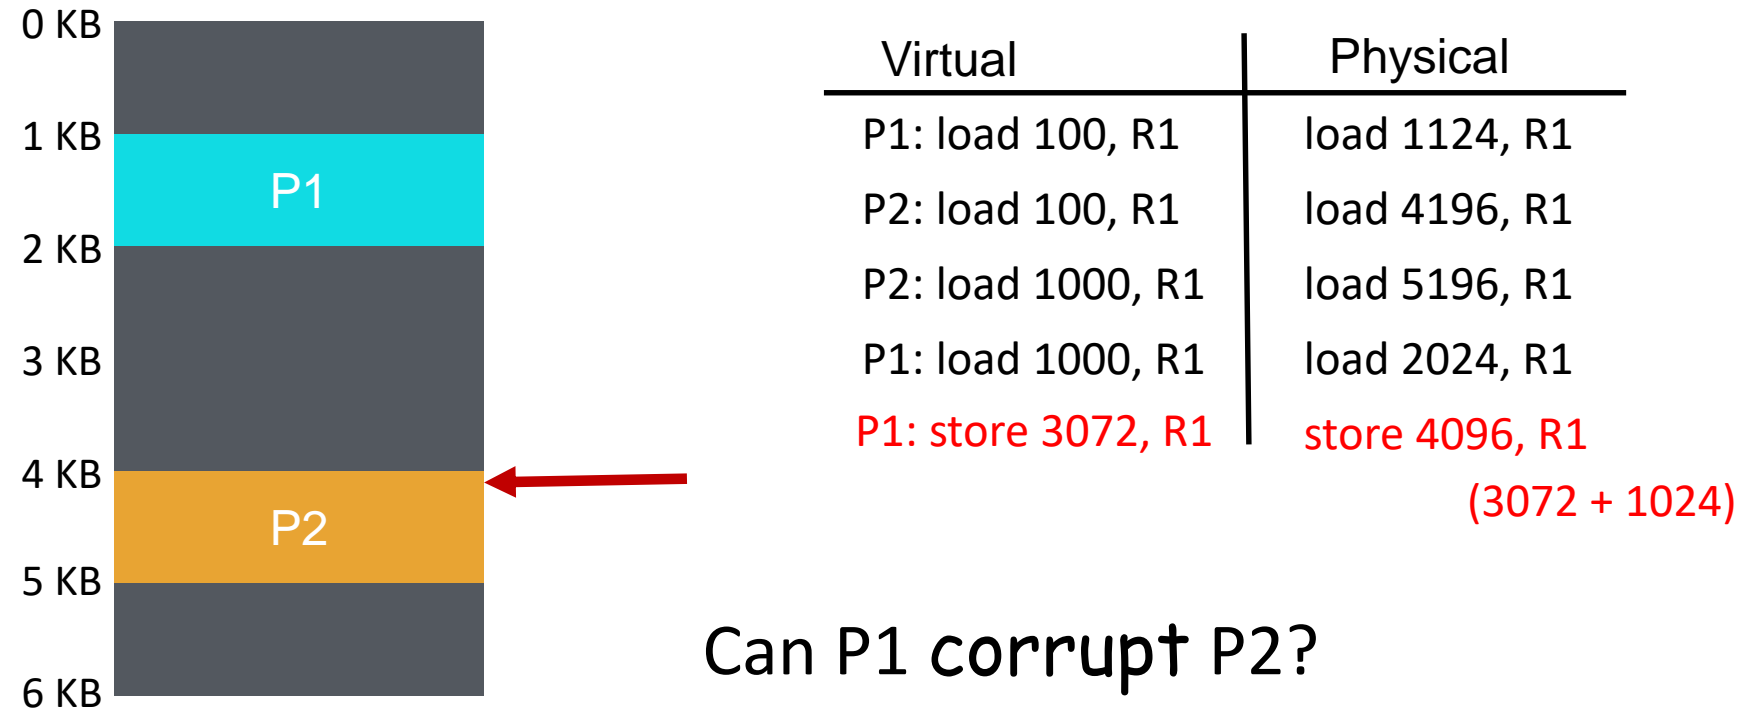
Dynamic relocation: Protection



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1

Can P1 corrupt P2?

Dynamic relocation: Protection

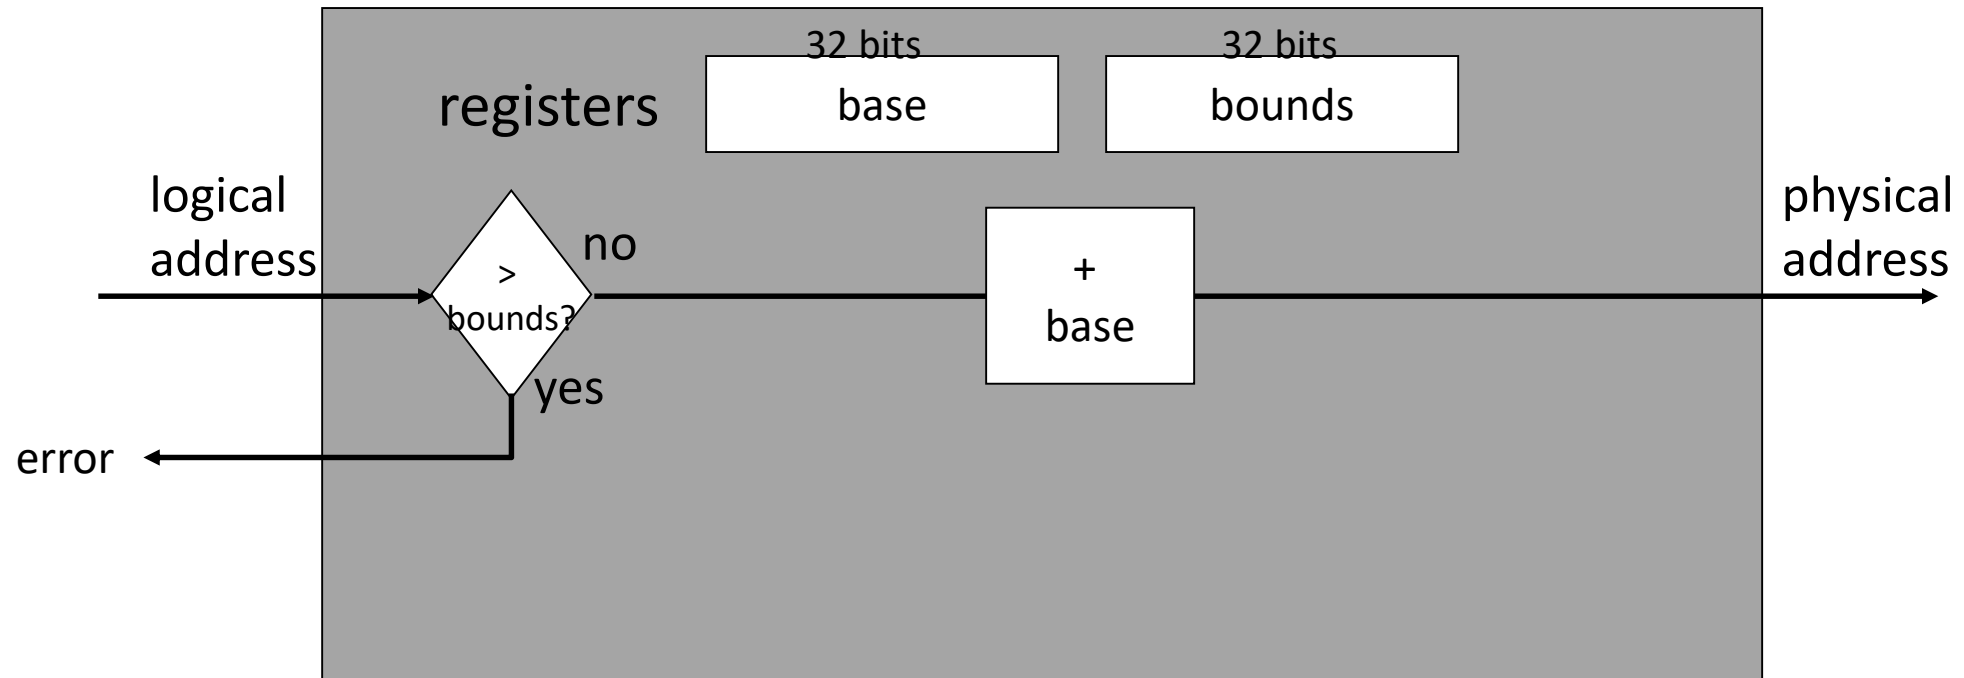


Dynamic relocation with base register is not adequate for protection.


Dynamic relocation: Bounds register

- Limit the address space with bounds register
- Base register: starting location (smallest physical addr)
- Bounds register: size of the process virtual space
 - could be largest physical address (base + size)
- What should OS do if a process load/stores beyond bounds?

MMU with base and bounds register



Dynamic relocation: Protection with bounds



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	interrupt OS
	(3072 > 1024)

Can P1 corrupt P2?

Context Switching

- What to do with base and bounds registers during context switch?
 - base and bound part of process' context
 - save them in PCB of old process
 - restore them from PCB of new process

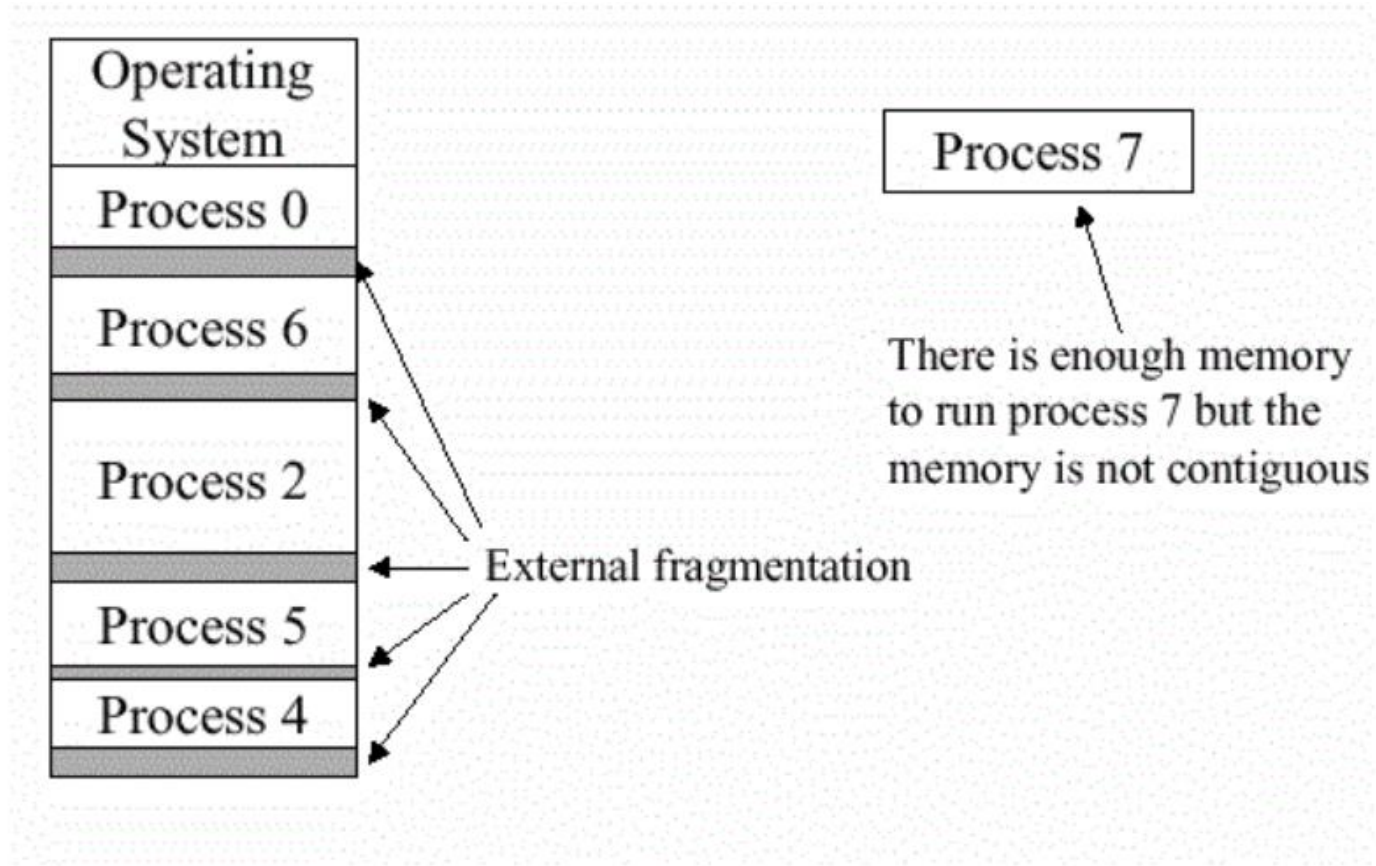
Base and Bounds: Advantages

- Dynamic relocation
- Protection
- Simple to implement in h/w

Base and Bounds: Disadvantages



- Process requires contiguous memory space
- A process may not use all the space allocated to it
 - internal fragmentation
- There may not be enough memory that are contiguous, but total free memory available is greater than what a process needs
 - external fragmentation

External Fragmentation





Segmentation

- Divide a process' address space into logical segments
 - each segment corresponds to a logical entity
- Each segment can be
 - independently placed separately in memory
 - grow and shrink (varied size)
 - be protected



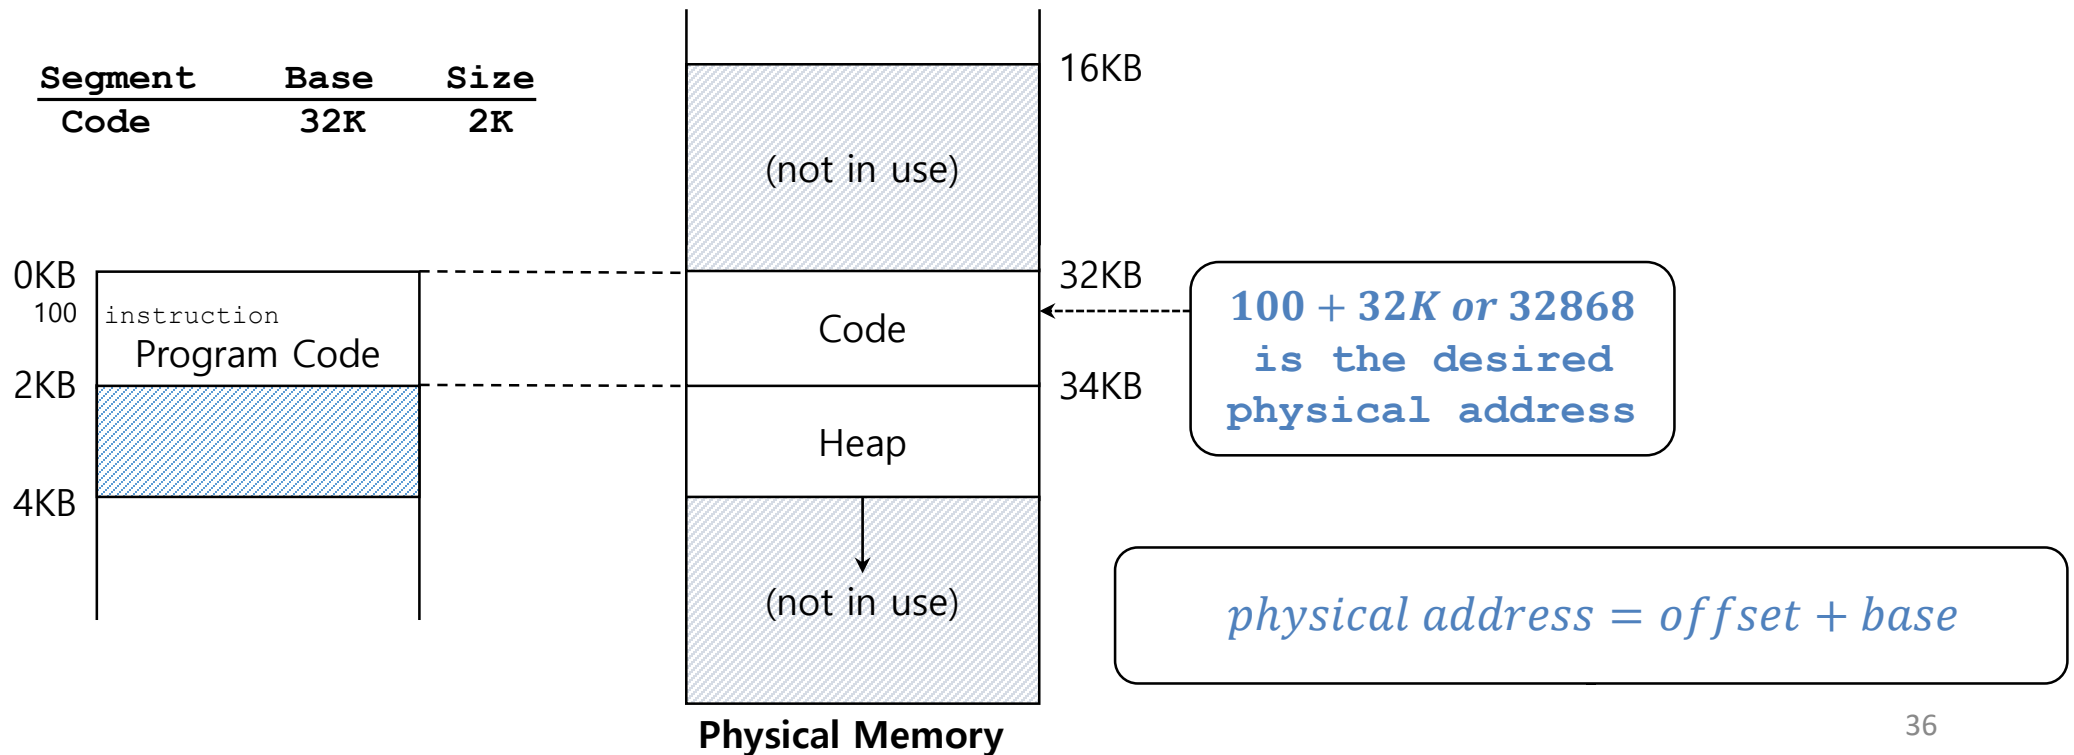
Placing Segment in Memory



Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

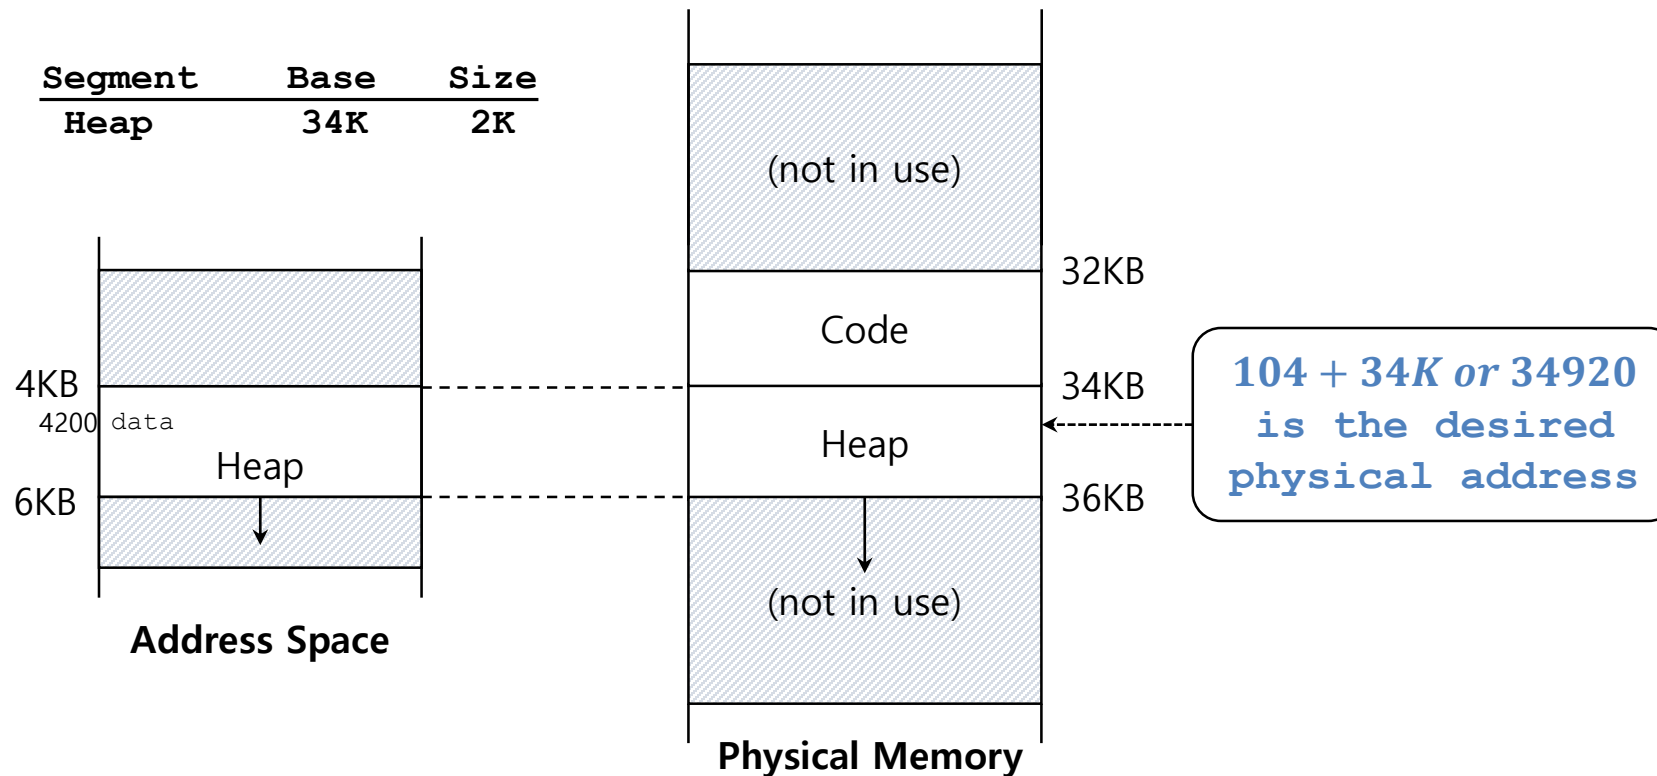
Address translation in Segmentation

- The offset of virtual address 100 is 100
 - The code segment starts at virtual address 0 in address space




Address translation in Segmentation

- The offset of virtual address 4200 is 104
 - The heap segment starts at virtual address 4096



Which segment are we referring to?


- Explicit approach: use part of logical address
 - The top bits specify the segment
 - The remaining bits specify the offset within segment



Which segment are we referring to?

- Example: virtual address 4200 (010000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11

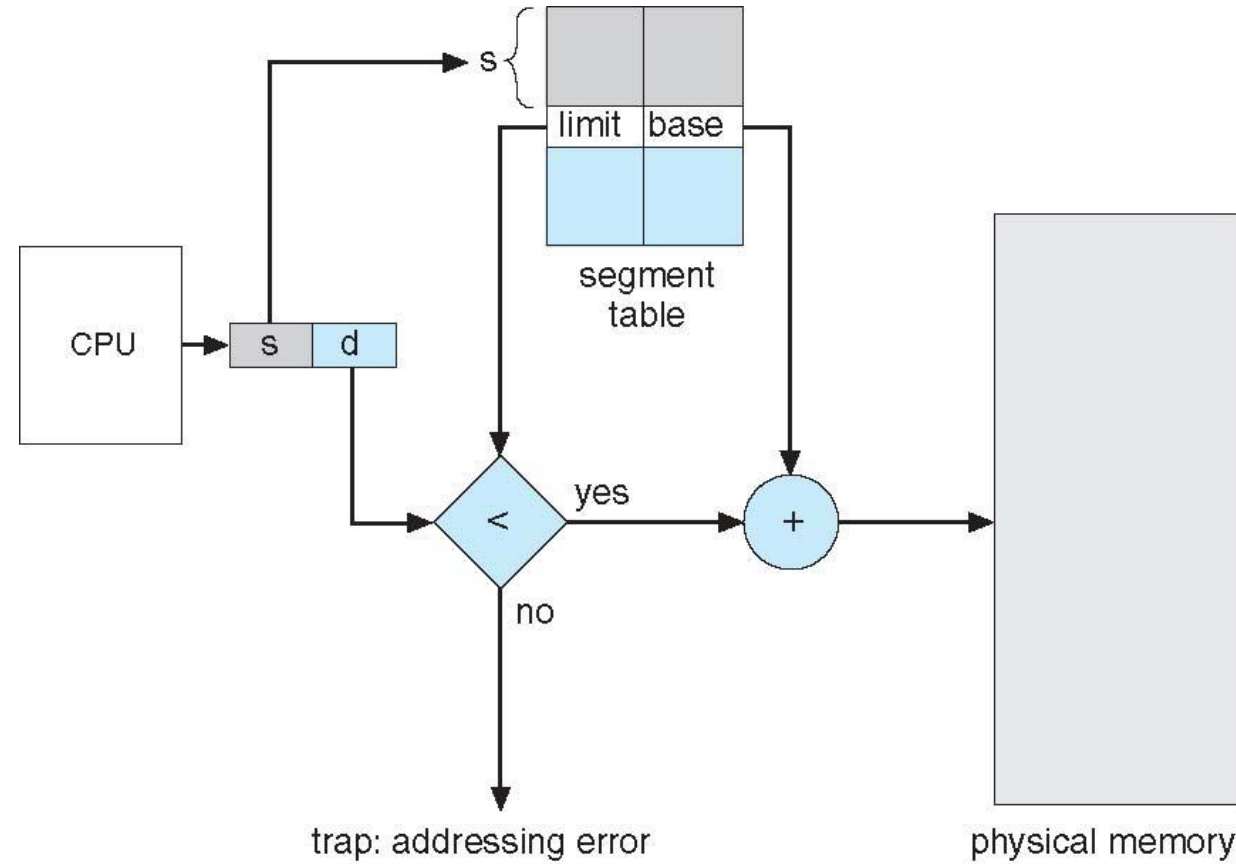


How to find segment and offset?

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- SEG_MASK = 0x3000 (1100000000000000)
- SEG_SHIFT = 12
- OFFSET_MASK = 0xFFF (0011111111111111)

Segmentation Hardware



Sharing segment

- Add permission/protection bits in segment table

Segment	Base	Size	Protection
Code	32K	2K	Read-Execute
Heap	34K	2K	Read-Write
Stack	28K	2K	Read-Write

Segmentation

- Advantages
 - Segment sharing
 - Easier to relocate segment than entire program
- Disadvantages
 - Segments have variable length
 - Segments can be large
 - Fragmentation not solved

Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.