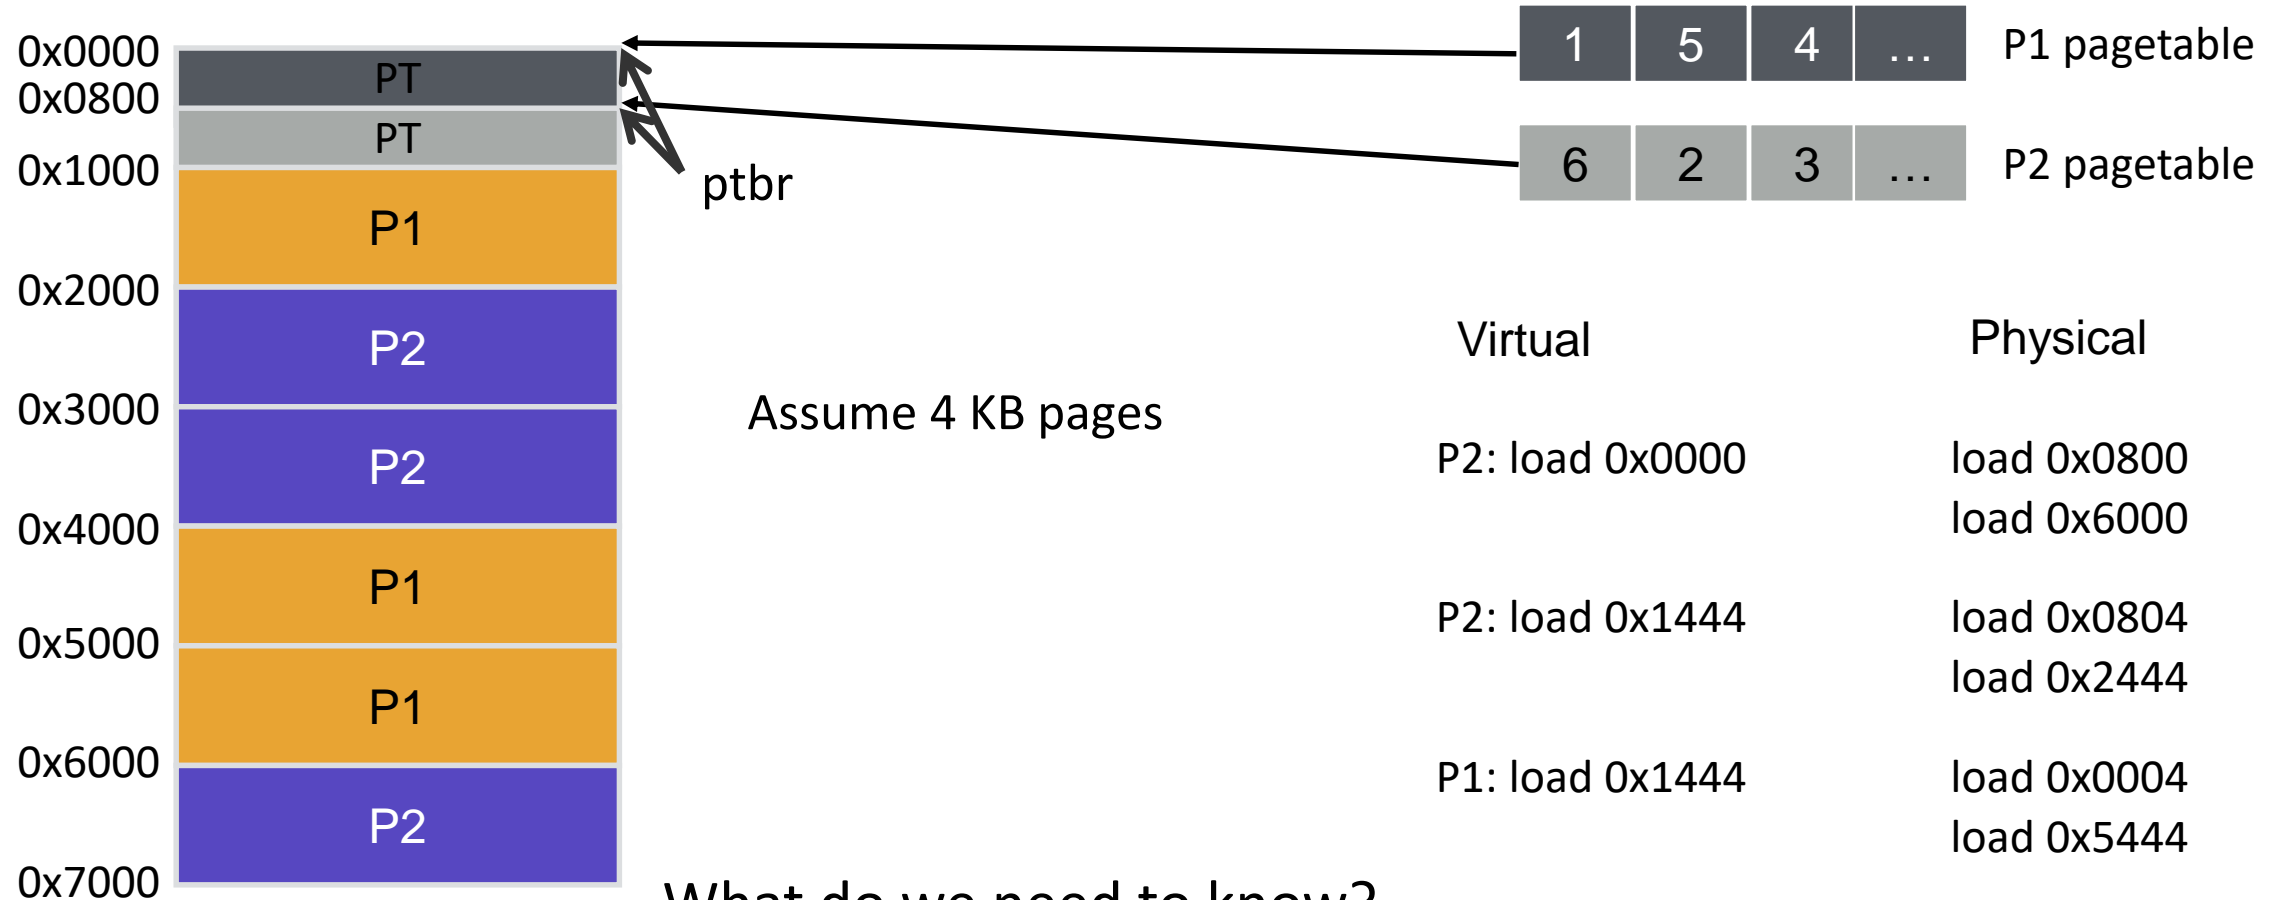


# VM: TLBs

Sridhar Alagar

# Review: Paging



What do we need to know?

Location of page table in memory (ptbr)

Size of each page table entry (assume 4 bytes)

# Review: Paging Pros and Cons

- Advantages
  - No external fragmentation
  - All free pages are equivalent
- Disadvantages
  - Page tables are too big
    - Must have one entry for every page of address space
  - Accessing page tables is too slow
  - Doubles number of memory references per instruction

# Example: Array

```
int sum = 0;
for (i=0; i<N; i++) {
    sum += a[i];
}
```

Assume 'a' starts at 0x3000

Ignore instruction fetches

virtual addresses

load 0x3000

load 0x3004

load 0x3008

load 0x300C

...

physical addresses?

load 0x100C

load 0x7000

load 0x100C

load 0x7004

load 0x100C

load 0x7008

load 0x100C

load 0x700C

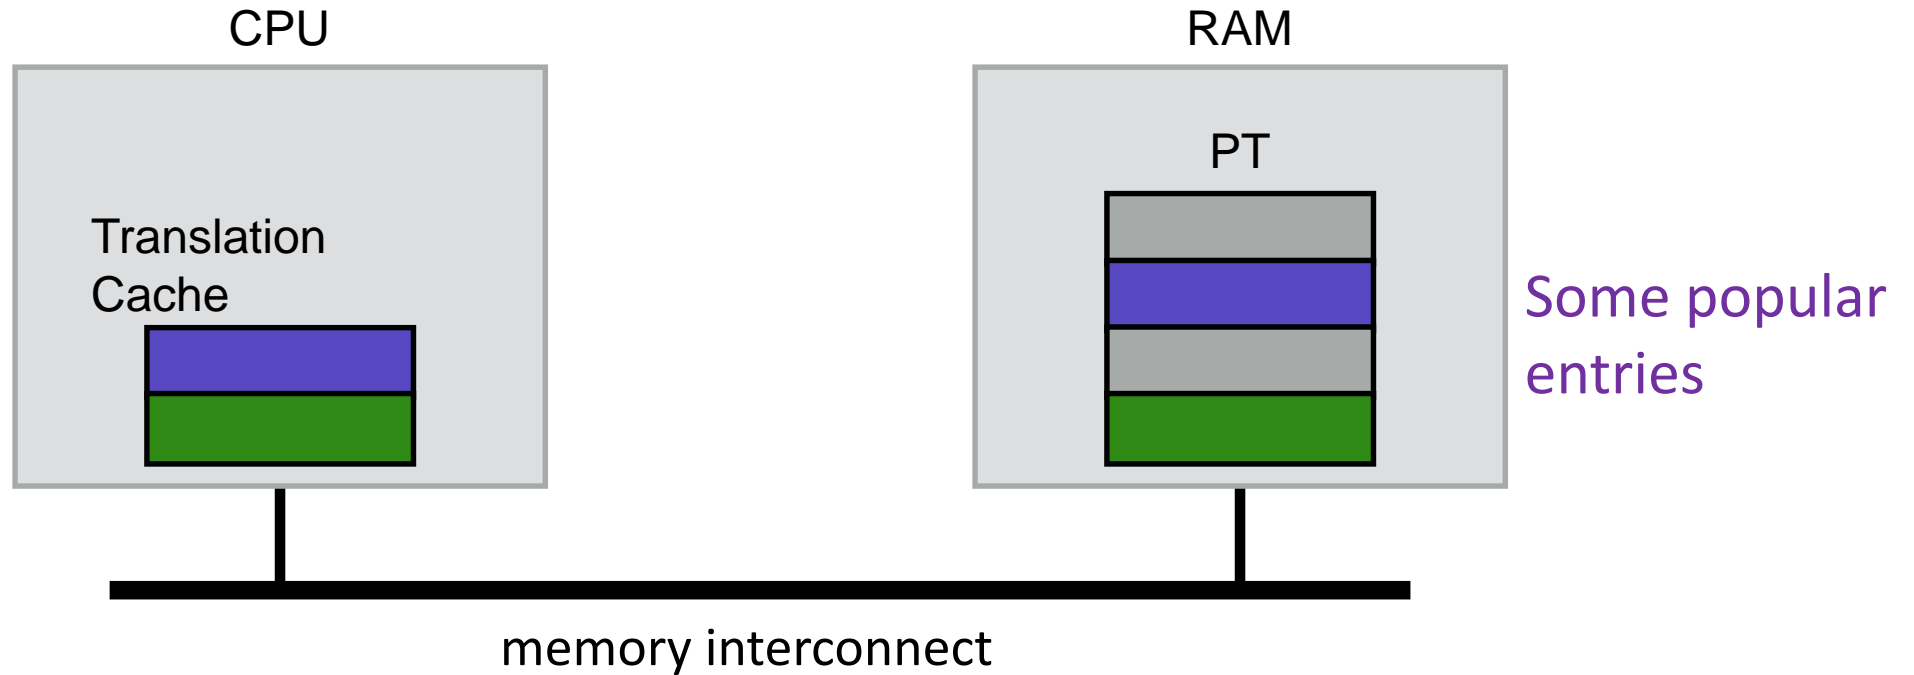
Observation:

Repeatedly access same PTE  
because program repeatedly  
accesses same virtual page

What can you infer?

- ptbr: 0x1000; PTE 4 bytes each
- VPN 3 -> PPN 7

# Cache Page Translations



## TLB: Translation Lookaside Buffer

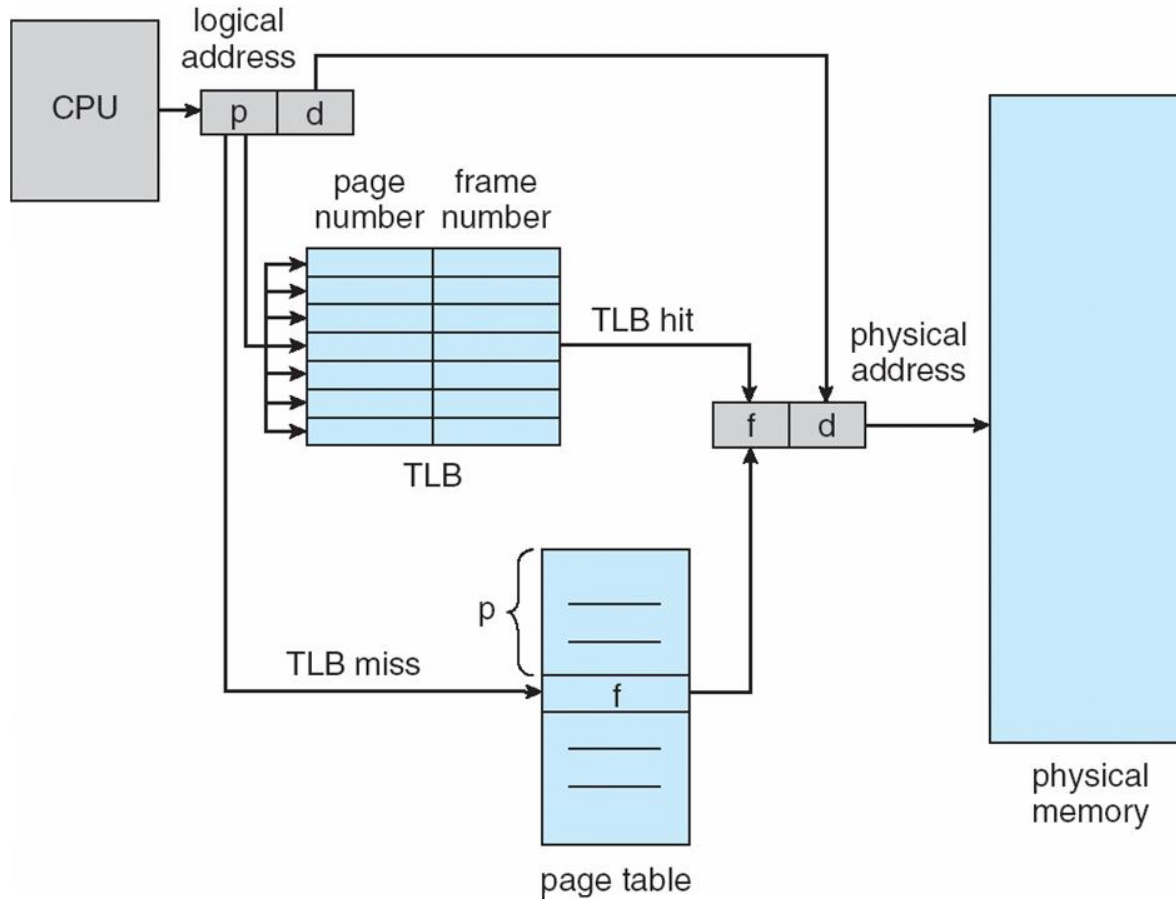
# TLB is Fully Associative

- Associative memory

Page #	Frame #
p1	f1

- Address translation (p, f)
  - If p1 is in associative register, get frame # f1 out
  - Do parallel search

# Paging with TLB Hardware



# TLB Example - Iterate Array

```
int sum = 0;  
for(i = 0; i < 2048; i++) {  
    sum += a[i];  
}
```

Assume 'a' starts at VA 0x1000:

load 0x1000

load 0x1004

load 0x1008

load 0x100C

...

What will TLB behavior look like?



# TLB Example

Pagetable @ 0x0000

1	5	4	...
0	1	2	3

CPU's TLB

VPN	PPN
1	5
2	4

Virt	Phys
load 0x1000	load 0x0004
load 0x1004	load 0x5000 (TLB hit)
	load 0x5004 (TLB hit)
load 0x1008	load 0x5008 (TLB hit)
load 0x100c	load 0x500C
...	...
load 0x2000	load 0x0008
	load 0x4000 (TLB hit)
load 0x2004	load 0x4004

# TLB Performance

- miss rate =  $\# \text{TLB miss} / \# \text{TLB Lookup}$
- In the sequential array access example:
  - # miss = 2
  - # lookup = 2048
  - miss rate =  $2/2048 = 0.1\%$
  - hit rate = 99.9%
- What access pattern will be bad for hit rate?
- TLB Performance vs Page size?

# Locality of reference

- Spatial Locality: future reference to nearby address
- Temporal Locality: future reference to the same data
- Both helps TLB hit rate

# Who handles LB miss?

- H/W
  - Should know where page table is stored (CR3 in x86)
  - PT structure fixed and agreed between H/W and OS
  - Walk the page table and fetch the appropriate entry
- OS. How will it know the TLB miss?
  - Interrupt
  - Software managed TLB

# Replacement Policies

- If TLB is full, which entry should be replaced?
- LRU: replace the least recently used
- Random: replace randomly chosen entry
- Which is better?


# Context Switching

- Any problem for TLB?
  - TLB contain previous process entries
- Flush the old entries during context switch
  - will loose all cached translation
  -
- Share TLB across context switch
  - Address space identifier (ASID)
  - Tag each entry with ASID

# Other bits in TLB entry

- valid/invalid

All 64 bits of this TLB entry(example of MIPS R4000)



Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support with up to 64GB of main memory( $2^{24} * 4KB$ pages ).
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.

# Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.