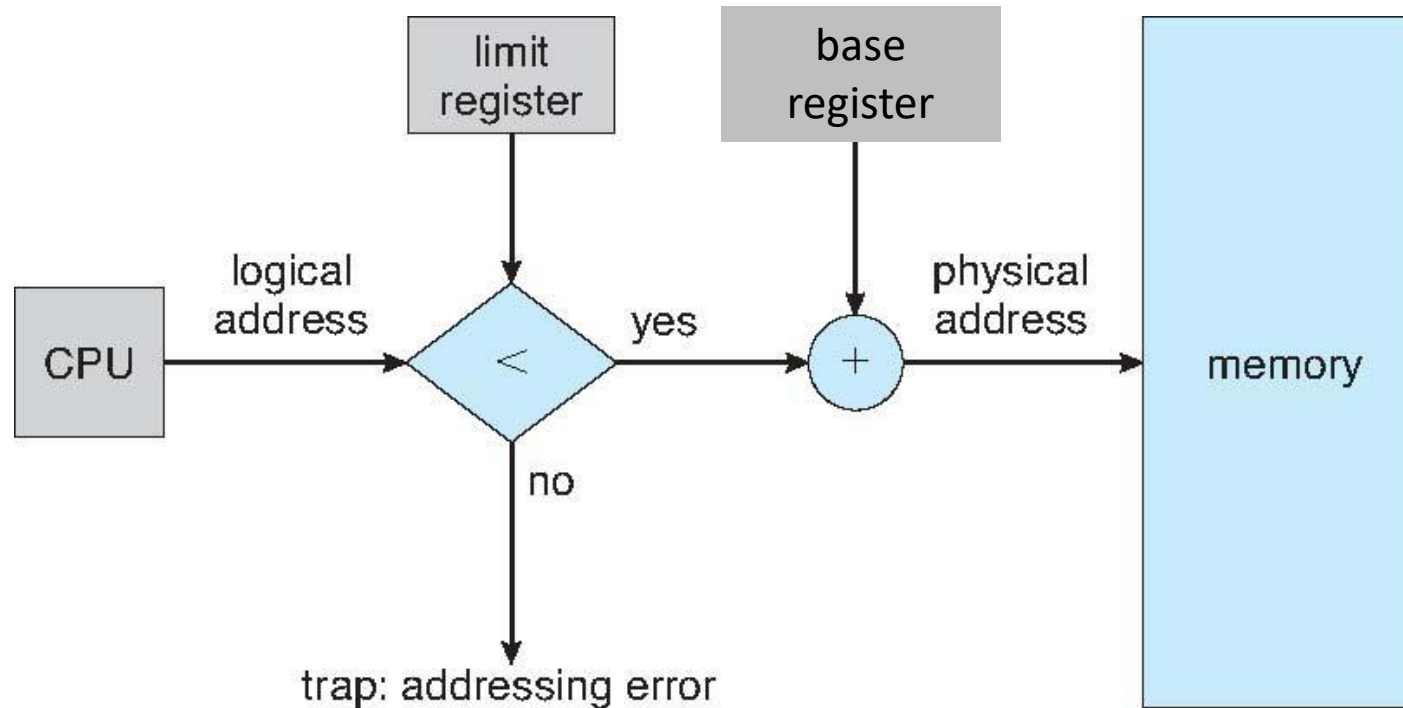


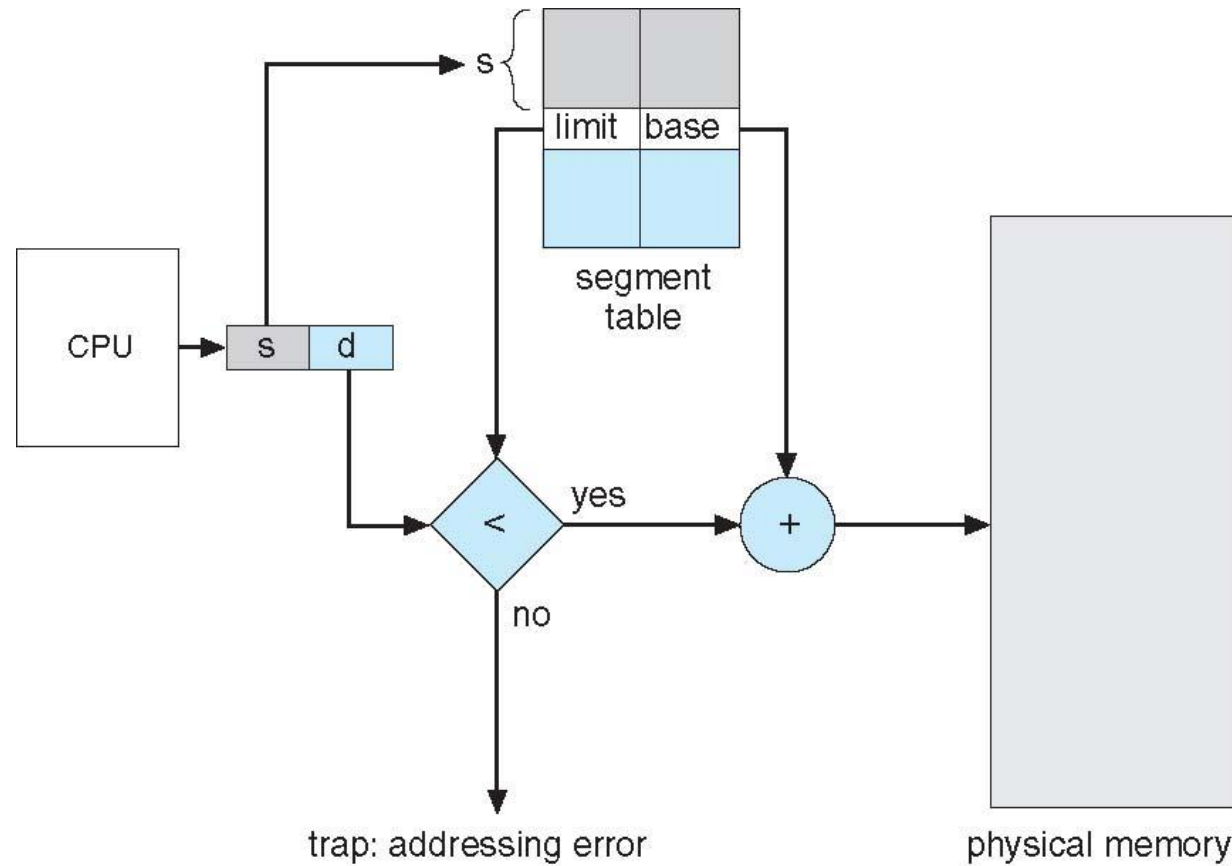
Virtualizing Memory: Smaller Page Tables

Sridhar Alagar

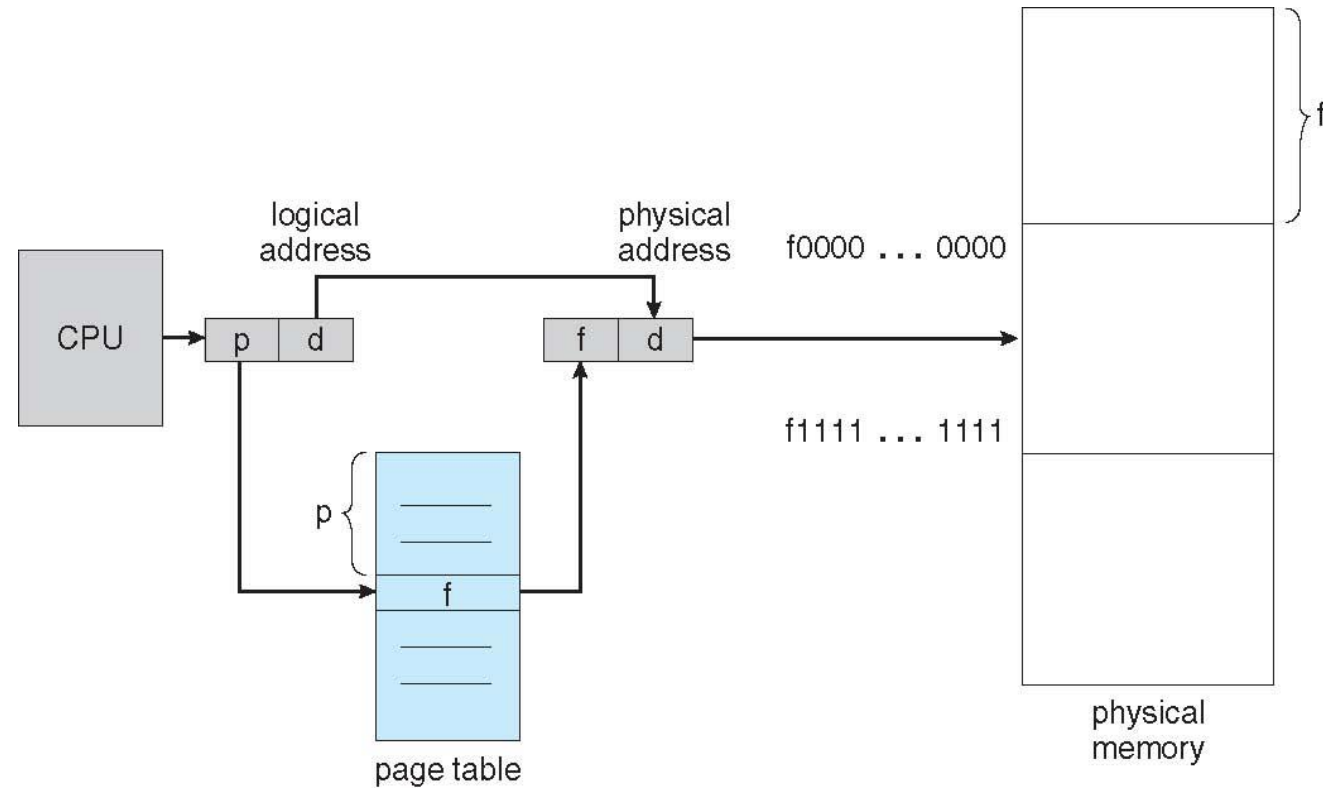
Translation with Base and Limit



Segmentation



Paging



Disadvantages of Paging

1. Additional memory reference to look up in page table (**inefficient**)
 - Page table must be stored in memory
 - MMU stores only base address of page table
 - Avoid extra memory reference for lookup with TLBs
2. Storage for page tables may be substantial
 - Simple page table: Requires PTE for all pages in address space
 - Entry needed even if page not allocated
 - Requires contiguous memory space

How big are page tables?

1. PTE's are **2 bytes**, and **32** possible virtual page numbers

$$32 * 2 \text{ bytes} = 64 \text{ bytes}$$

2. PTE's are **2 bytes**, virtual addr are **24 bits**, pages are **16 bytes**

$$2 \text{ bytes} * 2^{(24 - \lg 16)} = 2^{21} \text{ bytes (2 MB)}$$

3. PTE's are **4 bytes**, virtual addr are **32 bits**, and pages are **4 KB**


$$4 \text{ bytes} * 2^{(32 - \lg 4K)} = 2^{22} \text{ bytes (4 MB)}$$

4. PTE's are **4 bytes**, virtual addr are **64 bits**, and pages are **4 KB**

$$4 \text{ bytes} * 2^{(64 - \lg 4K)} = 2^{54} \text{ bytes}$$

Problem?

- Most of the page table is **unused**, full of invalid entries.



A 16KB Address Space with 1KB Pages

PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

A Page Table For 16KB Address Space

Avoid storing invalid PTEs

PFN	valid	prot
10	1	r-x
-	0	-
23	1	rw-
-	0	-
-	0	-
-	0	-
-	0	-
...many more invalid...		
-	0	-
-	0	-
-	0	-
-	0	-
28	1	rw-
4	1	rw-


how to avoid
storing these?

How did OS avoid
allocating holes in
physical memory?

Segmentation

Combine Paging and Segmentation

- Divide address space into segments (code, heap, stack)
- Divide each segment into fixed sized pages
- Logical address divided into 3 parts



Implementation Paging and Segmentation

- Base of segment
 - points to beginning of page table of that segment
- Bounds of segment
 - end/size of page table

Quiz: Paging and Segmentation

seg # (4 bits)	page number (8 bits)	page offset (12 bits)
-------------------	----------------------	-----------------------

seg	base	bounds	R W
0	0x002000	0xff	1 0
1	0x000000	0x00	0 0
2	0x001000	0x0f	1 1

0x002070 read: 0x004070
 0x202016 read: 0x003016
 0x104c84 read: error
 0x010424 write: error
 0x203568 read: 0x02a568

...	
0x01f	0x001000
0x011	
0x003	
0x02a	
0x013	
...	
0x00c	0x002000
0x007	
0x004	
0x00b	
0x006	
...	

Advantages of Paging and Segmentation

- Advantages of Segments
 - Supports sparse address spaces
 - Decreases size of page tables
 - If segment not used, no need for page table
- Advantages of Pages
 - No external fragmentation
 - Segments can grow without any reshuffling
- Advantages of Both
 - Increases flexibility of sharing
 - Share either single page or entire segment. How?

Disadvantages of Paging and Segmentation

- Potentially large page tables (for each segment)
 - Must allocate each page table contiguously
 - More problematic with more address bits
 - Page table size?
 - Assume 2 bits for segment, 18 bits for page number, 12 bits for offset

Each page table size:

= Number of entries * size of each entry

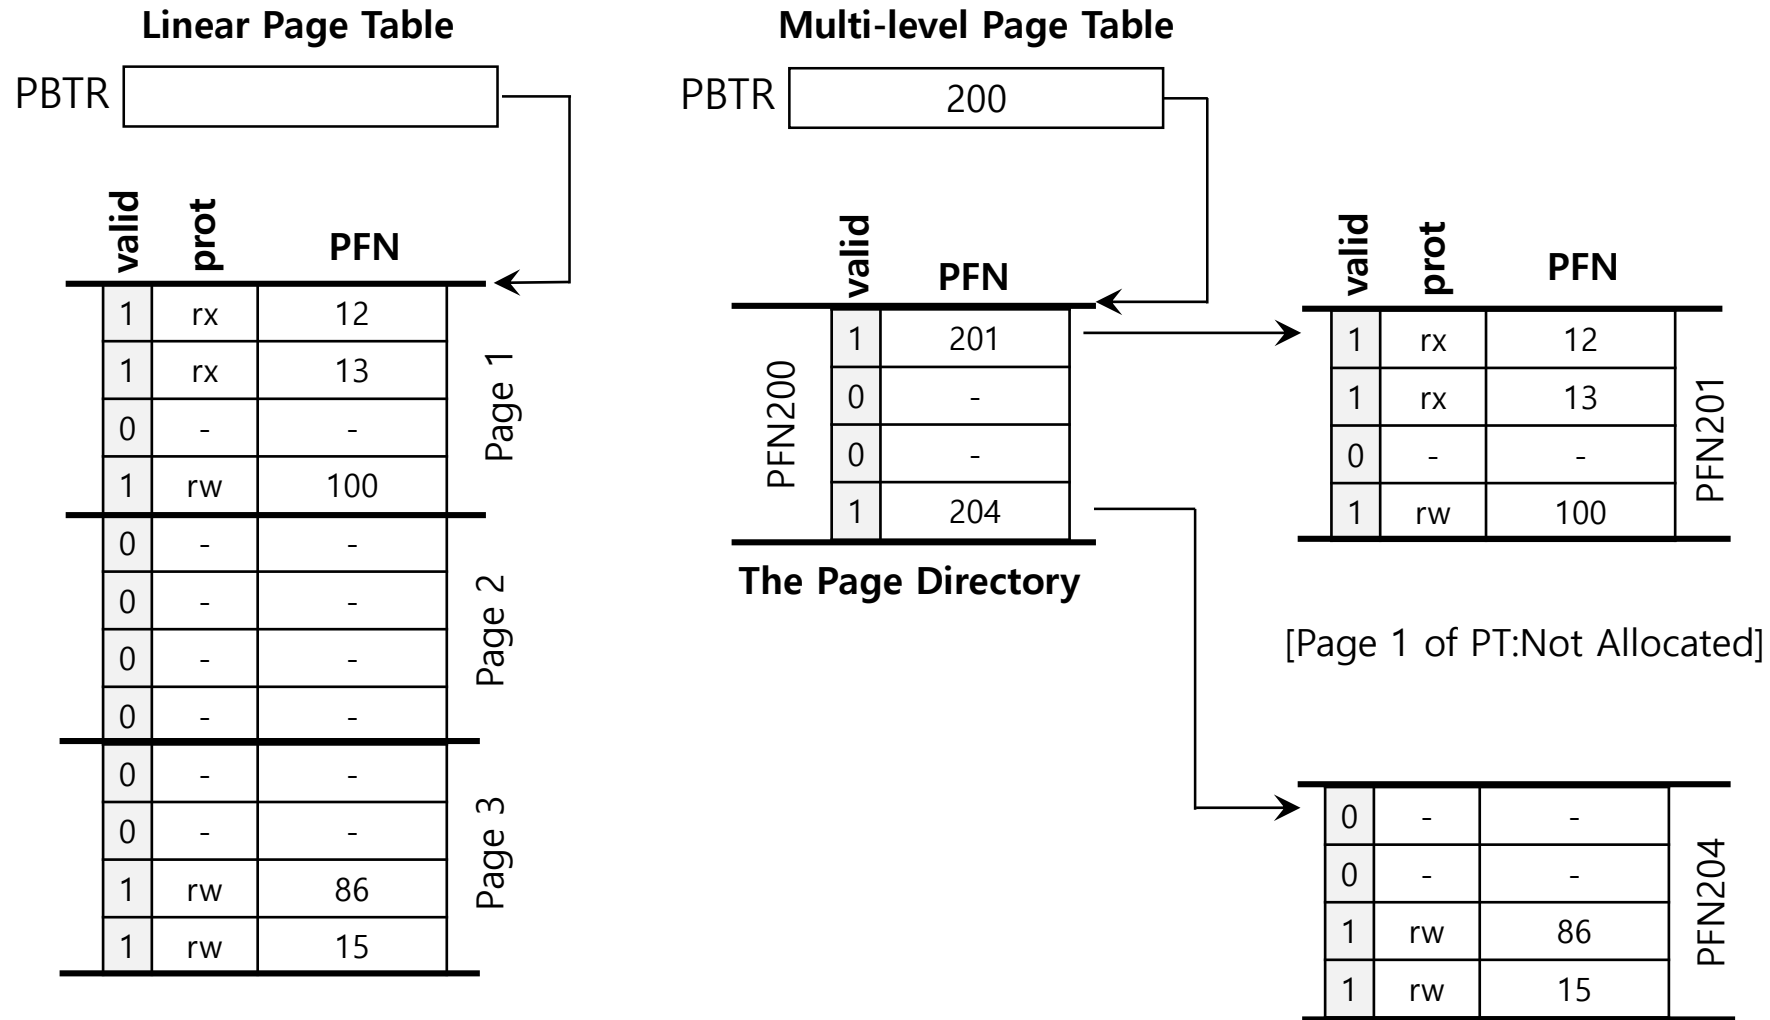
= Number of pages * 4 bytes

= $2^{18} * 4 \text{ bytes} = 2^{20} \text{ bytes} = 1 \text{ MB!!!}$

How to allocate memory for page table?


- Physical page size is 4K
- Page table size can be > 1MB
- How to allocate memory non-contiguously?
- Page the page table

Page the page table



Linear (Left) And Multi-Level (Right) Page Tables

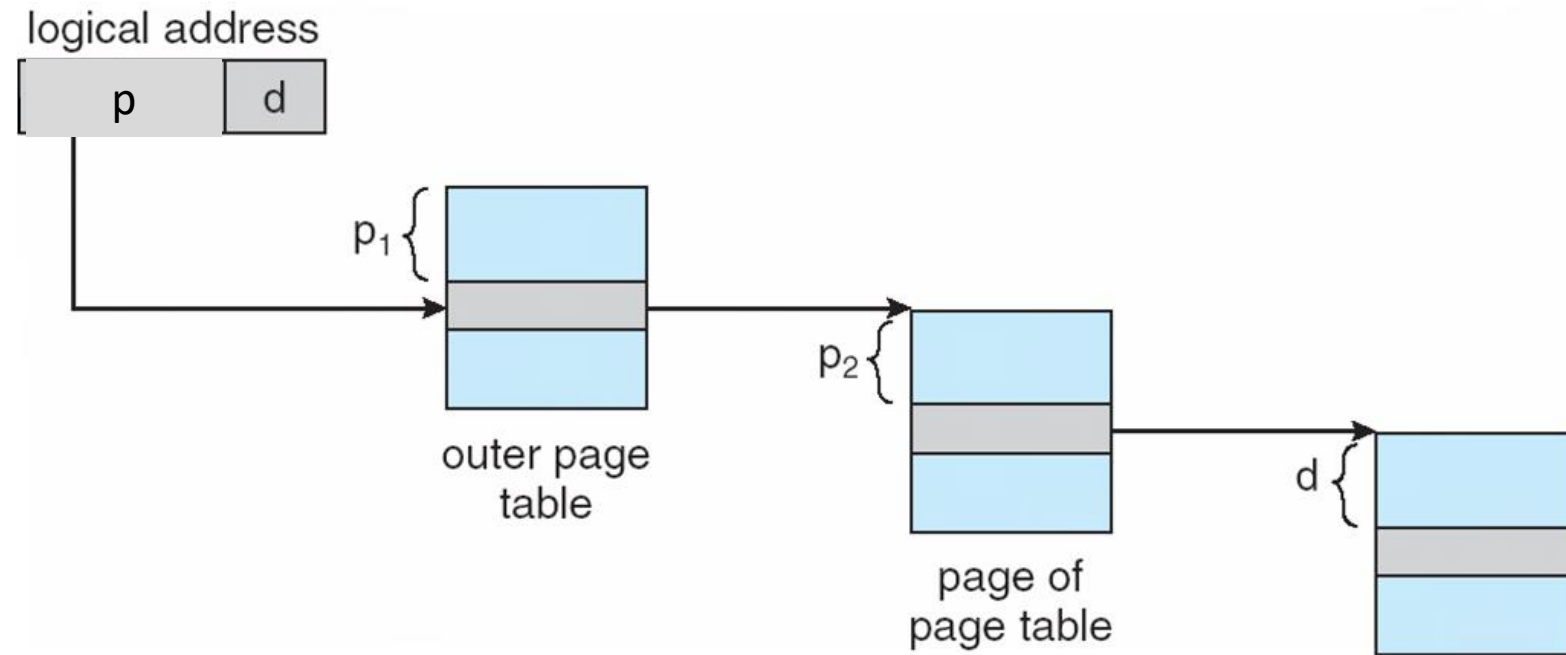
Two-level Page Table Scheme



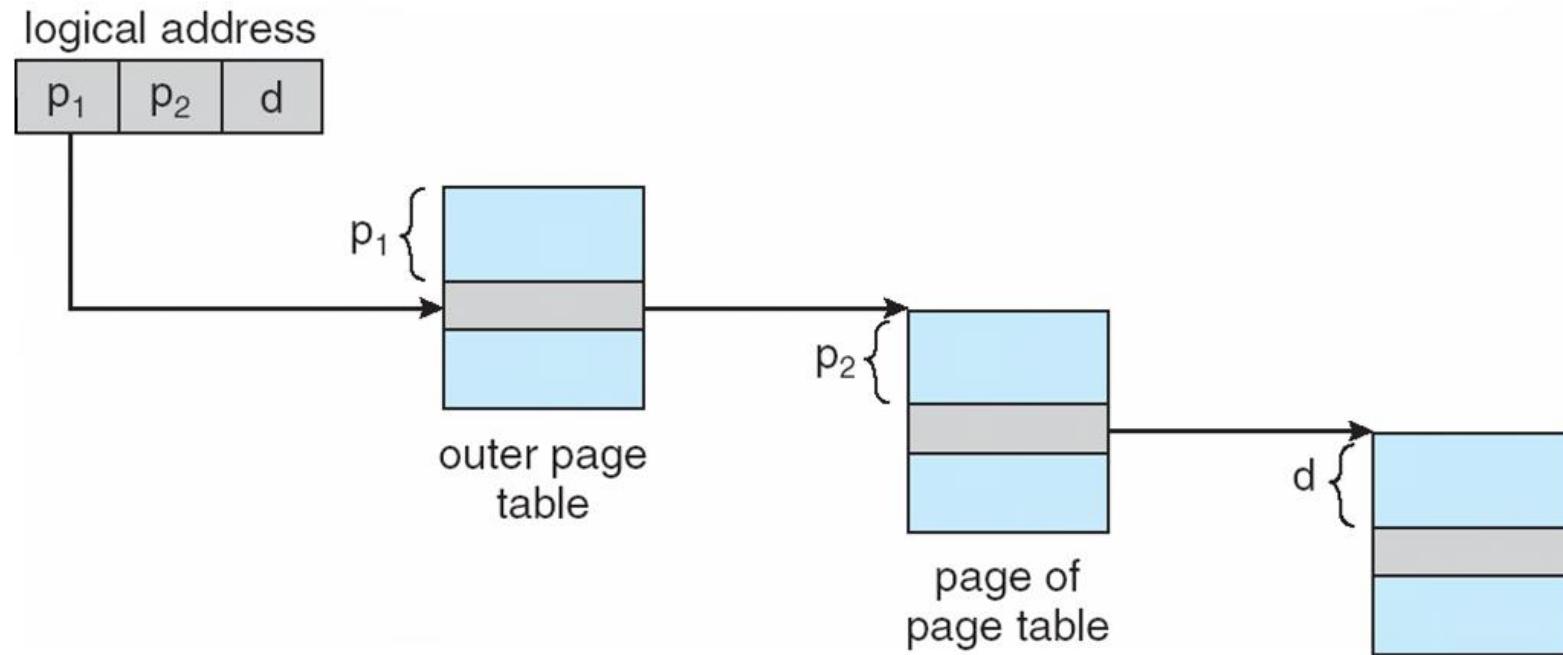
Multi-level Page Table

- Turns a linear page table to a tree
- Divide the page table into multiple pages (page-sized page tables)
- Create a page directory (outer level PT) to track the pages of the page table (inner-level PTs)

Multi-level Address Translation Scheme




Multi-level Address Translation Scheme



How many bits for each paging level?

32-bit address:




- Each inner page table fits within a page
- number of PTE = page size / PTE size
 - Assume PTE size = 4 bytes
 - Page size = 2^{12} bytes = 4KB
 - number PTE = 2^{12} bytes / 2^2 bytes
 - \rightarrow number PTE = 2^{10}
 - \rightarrow # bits for selecting inner page = 10

Remaining bits for outer page:

- $32 - 10 - 12 = 10$ bits

How many bits for each paging level?

32-bit address:




- Outer page table has one entry per inner page table
- Number of inner page tables = linear page table size / page size
 - Assume PTE size = 4 bytes
 - Linear page table size = # PTE * PTE size
 - Linear page table size = $2^{20} * 4 \text{ bytes} = 2^{22} \text{ bytes}$
- Number of inner page tables = $2^{22} \text{ bytes} / 2^{12} \text{ bytes}$
 - \rightarrow # bits for selecting outer page = 10

Remaining bits for inner page:

- $32 - 10 - 12 = 10 \text{ bits}$

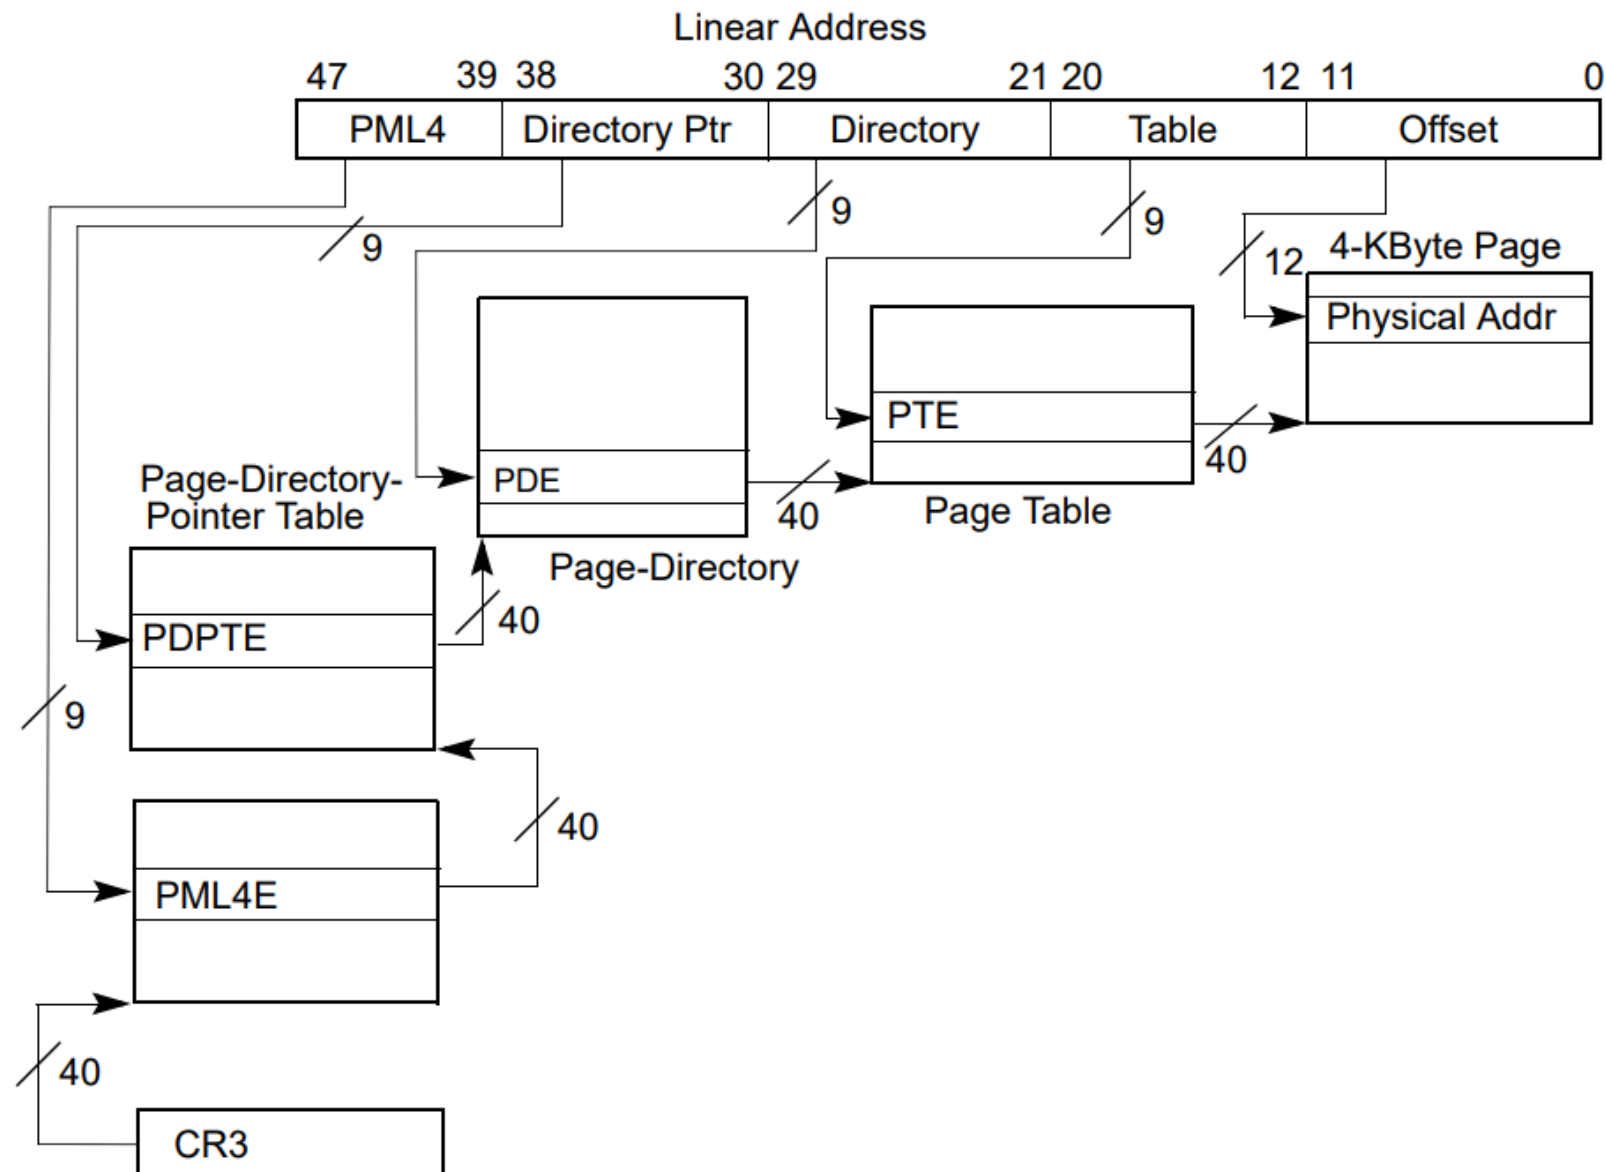
64-bit Address Space

- Page size is 4 KB (2^{12})
 - page table has 2^{52} entries
 - In two level scheme, inner page tables could be 2^{10} 4-byte entries

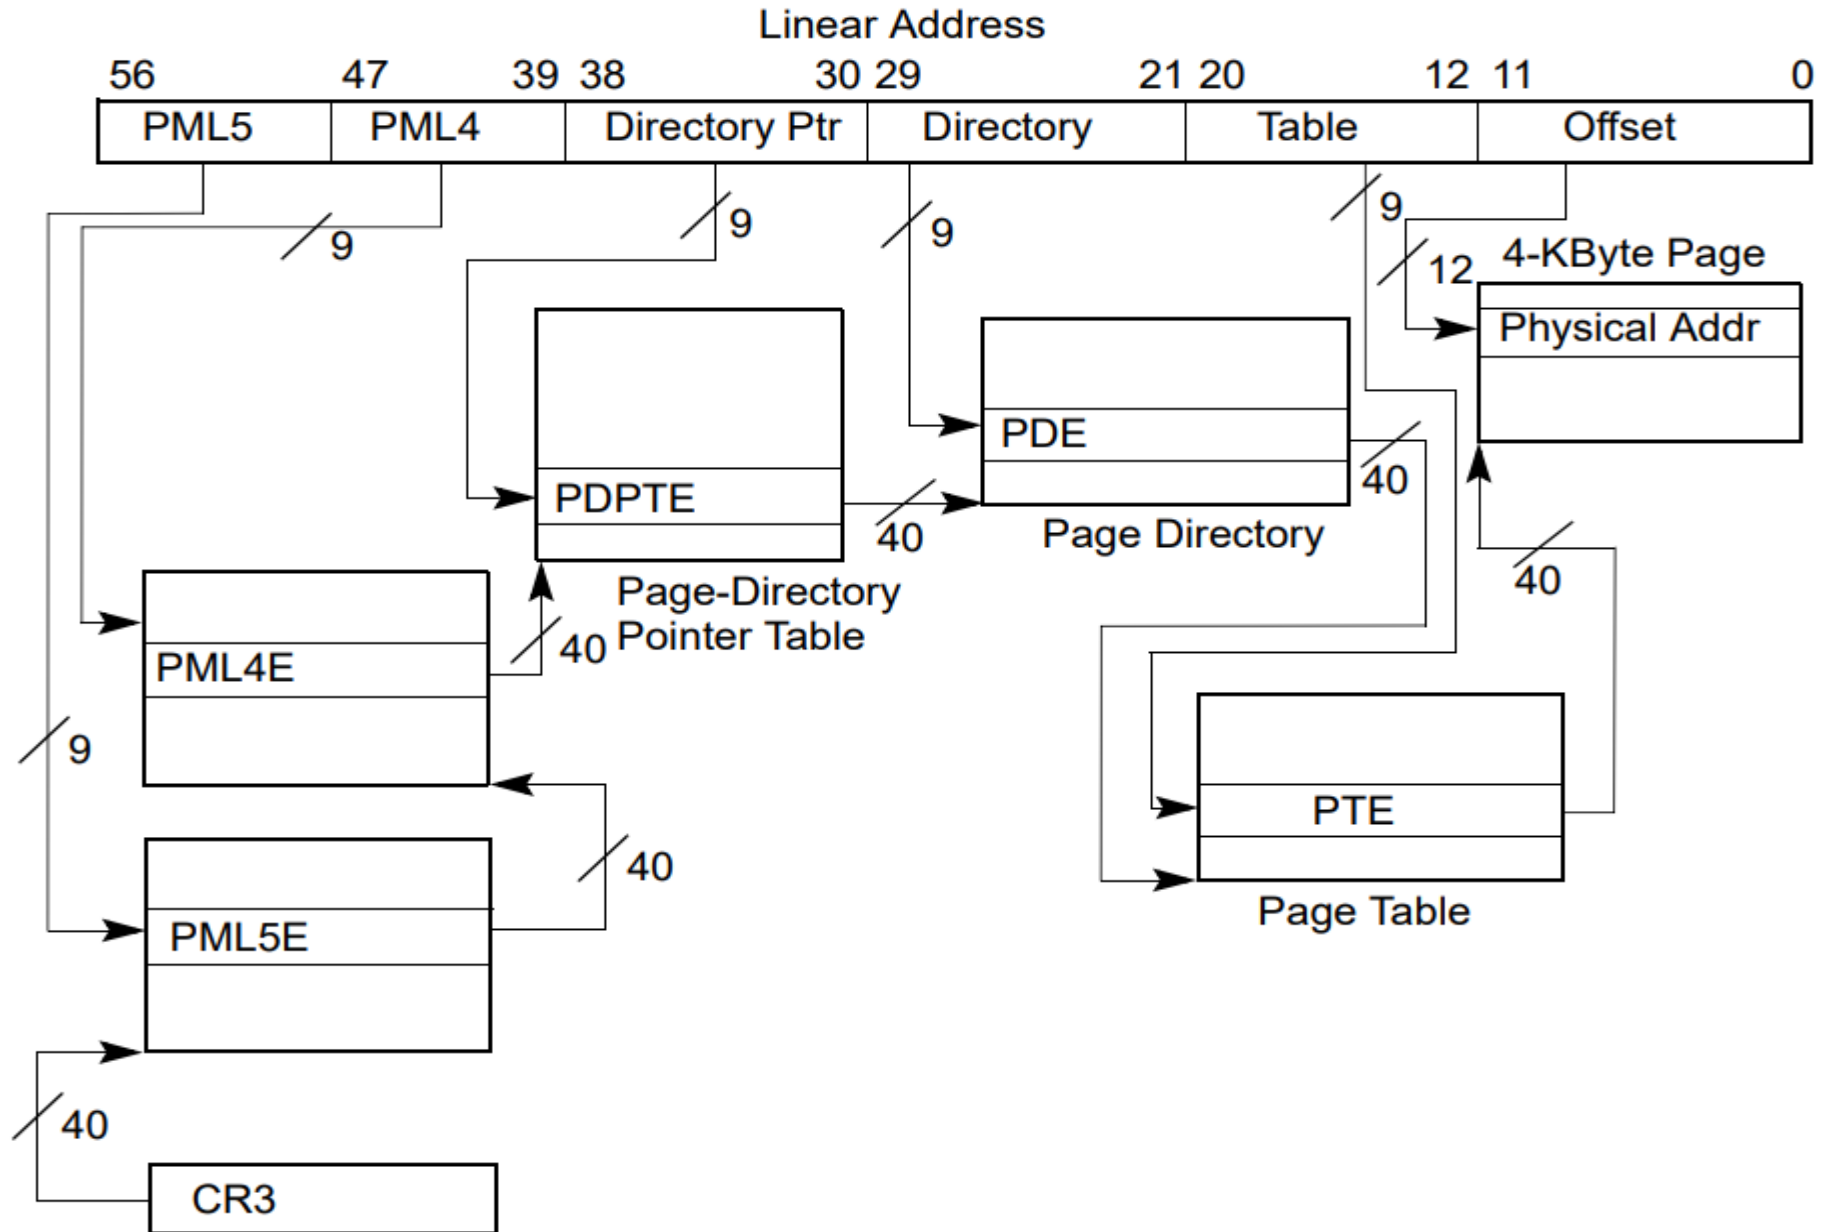


- Outer page table has 2^{42} entries or 2^{44} bytes
- One solution is to add a 2nd outer page table
- The 2nd outer page table is still 2^{34} bytes in size
 - And possibly 4 memory access to get to one physical memory location

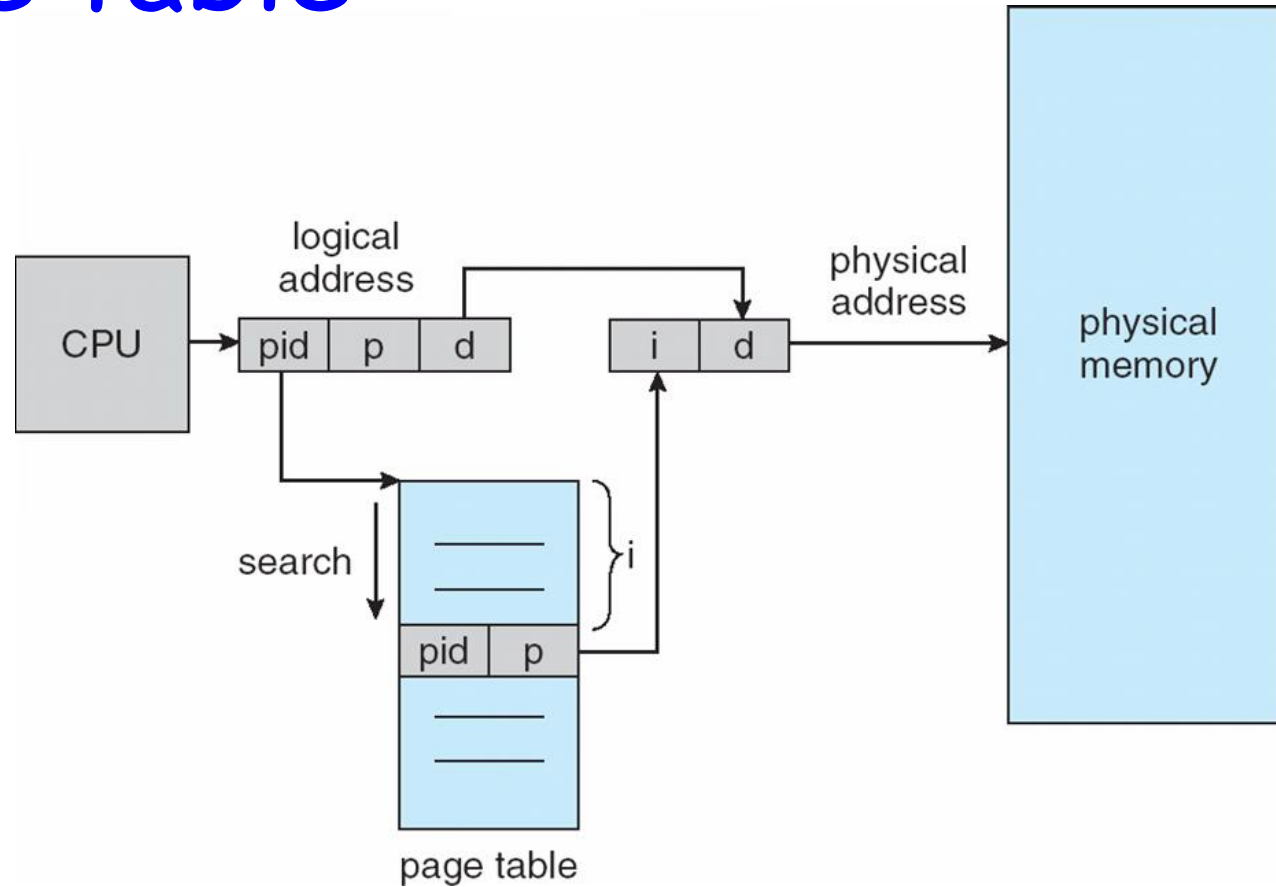
Intel x86-64 (4 level)



Intel x86-64 (5-level)

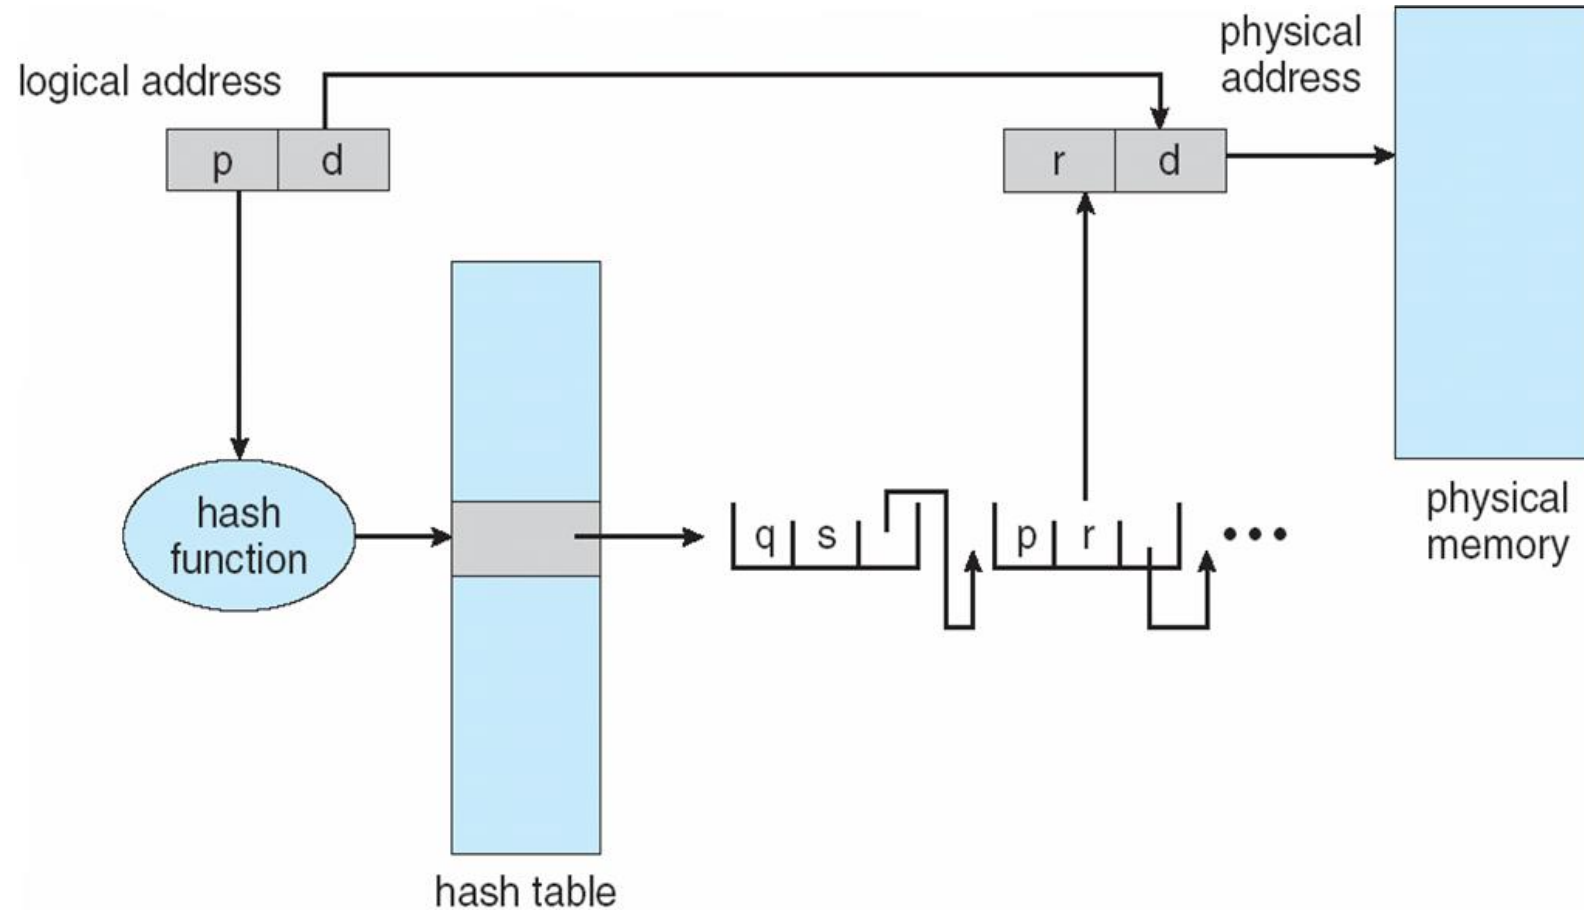


Inverted Page table



One entry for each physical page

Hashed Page Table



Avoid Simple linear page table

- Use complex page table instead of big array
- Inverted page table
- Hashing
- TLB should be software managed