

# VM: Beyond Physical

Sridhar Alagar

# Not enough Physical Memory

- Single process' address space can be larger than the physical memory
  - Need this for flexibility in programming
- Combined address spaces of all the processes larger than physical memory
- OS needs to provide illusion of large physical memory

# Not enough Physical Memory

A process request memory for a page and the memory is full.  
How to handle this request?

# Demand Paging

- Load a page into memory on demand
- If page needed is not in memory
  - Need to detect and load the page into memory from storage
    - without changing program behavior
    - without programmer needing to change code
- Need hardware support


# Demand Paging

- Mechanism
  - How to identify whether a page is in memory or on disk?
  - How to fetch page from disk transparently?
- Policy
  - Which page(s) should be replaced?

# Mechanism


- Each page in virtual address space maps to one of three locations
  - Physical main memory: Small, fast, expensive
  - Disk (backing store): Large, slow, cheap
  - Nothing (error): Free
- Extend page tables with an extra bit: **present**
  - permissions (r/w), valid, **present**
  - Page in memory: present bit **set** in PTE
  - Page on disk: present bit cleared
    - PTE points to block on disk
    - Causes trap into OS when page is referenced (**page fault**)

# Present Bit



What if vpn 0xb referenced?

# Handling Page Fault






# Aspects of demand paging

- Extreme case - start process with *no* pages in memory
  - Page fault on first access of every process pages
  - **Pure demand paging**
- An instruction could access multiple pages -> multiple page faults
- Hardware support needed for demand paging
  - Page table with protection bit
  - Secondary memory (swap device with **swap space**)
  - Instruction restart

# Frame Replacement

- No free frame to bring in a page
- Find a **victim** to evict
- Policy decision
- Plenty of time to make



# Replacement algorithms

- Goal
  - Reduce the number of page faults
- Several replacement algorithms

# The Optimal replacement Policy

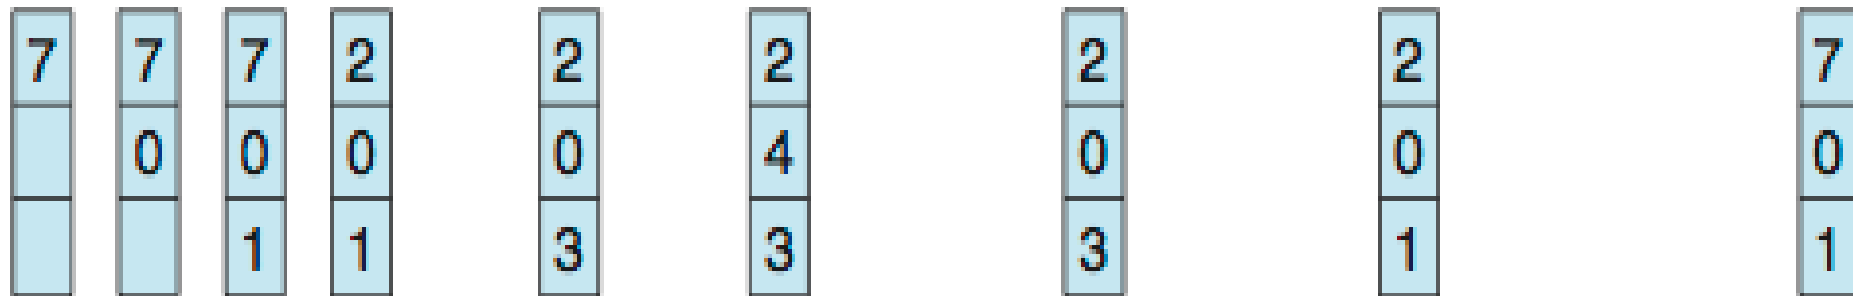
- Leads to the fewest number of faults overall
- Replaces the page that will be accessed furthest in the future
- Serve only as a comparison point, to know how close we are to perfect

# Optimal Policy

Assume 3 frames

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

9 faults

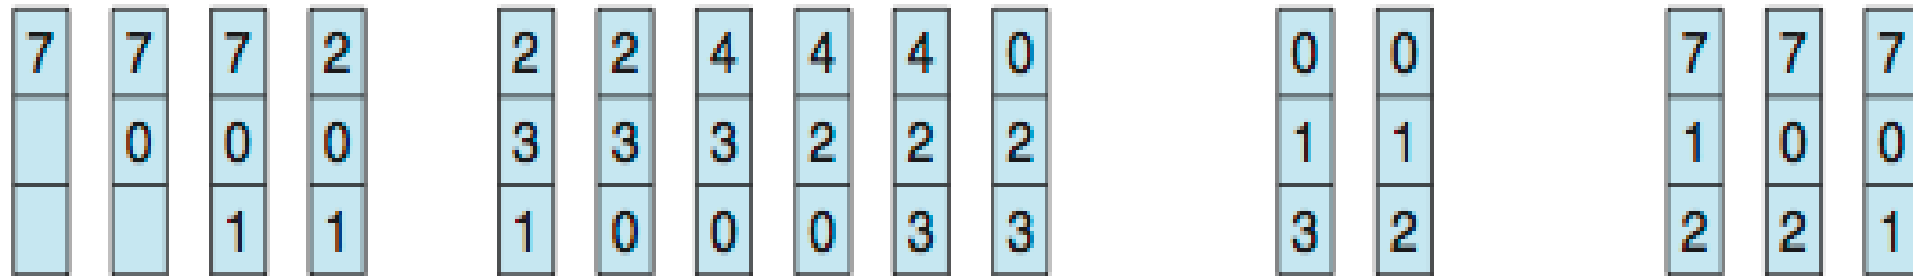
# FIFO

- Pages are placed in a queue when they enter the system.
- When a replacement occurs, the page on the tail of the queue(the "First-in" pages) is evicted.
- It is simple to implement.

# FIFO

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

15 faults

# Belady's anomaly in FIFO

- Increasing the number frames to 4 increases the page fault

Reference Row											
1	2	3	4	1	2	5	1	2	3	4	






# Least Recently Used (LRU)

- Since future is not known, use past to predict
- Replace the least recently used frame

# Least Recently Used (LRU)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

12 faults

# Implementing LRU

- Software Perfect LRU
- OS maintains ordered list of physical pages by reference time
  - When page is referenced: Move page to front of list
  - Victim: last page of the list
  - **Trade-off:** Slow on memory reference, fast on replacement

# Implementing LRU

- Hardware Perfect LRU
- Associate timestamp register with each frame
- When frame is referenced: Store system clock in register
- Victim: Scan through registers to find the oldest timestamp
- **Trade-off:** Fast on memory reference, slow on replacement (especially as size of memory grows)


# Approximate LRU

- In practice, do not implement Perfect LRU
- LRU is an approximation anyway, so approximate more
- Goal: Find an old page, but not necessarily the very oldest
- Hardware provides a single bit - **reference/use** bit
- The bit is set whenever the page is referenced


# Clock Algorithm

- OS executes page replacement algorithm:
  - Look for page with reference bit cleared (has not been referenced for awhile)
- Implementation:
- Keep a pointer to last replaced frame
- Traverse frames in circular fashion
- Clear reference bit if it is set
- Stop when found a page with already cleared bit, replace this frame

# Clock: Look for a frame




# Clock: Look For a Page






# Clock: Look For a Page




evict **frame 2** because it has not been recently used

# Clock: Look For a Page




frame 0 is accessed

# Clock: Look For a Page




frame 0 is accessed

# Clock: Look for a frame




New request for a frame


# Clock: Look for a frame



# Clock: Look for a frame



# Clock: Look for a frame



evict **frame 1** because it has not been recently used

# When to run replacement algorithm?

- Expensive to run for every page fault
- Maintain two thresholds for free frames:
  - High water mark
  - Low water mark
- When # free frames < low water mark run replacement algorithm
- Free frames till high water mark crossed



# Other extensions

- Avoid replacing **dirty** pages if possible
  - dirty bit part of PTE.
  - set whenever a page is modified
- Pre-fetch pages

# Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.