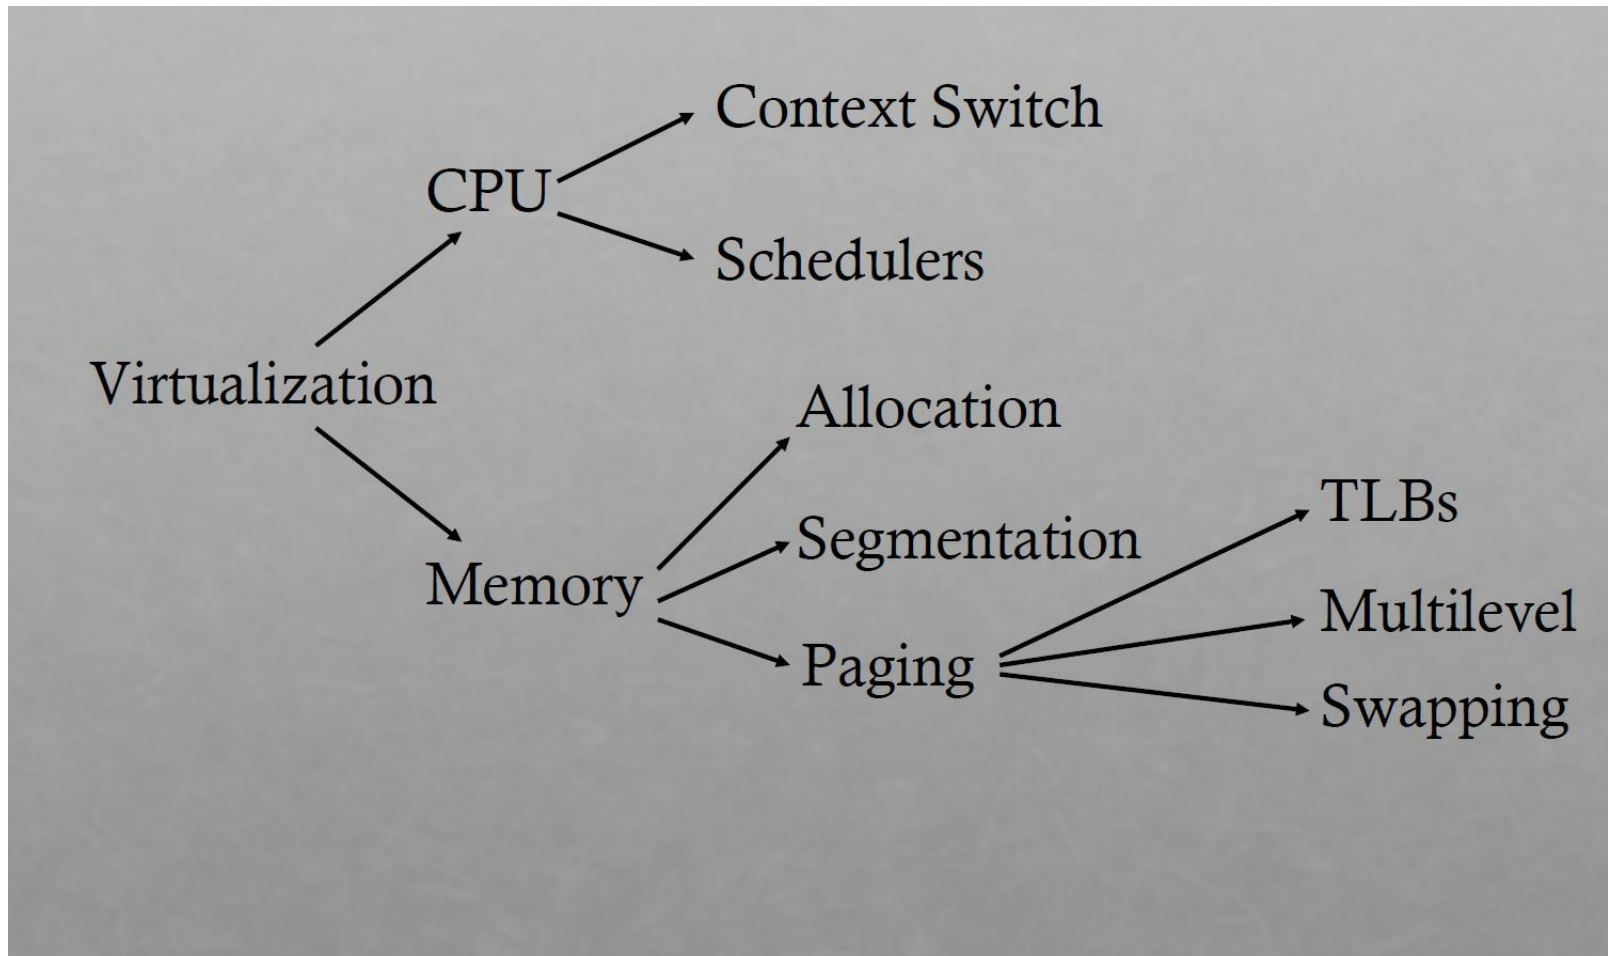


Big Picture so far...



Concurrency: Threads

Sridhar Alagar


Motivation

- Develop applications that utilizes many cores of CPU
- Build application using many processes
 - Example: Browser with one process per tab
 - Communicate using pipe or other IPC mechanisms
- Pros
 - No need to develop any new mechanisms
- Cons
 - Cumbersome programming
 - High communication overheads
 - Expensive context switch


New Abstraction: Threads

- Threads are like processes except that
 - threads of the same process share the same address space
- Threads are like procedures within a process except that they can be run in parallel
- Divide large tasks among multiple cooperating threads
 - Communicate through shared variables in shared address space


What state do threads share?



Share Address Space




Instruction Pointer?




Share code; but each executing independently
different parts of the code

Stack Pointer?



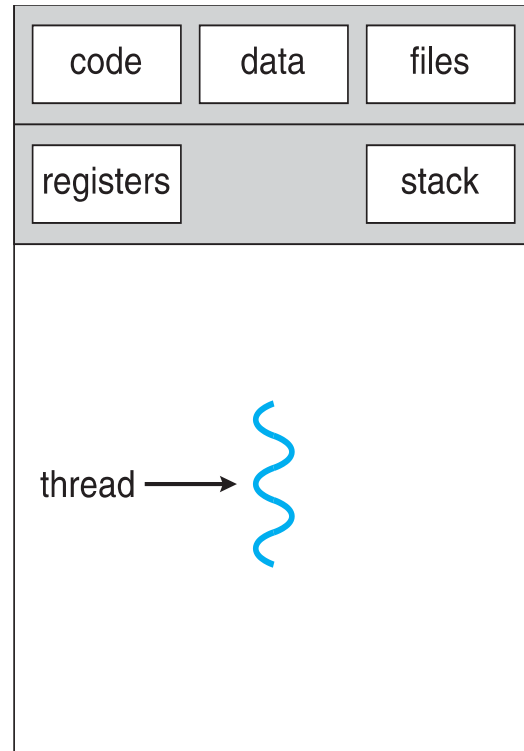
They execute different functions (instances)

Stack Pointer?

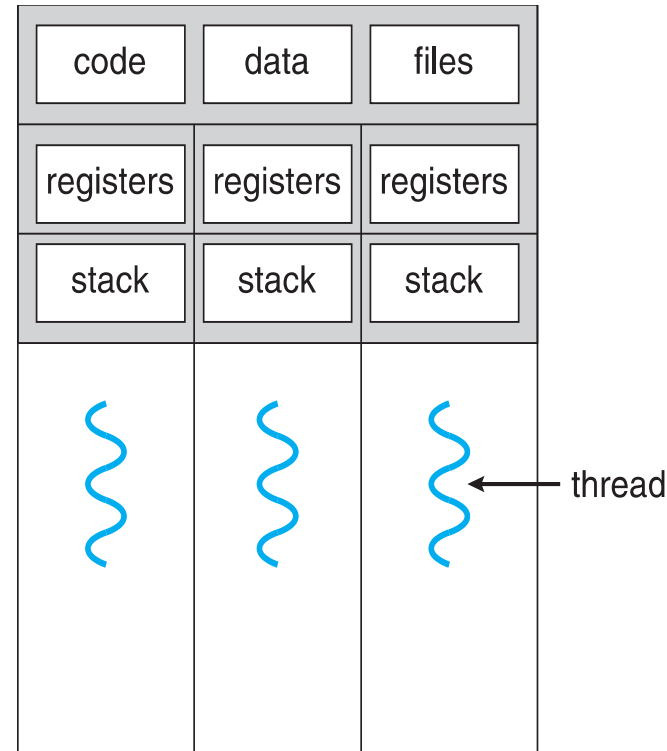


They execute different functions (instances)

Single and Multithreaded Processes



single-threaded process



multithreaded process

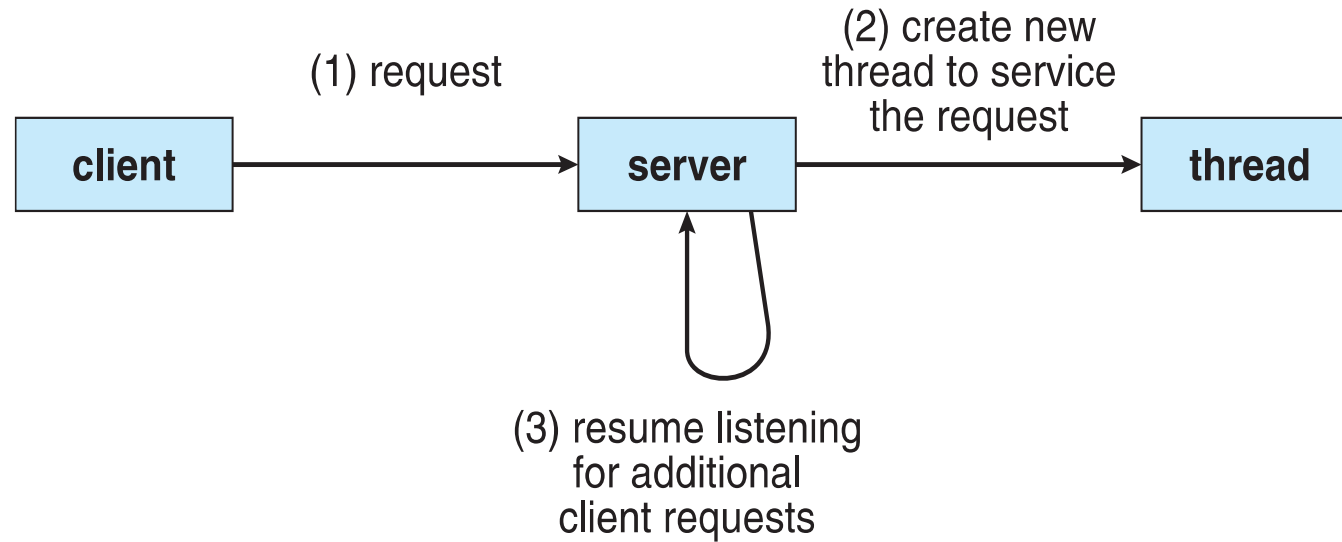
Threads vs Process

- Multiple threads within a same process share
 - Process ID
 - Address space
 - Code
 - Data (heap too)
 - Open file descriptors
 - Current Working directory
 - User and Group ID
- Each thread has its own
 - Thread ID
 - Registers including PC, and SP
 - Stack (in the same address space)

Programming Models

- Multiple tasks within the application can be implemented by separate threads.
- A word processor application uses several threads to
 - Get data from keyboard (foreground)
 - Spell checking and grammar (back ground)
 - Load images from file
 - Take a back up
- Another example is client server architecture

Multithreaded Server Architecture




Benefits

- Responsiveness
 - one thread may continue to execute while part of process is blocked, especially important for user interfaces
- Resource Sharing
 - threads share resources(code, memory, files) of process; easier than shared memory or message passing
- Economy
 - cheaper than process creation, thread switching lower overhead than context switching
-


OS Support: Many-to-One

- Many user-level threads mapped to single kernel thread
- Implemented by the user level runtime libraries
 - Create, schedule, synchronize at user level
- Kernel not aware of user level threads
 - Thinks each process contain single thread



OS support: One-to-One

- Each user-level thread mapped to a kernel thread by the OS
- Each kernel thread scheduled independently
- Thread operation performed by the OS




OS Support: Many-to-One

- Pros

- Does not require OS support; Portable
- Lower overhead thread operation since no system call


- Cons

- One thread blocking causes all to block
- Multiple threads cannot run in parallel on muticore system



OS support: One-to-One

- No blocking of other threads
- Can in run in parallel on a multiprocessor
- Higher overhead
- OS should be scalable or limit # threads
- Examples
 - Windows
 - Linux



Thread API

- Thread library provides the API for creating and managing threads
- Several exists
 - POSIX Pthreads
- Common thread operations:
 - Create
 - Exit
 - Join (instead of wait for process)

Thread Creation

```
#include <pthread.h>
```

```
int
```

```
pthread_create(      pthread_t*      thread,  
                  const pthread_attr_t* attr,  
                  void* (*start_routine)(void*),  
                  void*      arg);
```

Exit a thread

```
int  
pthread_exit(void *value_ptr);
```

- `value_ptr`: **A pointer to the return value**

Wait for a thread to complete

`int`

```
pthread_join(pthread_t thread, void **value_ptr);
```

- `thread`: Specify which thread to wait for
- `value_ptr`: A pointer to the return value

Thread Schedule #1

counter = counter + 1; counter at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4A

Thread Schedule #1

State:

0x9cd4: 100
%eax: 100
%rip = 0x19a

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

- T1 →
- 0x195 mov 0x9cd4, %eax
 - 0x19a add \$0x1, %eax
 - 0x19d mov %eax, 0x9cd4

Thread Schedule #1

State:

0x9cd4: 100
%eax: 101
%rip = 0x19d


process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

State:

0x9cd4: 101
%eax: 101
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

Thread Schedule #1

State:

0x9cd4: 101
%eax: 101
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

Thread Context Switch

Thread Schedule #1

State:

0x9cd4: 101
%eax: ?
%rip = 0x195

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

State:

0x9cd4: 101
%eax: 101
%rip = 0x19a

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

State:

0x9cd4: 101
%eax: 102
%rip = 0x19d

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

State:

0x9cd4: 102

%eax: 102

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2 →

Thread Schedule #1

State:

0x9cd4: 102
%eax: 102
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2 →

Desired Result!

Another schedule

Thread Schedule #2

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 100
%eax: 101
%rip = 0x19d

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

T1



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Context Switch

Thread Schedule #2

State:

0x9cd4: 100

%eax: ?

%rip = 0x195


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 100
%eax: 100
%rip = 0x19a

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 100
%eax: 101
%rip = 0x19d

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4A

T2 →

Thread Schedule #2

State:

0x9cd4: 101
%eax: 101
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Context Switch

Thread Schedule #2

State:

0x9cd4: 101
%eax: 101
%rip = 0x19d


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: 101
%rip: 0x1a2

T1 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

State:

0x9cd4: 101
%eax: 101
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: 101
%rip: 0x1a2

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

Thread Schedule #2

State:

0x9cd4: 101
%eax: 101
%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: 101
%rip: 0x1a2

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

WRONG Result! Final value of counter is 101

Timeline View

Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

Thread 2

```
mov 0x123, %eax  
  
add %0x2, %eax  
  
mov %eax, 0x123
```

How much is added to shared variable? **3: correct!**

Timeline View

Thread 1

`mov 0x123, %eax`

`add %0x1, %eax`

`mov %eax, 0x123`

Thread 2

`mov 0x123, %eax`

`add %0x2, %eax`

`mov %eax, 0x123`

How much is added?

2: incorrect!

Timeline View

Thread 1

`mov 0x123, %eax`

`add %0x1, %eax`

`mov %eax, 0x123`

Thread 2

`mov 0x123, %eax`

`add %0x2, %eax`

`mov %eax, 0x123`

How much is added?

1: incorrect!

Timeline View

Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

Thread 2

```
mov 0x123, %eax  
add %0x2, %eax  
mov %eax, 0x123
```

How much is added?

3: correct!

Timeline View

Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

Thread 2

```
mov 0x123, %eax  
add %0x2, %eax
```

```
mov %eax, 0x123
```

How much is added?

2: incorrect!

Non-Determinism

- Concurrency leads to non-deterministic results
 - Non deterministic result: different results even with same inputs
 - race conditions
- Whether bug manifests depends on CPU schedule!
- Passing tests means little
- How to program: imagine scheduler is malicious
 - Assume scheduler will pick bad ordering at some point...

What do we want?

- Want 3 instructions to execute as an uninterruptable group
- That is, we want them to be atomic

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

—critical section

- Need mutual exclusion for critical sections
- if process A is in critical section, process B can't be in CS
(okay if other processes do unrelated work)

Solution using Locks

- Allocate and Initialize
 - `Pthread_mutex_t mylock;`
 - `Pthread_mutex_init(&mylock, NULL);`
- Acquire: `pthread_mutex_lock(&mylock)`
 - Acquire exclusion access to lock;
 - Wait if lock is not available (some other process in critical section)
 - Spin or block (relinquish CPU) while waiting
- Release: `pthread_mutex_unlock(&mylock)`
 - Release exclusive access to lock; let another process enter critical section

Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.