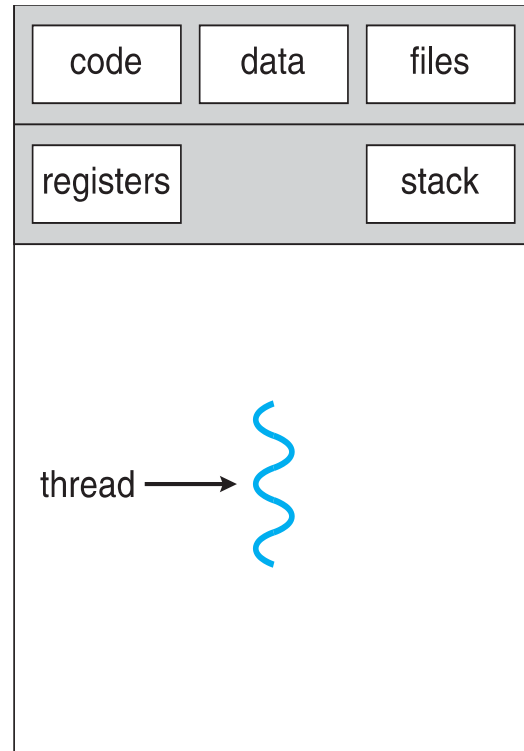


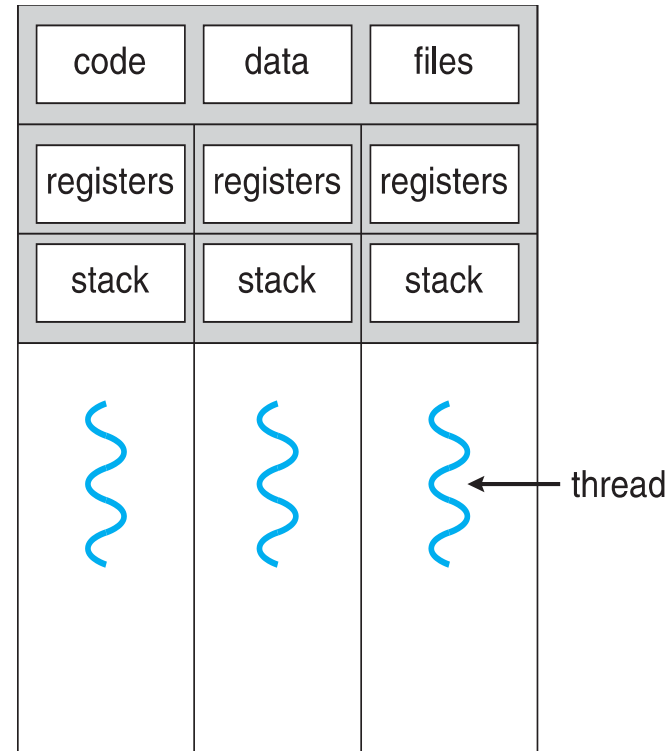
# Concurrency: Locks

Sridhar Alagar

# What do threads share?



single-threaded process



multithreaded process

# Concurrency Issues

- Concurrent threads access shared variable/memory
- While one thread is accessing the shared memory another thread accessing the shared memory
  - Due to context switching, or
  - Multi-processor environments
- Due to this race condition arises and results in non-deterministic and incorrect execution

# Handling Critical Section

counter = counter + 1; counter is a shared variable

- Instructions updating shared memory must execute as an uninterruptable group

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

—critical section

- **Need mutual exclusion** for critical sections
- if process A is in critical section, process B can't be in CS

# Criteria for a Solution to CS Problem

- Mutual Exclusion
  - only one thread in CS at a given time
- Fairness
  - Does each thread waiting to enter CS get a fair shot at entering CS? (Starvation free)
- Performance
  - The time overhead added by the solution

# Locks

- Use locks to ensure mutually exclusion for CS

```
1    lock_t mutex; // some globally-allocated lock 'mutex'
2    ...
3    lock(&mutex);
4    balance = balance + 1;
5    unlock(&mutex);
```

# Locks Semantics

- A variable holds the state of the lock
- available/free: no thread has acquired the lock
- acquired/locked: exactly one thread has acquired the lock
- lock()
  - Acquire the lock if it is free

# Pthread Locks

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Pthread_mutex_lock(&lock); // acquire the lock
4  balance = balance + 1;
5  Pthread_mutex_unlock(&lock);
```



# How to implement Locks?

# Control Interrupts

- Disable interrupts before entering CS and enable after exiting CS
  - one of the earliest solution

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

# Problems with disabling interrupts

- Application can monopolize the processor
- Does not work on multi-processors
- Interrupts can be lost
- Masking/unmasking very slow

# A software solution - use load/store

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7  void lock(lock_t *mutex) {
8      while (mutex->flag == 1)    // TEST the flag
9          ;    // spin-wait (do nothing)
10     mutex->flag = 1;    // now SET it !
11 }
12 void unlock(lock_t *mutex) {
13     mutex->flag = 0;
14 }
```

# Problems with software solution

- No mutual exclusion

Thread1	Thread2
<pre>call lock() while (flag == 1) interrupt: switch to Thread 2</pre>	<pre>call lock() while (flag == 1) flag = 1; interrupt: switch to Thread 1</pre>
<pre>flag = 1; // set flag to 1 (too!)</pre>	

- Spin-waiting wastes CPU cycles

# Hardware is your friend

- Need a hardware supported atomic instruction for test and set

```
1  int TestAndSet(int *ptr, int new) {  
2      int old = *ptr; // fetch old value at ptr  
3      *ptr = new;      // store 'new' into ptr  
4      return old;      // return the old value  
5  }
```

# Spin Lock using test-and-set

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ;    // spin-wait
13 }
14 void unlock(lock_t *lock) {
15     lock->flag = 0;
16 }
```

# Compare and Swap (H/W-atomic)

```
1  int CompareAndSwap(int *ptr, int expected, int new) {
2      int actual = *ptr;
3      if (actual == expected)
4          *ptr = new;
5      return actual;
6  }
```


```
1  void lock(lock_t *lock) {
2      while (CompareAndSwap(&lock->flag, 0, 1) == 1)
3          ; // spin
4  }
```



# Evaluating Spin Locks

- Correctness:
  - Yes: Only one thread is allowed to enter CS
- Fairness:
  - No: a thread may spin forever
- Performance
  - Can be bad, especially, in a single processor system

# Basic Spinlocks are Unfair



Scheduler is independent of locks/unlocks

# Fetch-and-add

- Atomically increment value and return old value

```
1  int FetchAndAdd(int *ptr) {  
2      int old = *ptr;  
3      *ptr = old + 1;  
4      return old;  
5  }
```

# Ticket lock using fetch-and-add

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5  void lock_init(lock_t *lock) {
6      lock->ticket = 0;
7      lock->turn = 0;
8  }
9  void lock(lock_t *lock) {
10     ...
11 }
12 void unlock(lock_t *lock) {
13     ...
14 }
```

# Ticket lock using fetch-and-add

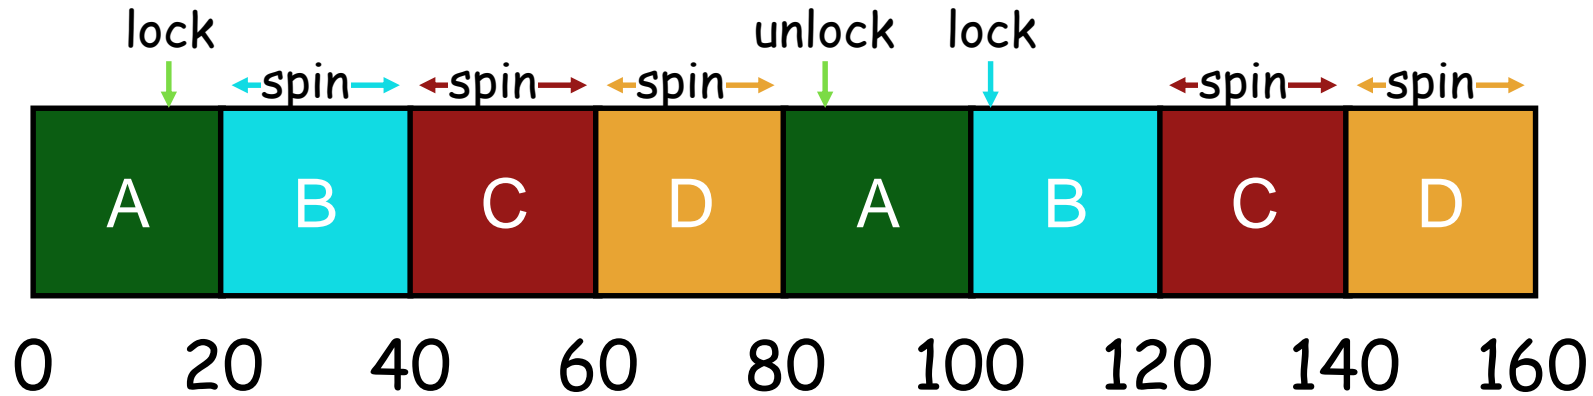
Ensures  
fairness

```
1 void lock(lock_t *lock) {  
2     int myturn = FetchAndAdd(&lock->ticket);  
3     while (lock->turn != myturn); // spin  
4 }  
  
5 void unlock(lock_t *lock) {  
6     FetchAndAdd(&lock->turn);  
7 }
```

# Too much spinning

- Hardware based spin locks are simple
- But *CPU* cycles wasted by spinning

# CPU Scheduler is Ignorant




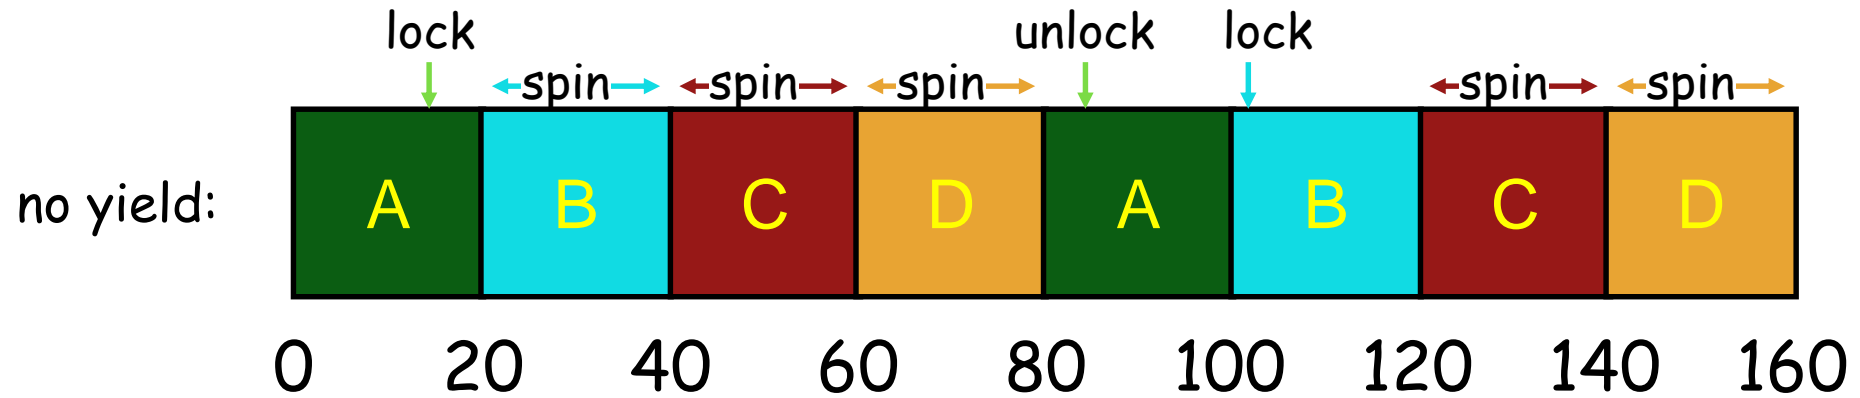
CPU scheduler may run **B** instead of **A**  
even though **B** is waiting for **A**

# Ticket lock with yield

```
1 void lock(lock_t *lock) {  
2     int myturn = FetchAndAdd(&lock->ticket);  
3     while (lock->turn != myturn)  
4         yield();  
5 }  
6 void unlock(lock_t *lock) {  
7     FetchAndAdd(&lock->turn);  
8 }
```



# Yield Instead of Spin



# Spinlock Performance

- Waste...
  - without yield:  $O(\text{threads} * \text{time\_slice})$
  - with yield:  $O(\text{threads} * \text{context\_switch})$
  - So even with yield, performance is slow with high thread contention
- Next improvement: Block and put thread on waiting queue instead of spinning


# Lock Implementation: Block when Waiting

- Lock implementation removes waiting threads from scheduler ready queue (e.g., `park()` and `unpark()`)
- **Scheduler** runs any thread that is **ready**
- *Good separation of concerns*

RUNNABLE: A, B, C, D

RUNNING: <empty>


WAITING: <empty>



RUNNABLE: B, C, D

RUNNING: A


WAITING: <empty>



RUNNABLE: C, D, A

RUNNING: B


WAITING: <empty>



RUNNABLE: C, D, A

RUNNING:


WAITING: B



RUNNABLE: D, A

RUNNING: C

WAITING: B






RUNNABLE: A, C

RUNNING: D


WAITING: B



RUNNABLE: A, C

RUNNING:


WAITING: B, D



RUNNABLE: C

RUNNING: A


WAITING: B, D



RUNNABLE: A

RUNNING: C


WAITING: B, D



RUNNABLE: C

RUNNING: A


WAITING: B, D



RUNNABLE: B, C

RUNNING: A


WAITING: D



RUNNABLE: B, C

RUNNING: A


WAITING: D



RUNNABLE: C, A

RUNNING: B

WAITING: D





# Lock Implementation: Block when Waiting

```
typedef struct {  
    bool lock = false;  
    queue_t q;  
} LockT;
```

Problem: q is shared

Solution: lock q before modifying it

```
void lock(LockT *l) {  
    if(testandset(&l->lock, 1)){  
        qadd(l->q, tid);  
        park();    // blocked  
    }  
}  
  
void unlock(LockT *l) {  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q)) ;  
}
```

# Lock Implementation: Block when Waiting

```
typedef struct {  
    bool guard = false;  
    bool lock = false;  
    queue_t q;  
} LockT;
```

- (a) use guard to lock q
- (b) Why okay to **spin** on guard?
- (c) In unlock(), why not set lock=false when unpark?
- (d) What is the race condition?

```
void lock(LockT *l) {  
    while (testandset(&l->guard, true));  
    if (l->lock) {  
        qadd(l->q, tid);  
        l->guard = false;  
        park();      // blocked  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}
```

```
void unlock(LockT *l) {  
    while (testandset(&l->guard, true));  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q)) ;  
    l->guard = false;  
}
```

# Race Condition

**Thread 1** (in lock)

```
if (l->lock) {  
    qadd(l->q, tid);  
    l->guard = false;
```

```
park();    // block
```

**Thread 2** (in unlock)

```
while (testandset(&l->guard, true));  
if (qempty(l->q)) // false!!  
else unpark(qremove(l->q));  
l->guard = false;
```

Problem: Guard not held when call park()  
Unlocking thread may unpark() before other park()

# Block when waiting: correct LOCK

```
typedef struct {  
    bool lock = false;  
    bool guard = false;  
    queue_t q;  
} LockT;
```

setpark() fixes race condition

```
void lock(LockT *l) {  
    while (TAS(&l->guard, true));  
    if (l->lock) {  
        qadd(l->q, tid);  
        setpark(); // notify of plan  
        l->guard = false;  
        park(); // unless unpark()  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}
```

```
void unlock(LockT *l) {  
    while (TAS(&l->guard, true));  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q));  
    l->guard = false;  
}
```

# Spin Waiting vs Blocking

- Each approach is better under different circumstances
- Uniprocessor
  - Waiting process is scheduled --> Process holding lock isn't
  - Waiting process should always relinquish processor
  - Associate queue of waiters with each lock (as in previous implementation)
- Multiprocessor
  - Waiting process is scheduled --> Process holding lock might be
  - Spin or block depends on how long,  $t$ , before lock is released
    - Lock released quickly --> Spin-wait
    - Lock released slowly --> Block
    - Quick and slow are relative to context-switch cost,  $C$

# Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.