

Concurrency: Condition Variables

Sridhar Alagar

Synchronization Objectives

- Mutual exclusion (e.g., A and B don't run at same time)
 - solved with locks
- Ordering (e.g., B runs after A does something)
 - solved with condition variables and semaphores

Ordering Example: Join

```
pthread_t p1, p2;  
Pthread_create(&p1, NULL, mythread, "A");  
Pthread_create(&p2, NULL, mythread, "B");  
// join waits for the threads to finish  
Pthread_join(p1, NULL);  
Pthread_join(p2, NULL);  
printf("main: done\n");  
return 0;
```

how to implement join()?

Condition Variables

- Condition Variable: queue of waiting threads (CV)
- B waits for condition CV before running
 - `wait(CV, ...)`
- A sends signal to CV when the condition is met
 - `signal(CV, ...)`

Join Implementation

Parent:

```
void thread_join() {  
    Cond_wait(&c); // y  
}
```

Child:

```
void thread_exit() {  
    Cond_signal(&c); // b  
}
```

Example schedule:

Parent: y

Child: a

Works!

Join Implementation

Parent:

```
void thread_join() {
```

```
    Cond_wait(&c); // y
```

```
}
```

Child:

```
void thread_exit() {
```

```
    Cond_signal(&c); // b
```

```
}
```

Can you construct ordering that does not work?

Example broken schedule:

Parent:

y

Child:

b

parents waits forever

Rule of Thumb 1

- Keep state in addition to CV's!
- CV's are used to signal threads when state changes
- If state is already as needed, thread doesn't wait for a signal!

Join Implementation (Attempt 2)

Parent:

```
void thread_join() {
```

```
    if (done == 0)           // x
        Cond_wait(&c); // y
```

```
}
```

Fixes previous broken ordering

Example schedule:

Parent:

x

y

Child:

b

c

Child:

```
void thread_exit() {
```

```
    done = 1           // b
    Cond_signal(&c);    // c
```

```
}
```


Join Implementation (Attempt 2)

Parent:

```
void thread_join() {
```

```
    if (done == 0)           // x
        Cond_wait(&c); // y
```

```
}
```

Child:

```
void thread_exit() {
```

```
    done = 1           // b
    Cond_signal(&c);    // c
```

```
}
```

Parent:

x

y

Child:

b

c

Use mutex to ensure no race between "interacting with state" and wait/signal

Join Implementation (Attempt 3)

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c);       // y  
    Mutex_unlock(&m);        // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                // b  
    Cond_signal(&c);         // c  
    Mutex_unlock(&m);        // d  
}
```

Parent: w x y

Child: a

Both parent and child
will wait

Join Implementation (Attempt 3)

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1                 // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

Parent: w x y z

Child: a b c d

Rule of Thumb 2

- Acquire lock before checking/updating state and subsequent calling of wait/signal

Condition Variables

- `Pthread_cond_wait(cond_t *cv, mutex_t *lock)`
 - assumes the lock is held when `wait()` is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- `Pthread_cond_signal(cond_t *cv)`
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return, doing nothing

Producer Consumer Problem

- Class of problems where producer generates data/jobs and consumer consumes/services
- Synchronization is required among producers and consumers

Example UNIX Pipes

- A pipe may have many writers and readers
- Internally, there is a finite-sized buffer
- Writers add data to the buffer
 - Writers have to wait if buffer is full
- Readers remove data from the buffer
 - Readers have to wait if buffer is empty

Example UNIX Pipes

Implementation Outline:

- reads/writes to buffer require locking
- when buffers are full, writers must wait
- when buffers are empty, readers must wait

Producer Consumer Solution

Simple case:

- One producer thread
- One consumer thread
- Shared buffer of size 1

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {
```

```
        do_fill(i);
```

```
    }
```

```
}
```

```
void *consumer(void *arg) {  
    while(1) {
```

```
        int tmp = do_get();
```

```
        printf("%d\n", tmp);
```

```
    }
```

```
}
```

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
  
        do_fill(i);  
  
        Mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
  
        int tmp = do_get();  
  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
  
        Mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    → for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    → while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    → for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        → Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    → for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        → Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    → for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]


```
void *producer(void *arg) {  
    → for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]


```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
         Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
             Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        → if(numfull == max)  
           Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        → do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=1

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        → Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=1


[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```




[RUNNABLE]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```




numfull=1

[RUNNING]


```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
         Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNABLE]


```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
             Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
         if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNABLE]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
             Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=1

[BLOCKED]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            → Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNABLE]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=1

[BLOCKED]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            → Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        → int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[BLOCKED]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            → Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        → Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            → Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNING]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        → Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNABLE]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            → Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

numfull=0

[RUNNING]

```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        → do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[BLOCKED]

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

How about 2 consumers?

- Will the previous code work?


```


void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        Mutex_lock(&m); // p1
        if(numfull == max) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(i); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}

```

```

void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        if(numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        printf("%d\n", tmp); // c7
    }
}

```



does the last signal wake producer or consumer2?

Producer Consumer: Two CVs

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        Mutex_lock(&m); // p1  
        if (numfull == max) // p2  
            Cond_wait(&empty, &m)  
        do_fill(i); // p4  
        Cond_signal(&fill); // p5  
        Mutex_unlock(&m); //p6  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&fill, &m)  
        int tmp = do_get();  
        Cond_signal(&empty);  
        Mutex_unlock(&m);  
    }  
}
```

Producer Consumer: Two CVs

```
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        Mutex_lock(&m); // p1
        if (numfull == max) // p2
            Cond_wait(&empty, &m)
        do_fill(i); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        if (numfull == 0)
            Cond_wait(&fill, &m)
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

Can you find another bad schedule?

1. consumer1 waits because numfull == 0
2. producer increments numfull, wakes consumer1
3. before consumer1 runs, consumer2 runs, grabs entry, sets numfull=0.
4. consumer1 then reads bad data.

Producer Consumer: Two CVs and while

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        Mutex_lock(&m);  
        while(numfull == max)  
            Cond_wait(&empty, &m);  
        do_fill(i);  
        Cond_signal(&fill);  
        Mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        while (numfull == 0)  
            Cond_wait(&fill, &m);  
        int tmp = do_get();  
        Cond_signal(&empty);  
        Mutex_unlock(&m);  
    }  
}
```

Is this correct?

Correct

- no concurrent access to shared state
- every time lock is acquired, assumptions are reevaluated
- a consumer will get to run after every do_fill()
- a producer will get to run after every do_get()

Good Rule of Thumb 3

- Whenever a lock is acquired, recheck assumptions about state!
- Possible for another thread to grab lock in between signal and wakeup from wait
- Note that some libraries also have "spurious wakeups" (may wake multiple waiting threads at signal or at any time)

RULES OF THUMB FOR CVs

- Keep state in addition to CV's
- Always do wait/signal with lock held
- Whenever thread wakes from waiting, recheck state

Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.