

# Concurrency: Semaphores

Sridhar Alagar

# Synchronization Objectives

- Mutual exclusion (e.g., A and B don't run at same time in CS)
  - solved with locks
- Ordering (e.g., B runs after A does something)
  - solved with condition variables

# Condition Variables

- `Pthread_cond_wait(cond_t *cv, mutex_t *lock)`
  - assumes the lock is held when `wait()` is called
  - puts caller to sleep + releases the lock (atomically)
  - when awoken, reacquires lock before returning
- `Pthread_cond_signal(cond_t *cv)`
  - wake a single waiting thread (if  $\geq 1$  thread is waiting)
  - if there is no waiting thread, just return, doing nothing

# Join Implementation

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1                 // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

# Producer Consumer Problem

- Class of problems where producer generates data/jobs and consumer consumes/services
- Synchronization is required among producers and consumers

# Producer Consumer Solution

Simple case:

- One producer thread
- One consumer thread
- Shared buffer of size 1

# Producer Consumer: Two CVs and while

```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        while(numfull == max)  
            Cond_wait(&empty, &m);  
        do_fill();  
        Cond_signal(&fill);  
        Mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        while (numfull == 0)  
            Cond_wait(&fill, &m);  
        int tmp = do_get();  
        Cond_signal(&empty);  
        Mutex_unlock(&m);  
    }  
}
```

Is this correct?

Correct

- no concurrent access to shared state
- every time lock is acquired, assumptions are reevaluated
- a consumer will get to run after every do\_fill()
- a producer will get to run after every do\_get()

# Semaphores - State variable + Queue

- CVs have no state variable other than a queue
  - Programmer has to maintain it
- Semaphores have state - an integer variable
  - User cannot access state directly other than initializing it
  - Behaves based on its state



# Semaphores - operations

- Initialization

```
sem_t sem;
```

```
sem_init(sem_t *s, int initval) {  
    s->value = initval;  
}
```

- user cannot read or write after initialization

# Semaphores - operations

- Wait or P()

```
sem_wait(sem_t *s) {
```

```
    Decrements the value of semaphore s by 1;  
    wait if the value of semaphore s is -ve
```

```
}
```

- When value of semaphore is -ve, the absolute value denotes the number of waiting thread

# Semaphores - operations

- Signal or V()

```
sem_post(sem_t *s) {  
    increments the value of semaphore s by 1;  
    wakeup one thread if the value of  
    semaphore s is  $\leq 0$   
}
```

# Locks using Semaphores (binary)

```
sem_t m;  
sem_init(&m, 0, X); // initialize semaphore to X;  
  
sem_wait(&m);  
//critical section here  
sem_post(&m);
```

What should be the value of X?

# Locks using Semaphores (binary)

```
sem_t m;  
sem_init(&m, 0, 1); // initialize semaphore to X;  
  
sem_wait(&m);  
//critical section here  
sem_post(&m);
```

# Trace: two threads - binary semaphore

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() returns	Running		Ready
0	(crit sect: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	<i>Switch → T0</i>	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

# Join with CVs vs Semaphores

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                 // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

Semaphores:

```
sem_t s;  
sem_init(&s, ??)  
void thread_join() { sem_wait(&s); }  
Void thread_exit() { sem_post(&s); }
```

# Trace 1: Parent waiting for Child

Parent calls `sem_wait()` before child calls `sem_post()`

Value	Parent	State	Child	State
0	Create Child	Running	<i>(Child exists; is runnable)</i>	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(sem < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch</i> →Child	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch</i> →Parent	Ready
0	<code>sem_wait()</code> returns	Running		Ready



# Trace 2: Parent waiting for Child

Child calls `sem_post()` before parent calls `sem_wait()`

Value	Parent	State	Child	State
0	Create (Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake (nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	(sem>=0)→awake	Running		Ready
0	<code>sem_wait()</code> returns	Running		Ready

# Producer/Consumer: Semaphores #1

Simplest case:

- Single producer thread, single consumer thread
- Single shared buffer between producer and consumer

Requirements

- Consumer must wait for producer to fill buffer
- Producer must wait for consumer to empty buffer (if filled)

Requires 2 semaphores

- `emptyBuffer`: Initialize to ???  $1 \rightarrow 1$  empty buffer; producer can run 1 time first
- `fillBuffer`: Initialize to ???  $0 \rightarrow 0$  full buffer; consumer can run 0 times first

Producer

```
While (1) {  
  
    Fill(&buffer);  
  
}
```

Consumer

```
While (1) {  
  
    Use(&buffer);  
  
}
```

# Producer/Consumer: Semaphores #1

Simplest case:

- Single producer thread, single consumer thread
- Single shared buffer between producer and consumer

Requirements

- Consumer must wait for producer to fill buffer
- Producer must wait for consumer to empty buffer (if filled)

Requires 2 semaphores

- `emptyBuffer`: Initialize to ???  $1 \rightarrow 1$  empty buffer; producer can run 1 time first
- `fillBuffer`: Initialize to ???  $0 \rightarrow 0$  full buffer; consumer can run 0 times first

Producer

```
While (1) {  
  
    sem_wait(&emptyBuffer);  
    Fill(&buffer);  
  
    sem_post(&fillBuffer);  
  
}
```

Consumer

```
While (1) {  
  
    sem_wait(&fillBuffer);  
    Use(&buffer);  
  
    sem_post(&emptyBuffer);  
  
}
```

# Producer/Consumer: Semaphores #2

Next case: Circular buffer

- Single producer thread, single consumer thread
- Shared buffer of size N between producer and consumer

Requires 2 semaphores

- `emptyBuffer`: Initialize to ???  $N \rightarrow N$  empty buffer; producer can run N times first
- `fillBuffer`: Initialize to ???  $0 \rightarrow 0$  full buffer; consumer can run 0 time first

## Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer);
    i = (i+1)%N;
    sem_post(&fillBuffer);
}
```

## Consumer

```
j = 0;
while (1) {
    sem_wait(&fillBuffer);
    Use(&buffer);
    j = (j+1)%N;
    sem_post(&emptyBuffer);
}
```

# Producer/Consumer: Semaphores #3

Final case:

- Multiple producer thread, Multiple consumer thread
- Shared buffer of size N between producer and consumer

Requires 2 semaphores

- emptyBuffer: Initialize to ???  $N \rightarrow N$  empty buffer; producer can run N times first
- fillBuffer: Initialize to ???  $0 \rightarrow 0$  full buffer; consumer can run 0 time first

Why will the previous code (shown below) not work? **Need Mutex**

Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer);
    i = (i+1)%N;
    sem_post(&fillBuffer);
}
```

Consumer

```
j = 0;
while (1) {
    sem_wait(&fillBuffer);
    Use(&buffer);
    j = (j+1)%N;
    sem_post(&emptyBuffer);
}
```

# Producer/Consumer: Semaphores #3

Final case:

- Multiple producer thread, Multiple consumer thread
- Shared buffer of size N between producer and consumer

Mutex added. Will this work?

Deadlock

Producer

```
i = 0;
while (1) {
    sem_wait(&mutex);
    sem_wait(&emptyBuffer);
    Fill(&buffer);
    i = (i+1)%N;
    sem_post(&fillBuffer);
    sem_post(&mutex);
}
```

Consumer

```
j = 0;
while (1) {
    sem_wait(&mutex);
    sem_wait(&fillBuffer);
    Use(&buffer);
    j = (j+1)%N;
    sem_post(&emptyBuffer);
    sem_post(&mutex);
}
```

# Producer/Consumer: Semaphores #3

Final case:

- Multiple producer thread, Multiple consumer thread
- Shared buffer of size N between producer and consumer

Another version with mutex. Will this work? Works

Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    sem_wait(&mutex);
    Fill(&buffer);
    i = (i+1)%N;
    sem_post(&mutex);
    sem_post(&fillBuffer);
}
```

Consumer

```
j = 0;
while (1) {
    sem_wait(&fillBuffer);
    sem_wait(&mutex);
    Use(&buffer);
    j = (j+1)%N;
    sem_post(&mutex);
    sem_post(&emptyBuffer);
}
```

# Readers Writers Problem (RWP)

- A data set is shared among a number of concurrent processes/threads
  - Readers - only read the data set; they do *not* perform any updates
  - Writers - can both read and write
- Problem
  - allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time



# RW Problem: Solution structure

The structure of a writer process

```
do {  
  
    ...  
    // writing is performed  
    ...  
  
} while (true);
```

The structure of a reader process

```
do {  
  
    ...  
    // reading is performed  
    ...  
  
} while (true);
```

# RW Problem: Ensure ME

Writer process

```
do {  
  
    wait(rw_mutex);  
  
    ...  
    // writing is performed  
    ...  
    signal(rw_mutex);  
  
} while (true);
```

Reader process

```
do {  
  
    wait(rw_mutex);  
  
    ...  
    // reading is performed  
    ...  
    signal(rw_mutex);  
  
} while (true);
```

# RW Problem: Allow multiple readers

Writer process

```
do {  
  
    wait(rw_mutex);  
  
    ...  
    // writing is performed  
    ...  
  
    signal(rw_mutex);  
  
} while (true);
```

Reader process

```
do {  
  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
  
    ...  
    // reading is performed  
  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
  
} while (true);
```

# RW Problem: Allow multiple readers

Writer process

**Writer starves**

```
do {  
  
    wait(rw_mutex);  
  
    ...  
    // writing is performed  
    ...  
  
    signal(rw_mutex);  
  
} while (true);
```

Reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    // reading is performed  
  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
  
} while (true);
```

# RW Problem: Starvation free


Writer process

```
do {  
    wait(rw_line);  
    wait(rw_mutex);  
    signal(rw_line);  
  
    ...  
    // writing is performed  
  
    ...  
    signal(rw_mutex);  
  
} while (true);
```

Reader process

```
do {  
    wait(rw_line);  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(rw_line);  
    signal(mutex);  
  
    ...  
    // reading is performed  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
  
} while (true);
```

# Dining Philosophers Problem



- Philosophers spend their lives alternating between thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat
  - Need both to eat, then release both when done

# Dining Philosophers - Solution 1

```
sem_t chopstick[5];
```

Philosopher *i*:

```
do {  
    wait (chopstick[i]);  
    wait (chopstick[(i + 1) % 5]);  
  
    // eat  
  
    signal (chopstick[i]);  
    signal (chopstick[(i + 1) % 5]);  
  
    // think  
  
} while (TRUE);
```

- What is the problem with this algorithm? **deadlock**

# Dining Philosophers Problem

- Deadlock free solution:
  - Allow a philosopher to pick up the chopsticks only if both are available (picking must be done in a critical section).
  - Use an asymmetric solution
    - Odd-numbered philosopher picks up the left chopstick first and then the right chopstick.
    - Even-numbered philosopher picks up the right chopstick first and then the left chopstick.
  - Another asymmetric solution
    - The last philosopher picks up the left chopstick first and then the right chopstick.
    - Other philosophers pick up the right chopstick first and then the left chopstick.



# Semaphore Implementation

```
typedef struct{  
    int value;  
    struct process *q;  
} semaphore;
```

# Semaphore Implementation

```
wait(semaphore *S) {  
    disable interrupts;  
    S->value--;  
    if (S->value < 0) {  
        add this proc to S->q;  
        block();  
    }  
    enable interrupts;  
}
```

```
signal(semaphore *S) {  
    disable interrupts;  
    S->value++;  
    if (S->value < 0) {  
        remove proc P from S->q;  
        wakeup(p);  
    }  
    enable interrupts;  
}
```

# Semaphore Implementation

- Semaphore can be implemented using locks and CVs
  - Homework: Read the book

# Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.