


Concurrency: Common Bugs

Sridhar Alagar

Concurrency is hard



- Lu et al. study: Analyzed a sample of around 100 bugs from among > 500k bugs

Atomicity Violation

Thread 1:

```
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info, ...);  
    ...  
}
```

Thread 2:

```
thd->proc_info = NULL;
```

- What is wrong?
 - Test (if()) and Set (fputs()) should be atomic

Atomicity Violation: Fix with locks

Thread 1:

```
pthread_mutex_lock(&lock);
if (thd->proc_info) {
    ...
    fputs(thd->proc_info, ...);
pthread_mutex_unlock(&lock);
    ...
}
```

Thread 2:

```
pthread_mutex_lock(&lock);
thd->proc_info = NULL;
pthread_mutex_unlock(&lock);
```

Order Violation

Thread 1:

```
mThread = CreateThread(...);
```

...

Thread 2:

```
state = mThread->state;
```

- What is wrong?
 - `createThread()` should be executed before thread access

Order Violation: Fix

Thread 1:

```
mThread = CreateThread(...);
```

```
sem_post(s);
```

```
...
```

```
}
```

Thread 2:

```
sem_wait(s);
```

```
thd->proc_info = NULL;
```

Deadlock

No progress can be made because two or more threads are waiting for the other to take some action and thus neither ever does

Cooler name: DEADLY EMBRACE (Dijkstra)

Deadlock Bug

Thread 1:

```
lock (&L1) ;  
lock (&L2) ;
```


Thread 2:

```
lock (&L2) ;  
lock (&L1) ;
```

Circular wait

Wait For Graph(WFG)

Cycle \Leftrightarrow Deadlock



Necessary Conditions for Deadlock

- Mutual Exclusion
 - Only one thread can use a resource at a time
- Hold and Wait
 - Threads holding resources waiting for resources held by other threads
- No Preemption
 - Resources cannot be forcibly removed from threads holding it
- Circular wait
 - There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

Deadlock Prevention

- Prevent deadlock by ensuring that one of the conditions does not hold
- Constrains how resources are acquired. Incurs extra cost:
 - overhead
 - reduction in concurrency

Prevention - Eliminate Circular Wait

- Provide a total ordering of lock/resource acquisition
 - Cannot acquire a resource that is earlier in the order than any of the resources currently held

Thread 1:

```
lock (&L1) ;  
lock (&L2) ;
```

Thread 2:

```
lock (&L1) ;  
lock (&L2) ;
```

Prevention - Eliminate Hold and Wait

- Acquire all locks at once, atomically

```
1    lock(prevention);  
2    lock(L1);  
3    lock(L2);  
4    ...  
5    unlock(prevention);
```

- Problems:
 - Need to know the locks that will be acquired in the future
 - Reduces concurrency

Prevention - Preempt

- If cannot get the lock needed, release the locks held

```
1  top:
2      lock(L1);
3      if( tryLock(L2) == -1 ) {
4          unlock(L1);
5          goto top;
6      }
```

- Problems:
 - Livelock
 - use exponential back-off

Prevention - Wait Free Mutual Exclusion

- Using hardware atomic instruction build data structure that does not require explicit locking for update

```
1  int CompareAndSwap(int *address, int expected, int new) {
2      if(*address == expected) {
3          *address = new;
4          return 1; // success
5      }
6      return 0;
7  }
```

Prevention - Wait Free Mutual Exclusion

- Want to atomically increment a the value of a variable by a certain amount

```
1  void AtomicIncrement(int *value, int amount) {  
2      do{  
3          int old = *value;  
4      }while( CompareAndSwap(, , ) == 0 );  
5  }
```


Prevention - Wait Free Mutual Exclusion

- Want to atomically increment a the value of a variable by a certain amount

```
1  void AtomicIncrement(int *value, int amount) {  
2      do{  
3          int old = *value;  
4      }while( CompareAndSwap(value, old, old+amount)==0);  
5  }
```

- No lock acquired; so no deadlock

Prevention - Wait Free Mutual Exclusion

- A more complex example: list insertion

```
1  void insert(int value) {  
2      node_t * n = malloc(sizeof(node_t));  
3      assert( n != NULL );  
4      n->value = value ;  
5      n->next = head;  
6      head = n;  
7  }
```

- If called by multiple threads concurrently, race condition can arise

Prevention - Wait Free Mutual Exclusion

- Solution using locks

```
1  void insert(int value) {  
2      node_t * n = malloc(sizeof(node_t));  
3      assert( n != NULL );  
4      n->value = value ;  
5      lock(listlock); // begin critical section  
6      n->next  = head;  
7      head = n;  
8      unlock(listlock) ; //end critical section  
9  }
```

Prevention - Wait Free Mutual Exclusion

- Solution in wait-free manner

```
1  void insert(int value) {
2      node_t *n = malloc(sizeof(node_t));
3      assert(n != NULL);
4      n->value = value;
5      do {
6          n->next = head;
7      } while (!CompareAndSwap(&head, n->next, n));
8  }
```

Deadlock Avoidance: via Scheduling


- Requires global knowledge about thread locks/resource usage
- Schedule threads in such a way deadlocks can be avoided

| | T1 | T2 | T3 | T4 |
|----|-----|-----|-----|----|
| L1 | yes | yes | no | no |
| L2 | yes | yes | yes | no |

Deadlock Avoidance: via Scheduling

| | T1 | T2 | T3 | T4 |
|----|-----|-----|-----|----|
| L1 | yes | yes | no | no |
| L2 | yes | yes | yes | no |

- A deadlock avoiding schedule



Deadlock Avoidance: via Scheduling

| | T1 | T2 | T3 | T4 |
|----|-----|-----|-----|----|
| L1 | yes | yes | yes | no |
| L2 | yes | yes | yes | no |

- A deadlock avoiding schedule



Detect and Recover

- Allow deadlock to occasionally occur and then take some action.
 - Example: if an OS froze, you would reboot it.
- Many database systems employ deadlock detection and recovery technique.
 - A deadlock detector runs periodically
 - Build a resource graph and checking it for cycles
 - If in deadlock, restart the system

Disclaimer

- Some of the materials in this lecture slides are from the lecture slides by Prof. Andrea, Prof. Youjip, and other educators. Thanks to all of them.