

Persistence: Files

Sridhar Alagar

File

- Sequence of persistent bytes that can be read/written
- Logical storage unit with contiguous logical address space
- A **file system** organizes and stores file
- A file has many attributes

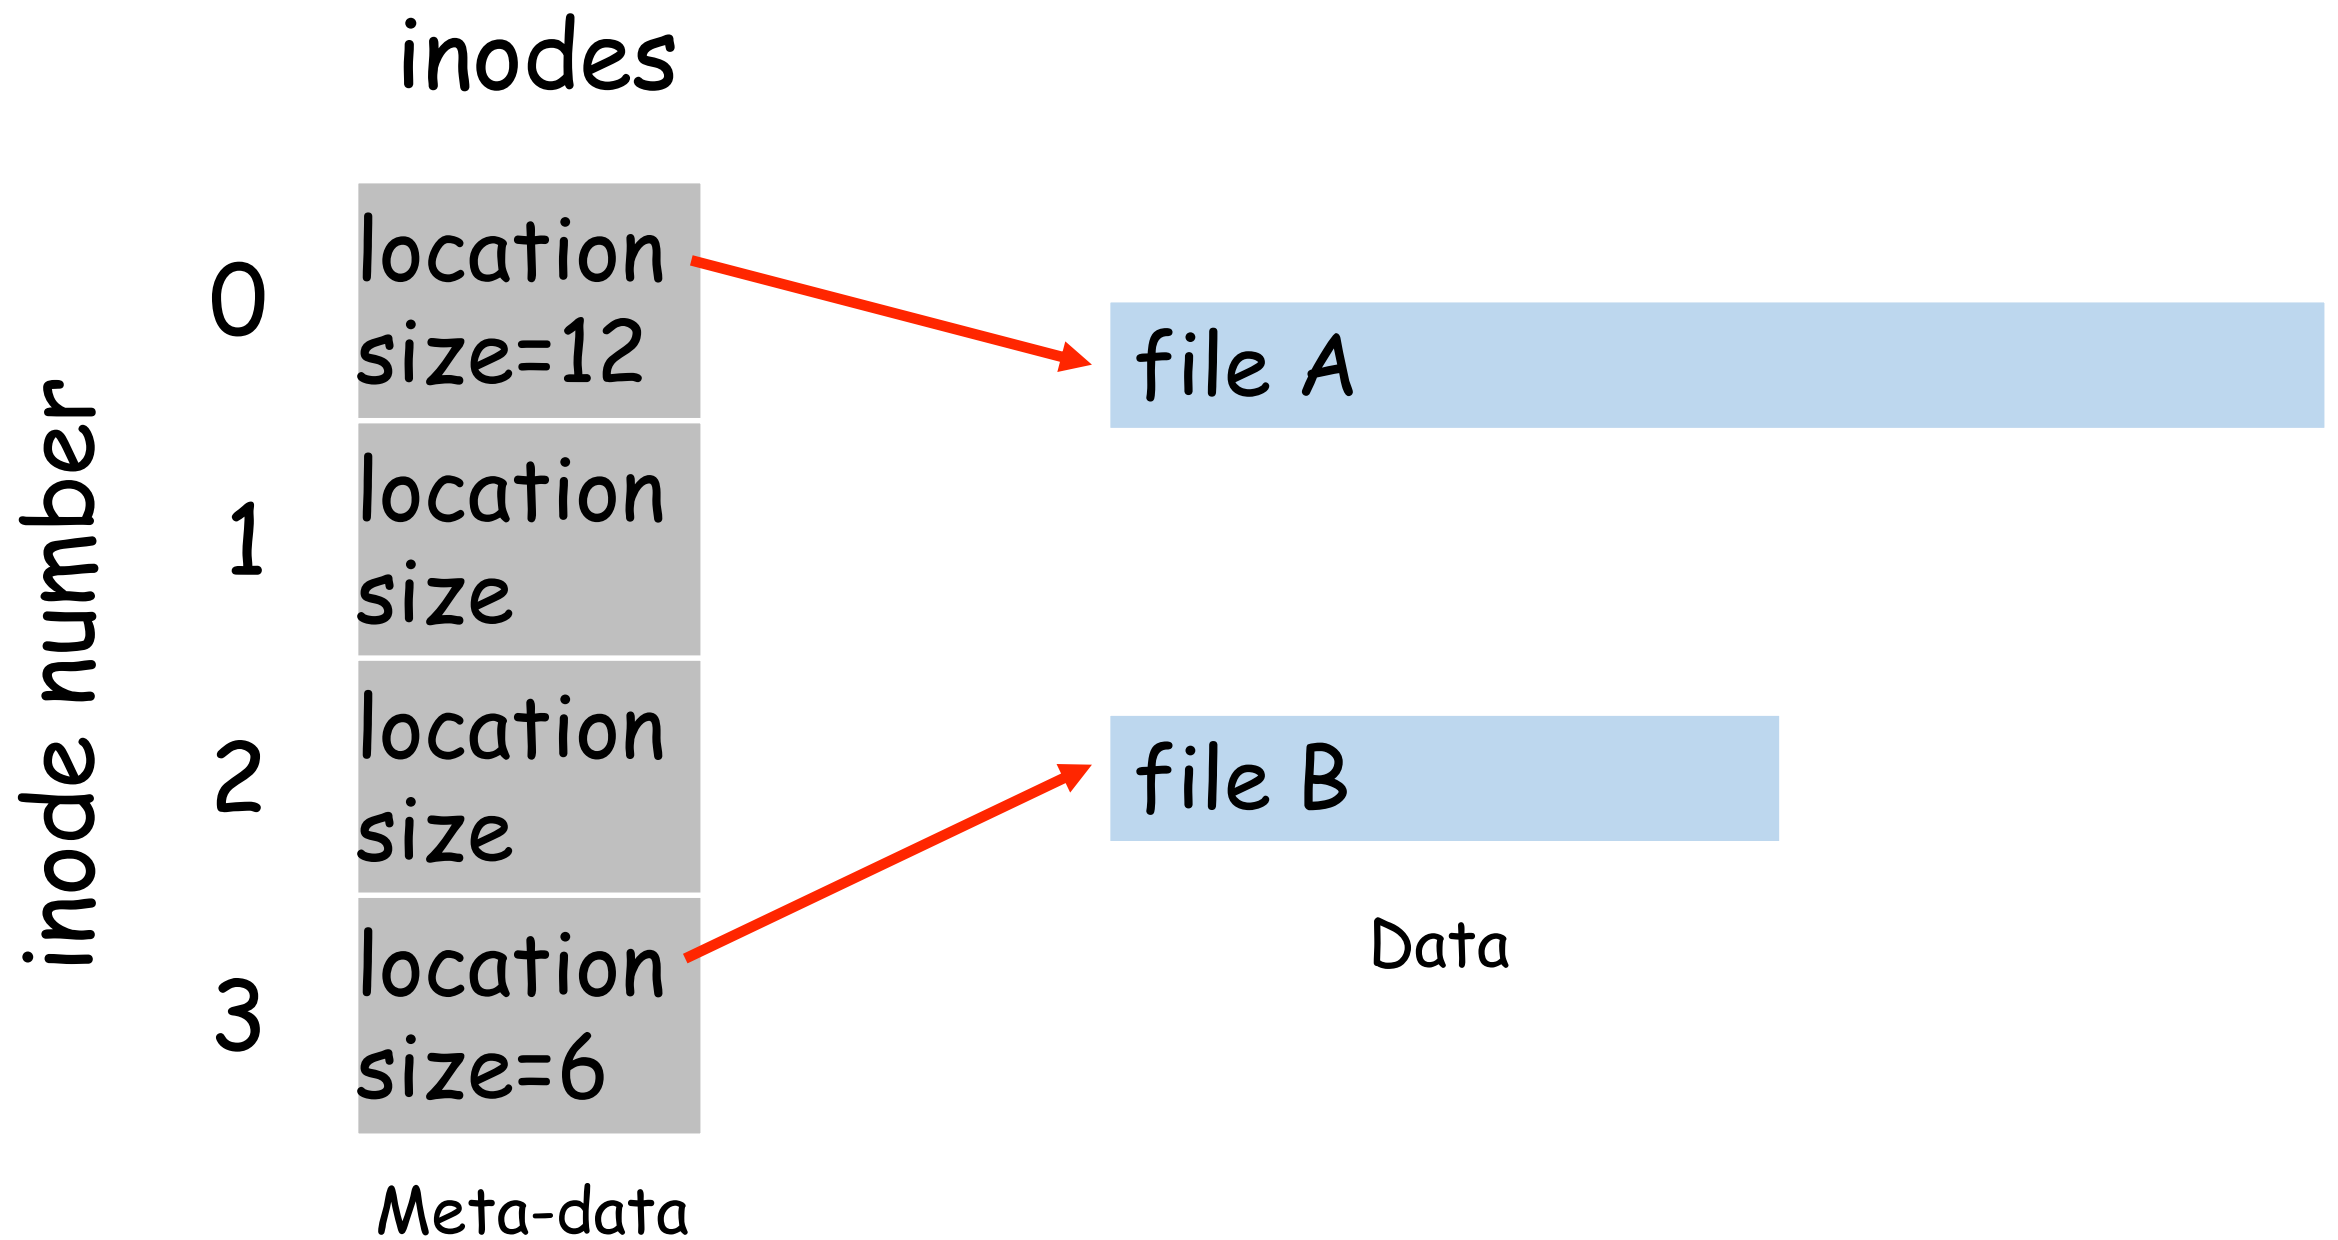
File attributes (meta data)

- **Name** - only information kept in human-readable form
- **Identifier** - unique number identifies file within file system
- **Location** - pointer to file location on device
- **Size** - current file size
- **Protection** - controls who can do reading, writing, executing
- **Time, date, and user identification** - data for protection, security, and usage monitoring

Meta data should be persistent

Inode

- Inode is an on-disk data structure that contains a file's attributes
- Every file has a unique inode
- Inodes are kept in inode table
- Inode number is the file identification number within the file system



How to access a file?


- Find its inode
- Need mapping between file name and file's inode number
- Where will this mapping be stored?

Directory

Try `ls -i`

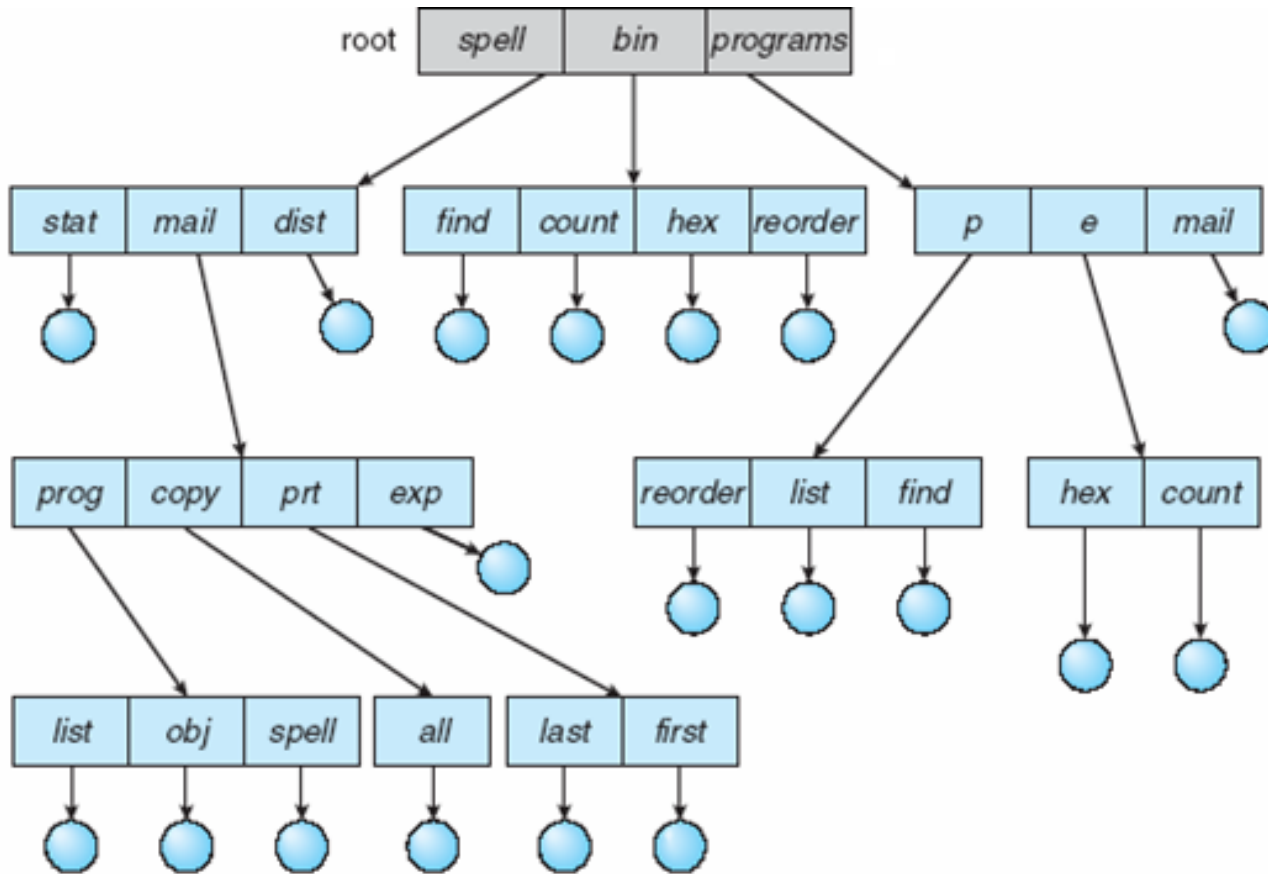
Directory Structure

- Single level



- Naming problem
- Grouping problem

Tree Structured Directories



- File name is specified with the path
/bin/count
/spell/mail/prog/list
- File (path) name is unique

Where is a directory stored?

- In the persistent store as a file
- Every entry in the directory is a file or a (sub)directory

File API (Attempt 1)

read(char* pathname, void *buf, int offset, size_t nbyte)

write(char* pathname, void *buf, int offset, size_t nbyte)

- Disadvantages?
 - Need to find the file every time read/write performed
 - traverse the directory tree, get the inode and then the file
 - User need to remember the current offset
- Solution
 - Open it once and maintain the state
 - in a file descriptor table

File Descriptor Table (per process)

- Is an array of file descriptors, maintained in proc structure
- xv6 definition of file descriptor, and proc structure:

```
struct file { // file descriptor
    ...
    struct inode *ip;
    uint off;
};
// Per-process state
struct proc {
    ...
    struct file *ofile[NOFILE]; // Open files
    ...
}
```

Open once

```
fd = open("file.txt", flags)
```

- Searches the directory once
- Fetch the inode and keep it in memory
- Find a free entry in file table of `ofile[]`
- Update `ofile[fd]` to point to inode
- Set the offset value in `ofile[fd]`
- Subsequent file operations use the `fd`

File API

```
int open(char* path_name, int flags)
int read(int fd, void *buf, size_t nbyte)
int write(int fd, void *buf, size_t nbyte)
close(int fd);
```

- fd is the file descriptor returned by open()
- No need to search for the file
- Read/write performed starting from the current offset

Sample Code

```
int fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
int fd2 = open("file.txt"); // returns 4
int fd3 = dup(fd2);          // returns 5
Close (fd2);
```

fd table

0 ■
1 ■
2 ■
3 ■
4 ■
5 ■

File descriptors(fds)

offset = 0
inode =

inode

location = ...
size = ...

```
int fd1 = open("file.txt"); // returns 3
```

fd table

0	■
1	■
2	■
3	■
4	■
5	■


fds

offset = 12
inode =


inode

location = ...
size = ...


```
fd1 = open("file.txt"); // returns 3  
read(fd1, buf, 12);
```

```
fd1 = open("file.txt"); // returns 3  
read(fd1, buf, 12);  
fd2 = open("file.txt"); // returns 4
```



```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
```



```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
close(fd2);
```

fd table

0 ■
1 ■
2 ■
3 ■
4 ■
5 ■

fds

offset = 0
inode =

inode

location = ...
size = ...

```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
close(fd2);
close(fd1);
```

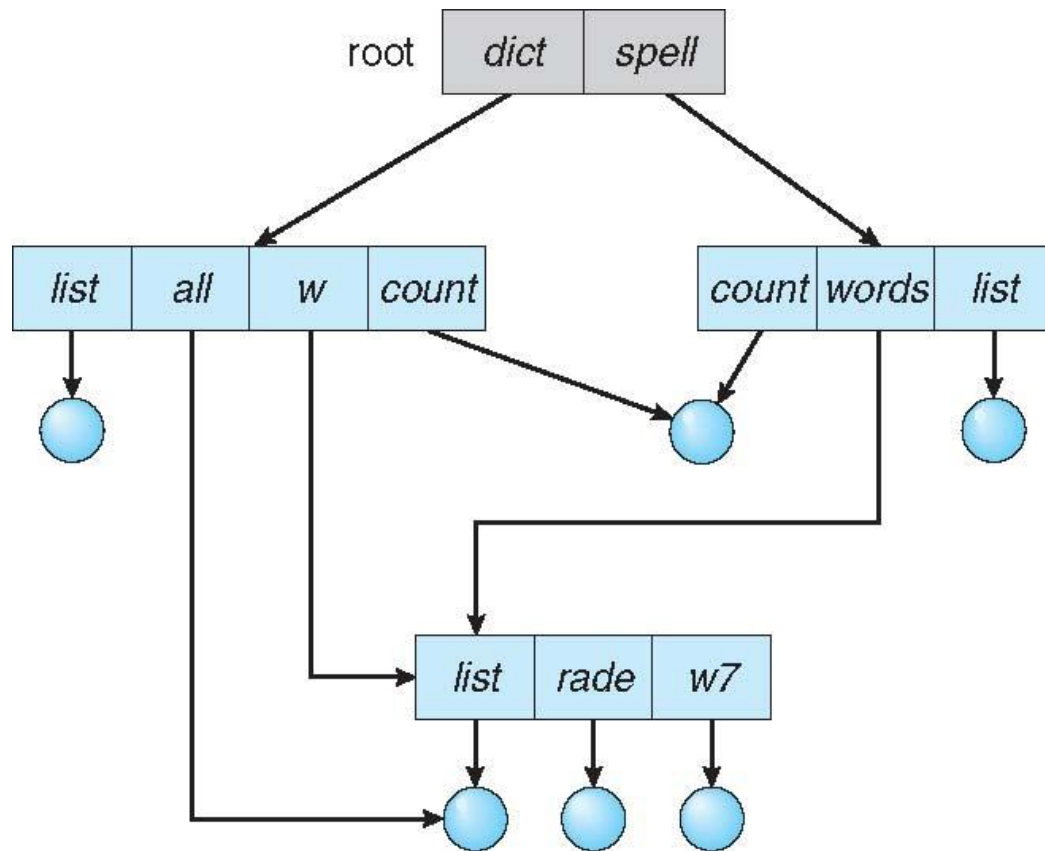
Links to Files

link(char* oldpath, char* newpath)

unlink(char* pathname)

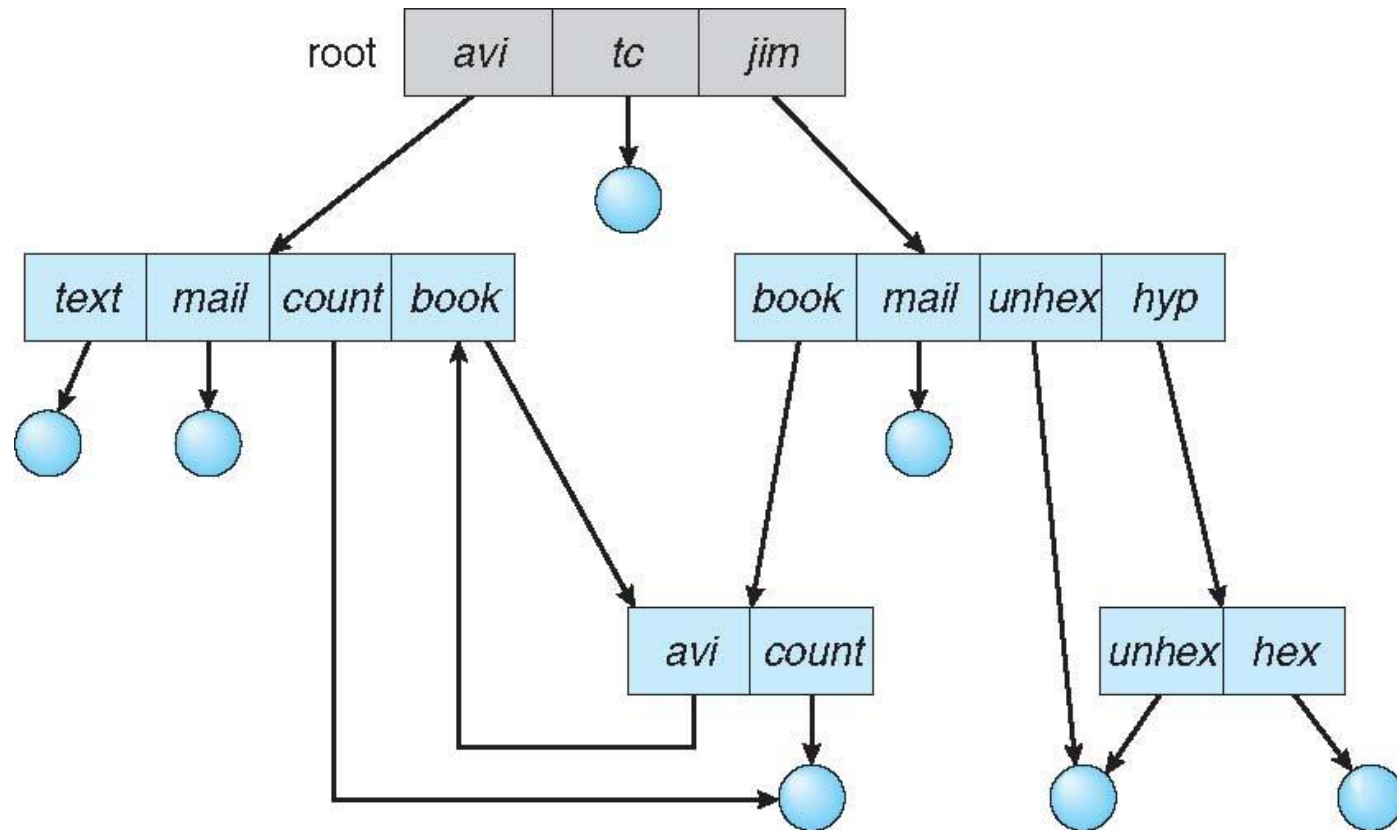
- link() creates a new (hard) link to existing file
 - inode not created
 - reference count increases
 - cannot link directories
- Unlink() removes link
 - If it is the last link to the file, deletes the file
- What is the directory structure now?

Acyclic Graph Directories



What is the directory structure when a symbolic link to a parent directory is created?

Directory structure - Graph (with cycles)



Files - More APIs

lseek(fd, offset, whence)

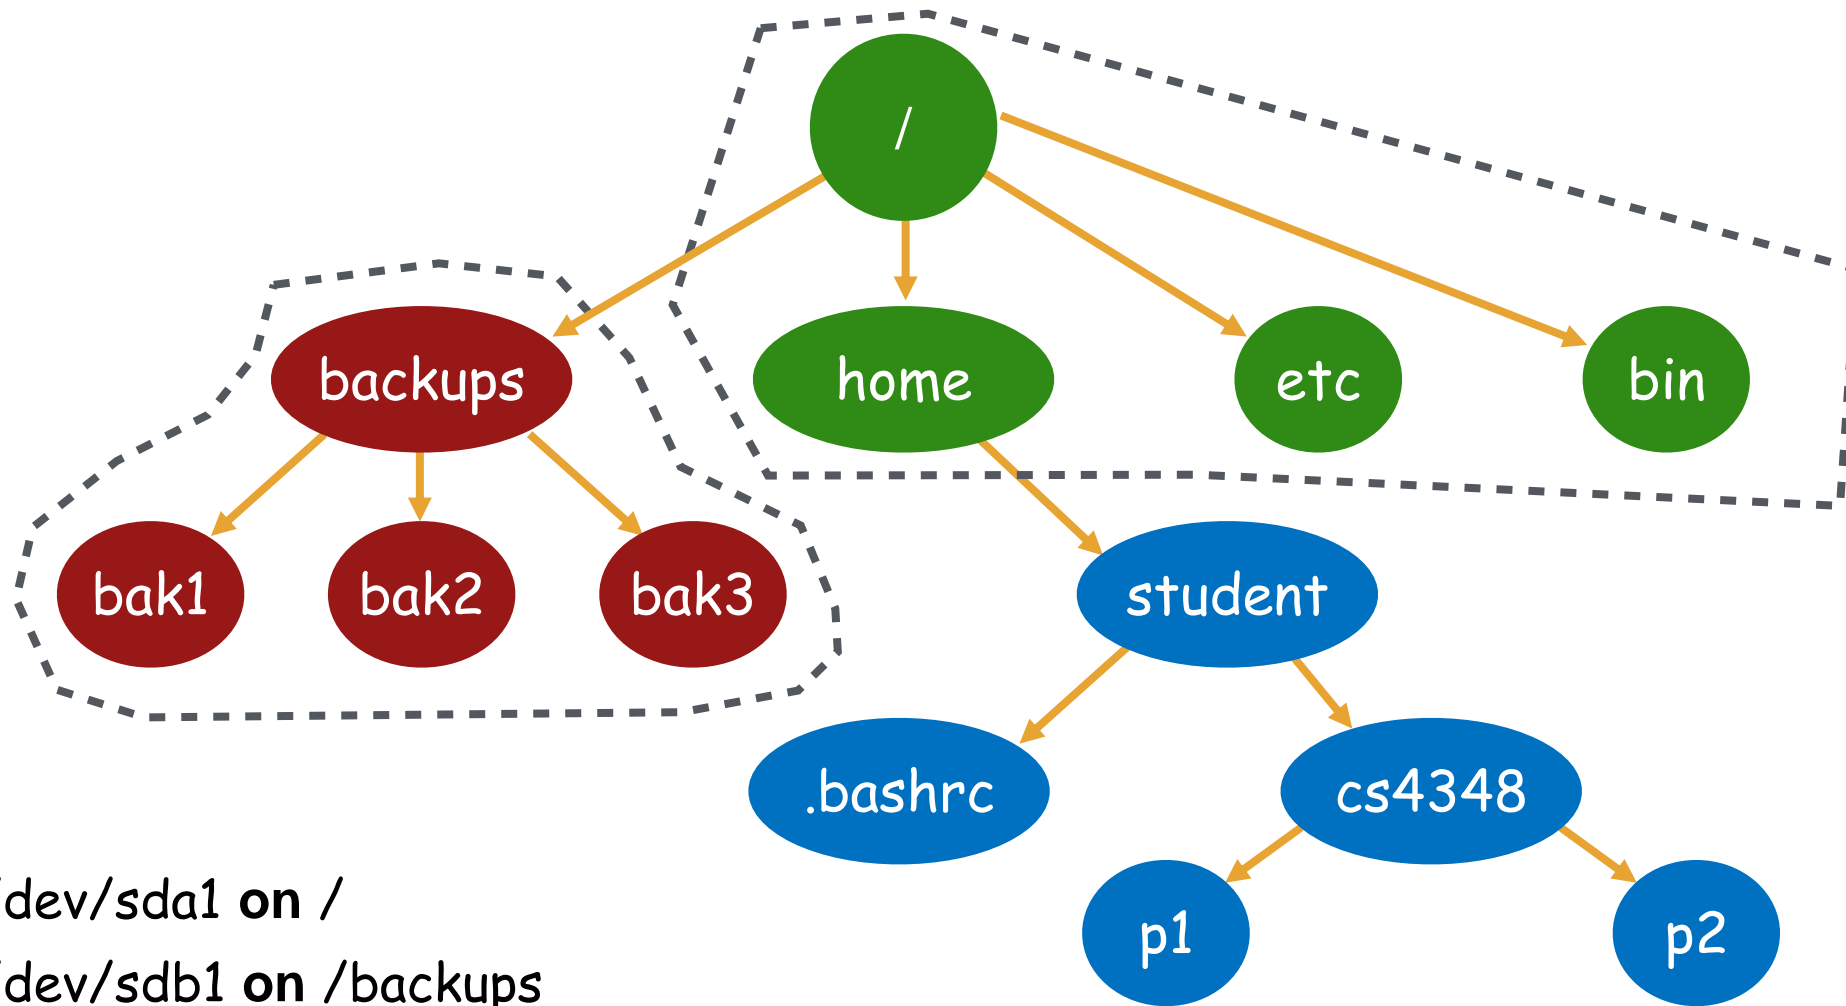
rename(oldname, newname)

stat(pathname, buf)

Fsync()

File System Tools

- `mkfs`: creates a new file system in a storage device
- `mount`: mounts another device with a file system onto an existing directory structure



- /dev/sda1 **on** /
- /dev/sdb1 **on** /backups
- hopper **on** /home