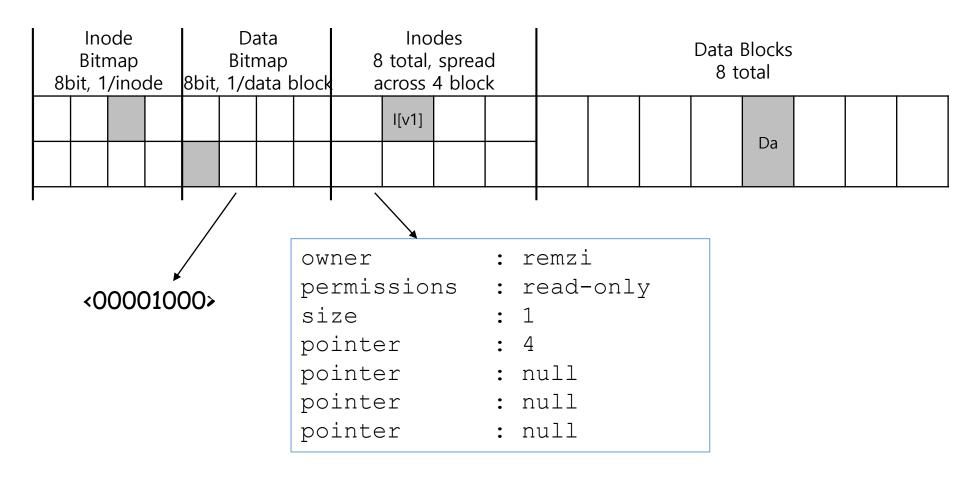
Persistence: Crash Consistency

Sridhar Alagar

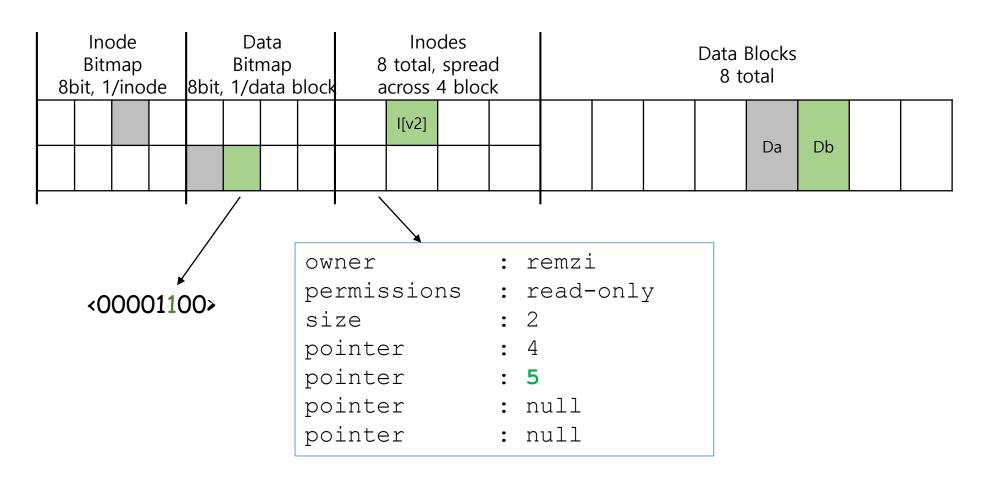
Crash Consistency

- An append to a file involves updating (non atomic)
 - Data blocks
 - Inode
 - Bitmap for data blocks
- What happens if the system crashes in the middle of updating the on-disk structures?
 - · Crash due to: power loss or kernel panic or reboot
- The file system can be in an inconsistent state

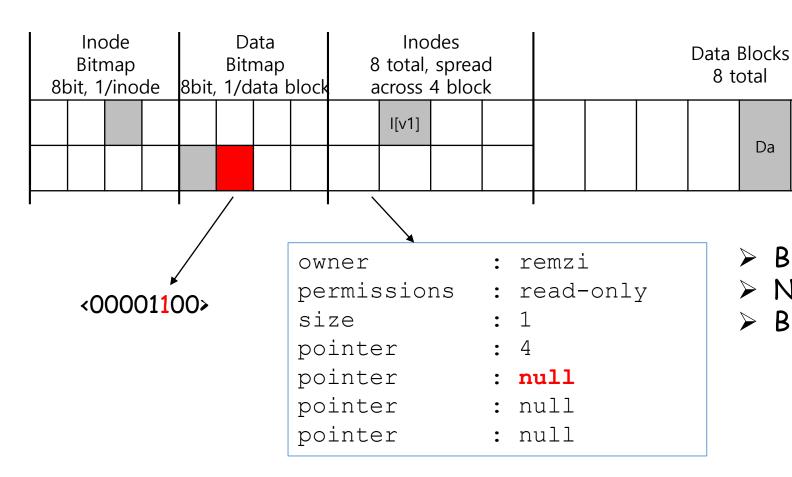
Example



Example: Append a new block of data



Updates bitmap and system crashes



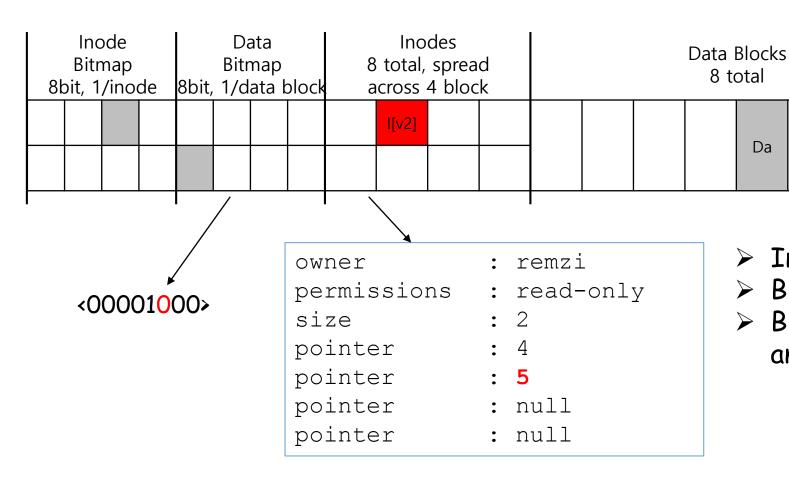
Bitmap says block 5 allocated

Inconsistent

- > None of the inodes point to 5
- ➤ Block 5 is lost

Da

Updates inode and system crashes



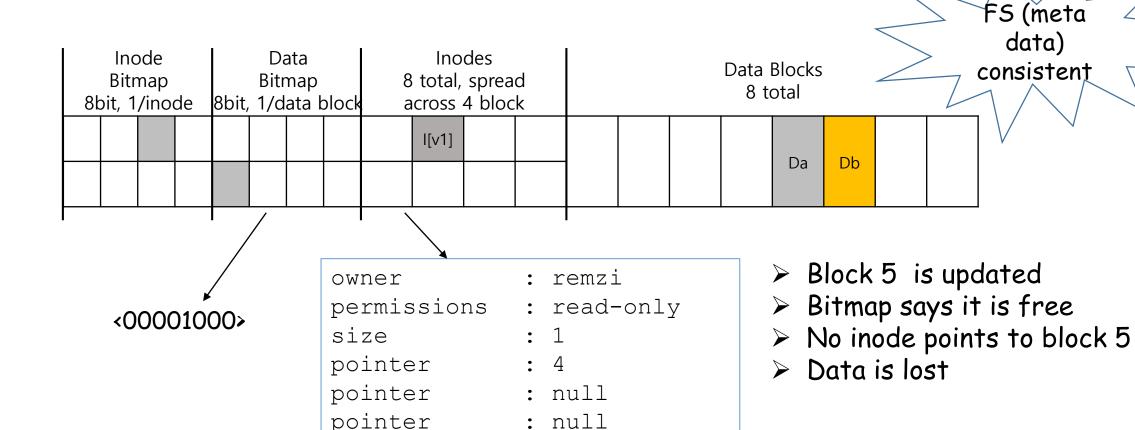
- > Inode says block 5 allocated
- > Bitmap says it is free

Da

> Block 5 may be allocated to another file; security breach

Inconsistent

Updates data block and system crashes



: null

pointer

Three more crash scenarios

 Inode and bitmap updated, but not data block

 Inode and data block updated but not bitmap

 Bitmap and data block updated not inode

- > Meta data consistent
- > But stale data in data block
- >FS inconsistent
- > Data block may be allocated to another file
- >Similar to update inode only
- >FS inconsistent
- > Data block lost
- > Similar to update bitmap only

Crash Consistency Problem

 Transition the FS from one state to another state atomically

- Not possible
 - · Disk commits one write at a time
 - Crash in the middle of update
- Need solutions to resolve inconsistencies after system recovers from crash

File System Checker (FSCK)

 After crash, scan the entire disk for inconsistencies and resolve them

- How to check if bitmap data is consistent?
 - · Read every valid inode's direct + indirect block pointers
 - If pointer to data block, the corresponding bit should be 1; else bit is 0

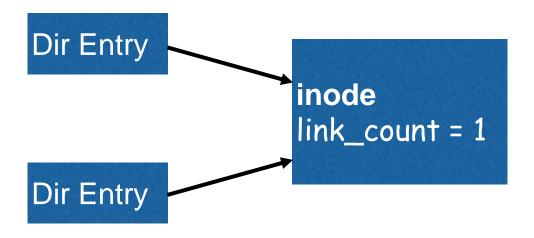
FSCK Checks

- Hundreds of types of checks over different fields...
- Do superblocks match?
- Do directories contain "." and ".."?
- Do number of dir entries equal inode link counts?
- Do different inodes ever point to same block?

• ...

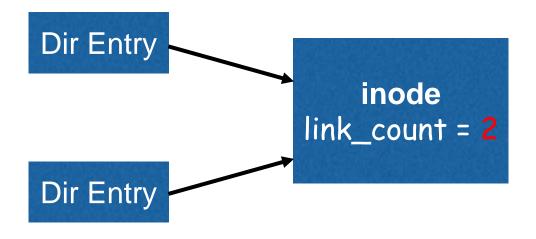
How to fix inconsistencies?

Link Count (example 1)



How to fix to have consistent file system?

Link Count (example 1)



How to fix to have consistent file system? Simple fix. Change link_count = 2

Link Count (example 2)

inode link_count = 1

No dir entry. How to fix?

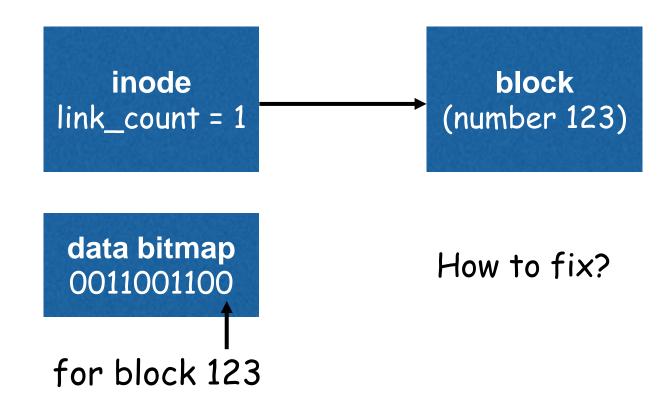
Link Count (example 2)

Create dir entry.
Where to put the entry?

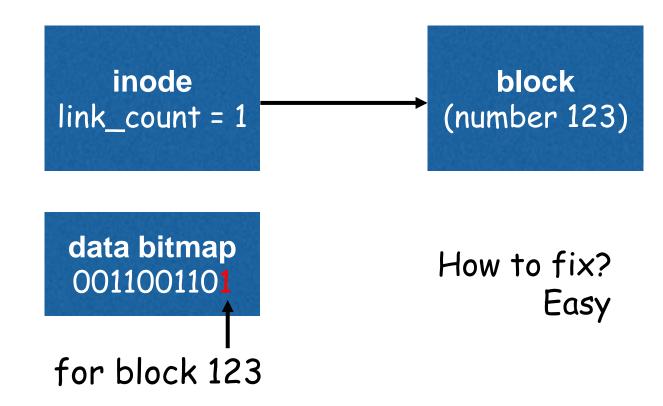
Dir Entry
inode
link_count = 1

```
ls -I /
total 150
drwxr-xr-x 401 18432 Dec 31 1969 afs/
drwxr-xr-x. 2 4096 Nov 3 09:42 bin/
drwxr-xr-x. 5 4096 Aug 1 14:21 boot/
dr-xr-xr-x. 13 4096 Nov 3 09:41 lib/
dr-xr-xr-x. 10 12288 Nov 3 09:41 lib64/
drwx----. 2 16384 Aug 1 10:57 lost+found/
```

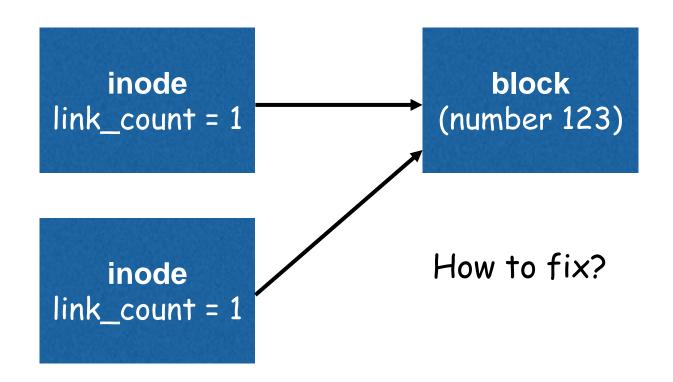
Data Bitmap



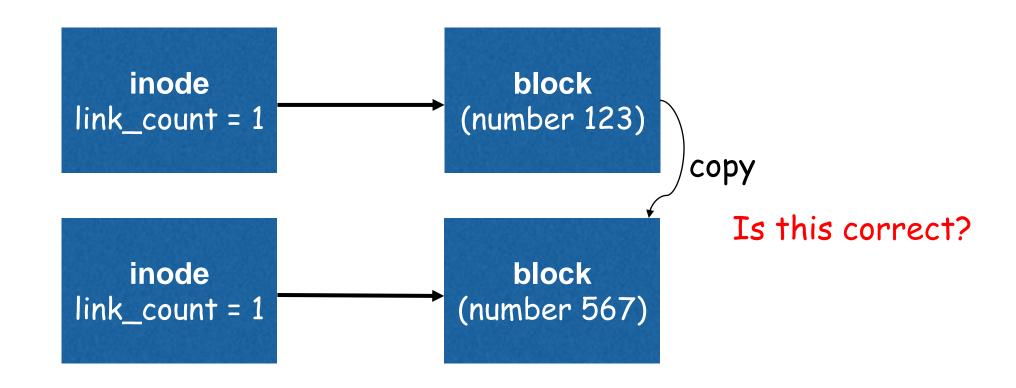
Data Bitmap



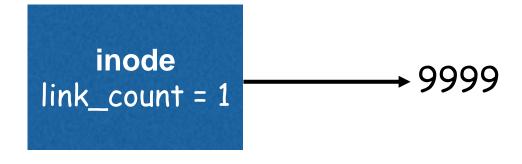
Duplicate Pointers



Duplicate Pointers



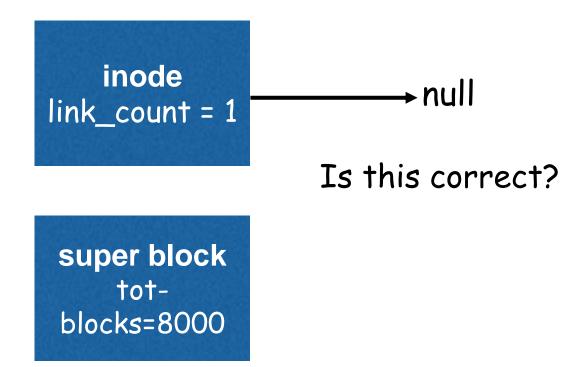
Bad Pointer



super block totblocks=8000

How to fix?

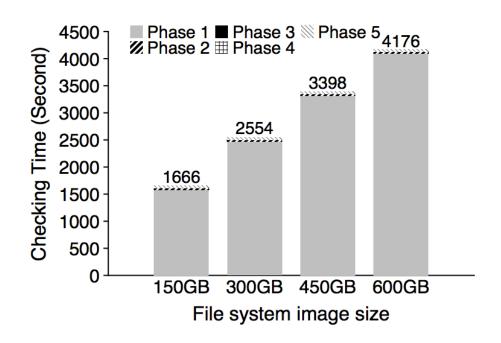
Bad Pointer



Problems with fsck

- Not always obvious how to fix file system image
- Don't know "correct" state, just consistent one
- Easy way to get consistency: reformat disk!

Bigger Problem: fsck is very slow



Checking a 600GB disk takes ~70 minutes

ffsck: The Fast File System Checker

Ao Ma, EMC Corporation and University of Wisconsin—Madison; Chris Dragga, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau, University of Wisconsin—Madison

Better Solution Goals

 Ok to do some recovery work after crash, but not to read entire disk

 Don't move file system to just any consistent state, get to correct state

Journaling meet the goals

Journaling

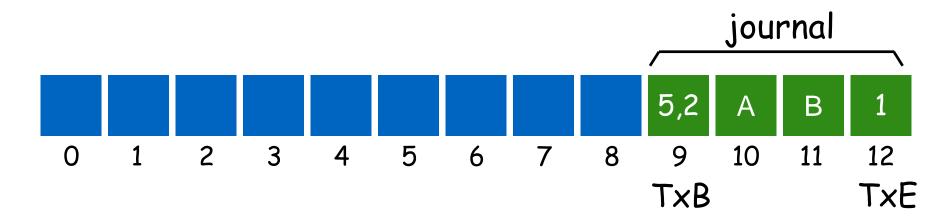
- Prepare a log with details of the update operations
- Write the log before the update operations on disk
 - Write-ahead-logging (on disk)
- If there is a crash during updates, on recovery, read the log and perform the updates
- What happens if system crashes while logging?

FS Structure with Journaling



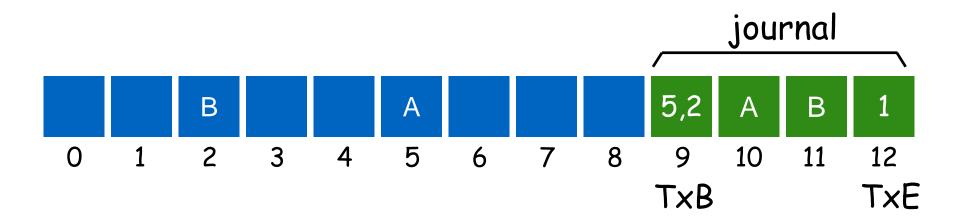
EXT3 FS structure

Journaling Example: Logging



transaction: write A to block 5; write B to block 2

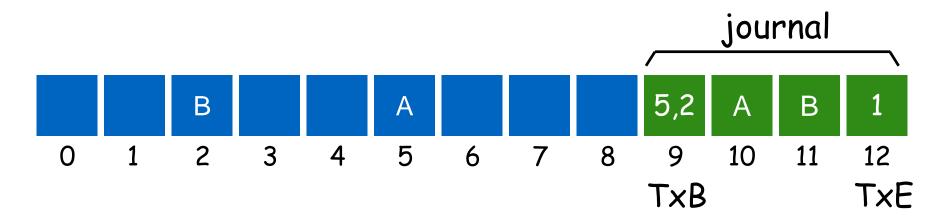
Journaling Example: Checkpointing



transaction: write A to block 5; write B to block 2

Checkpoint: Writing new data to in-place locations

Total Ordering



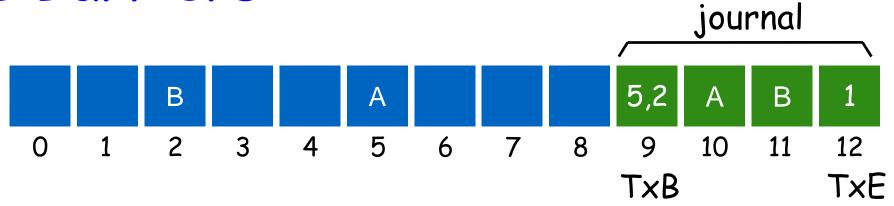
transaction: write A to block 5; write B to block 2

write order: 9, 10, 11, 12, 5, 2

Enforcing total ordering is inefficient. Why?

Random writes

Use Barriers



transaction: write A to block 5; write B to block 2

write order: 9, 10, 11 | 12 | 5, 2 |

Use barriers at key points in time:

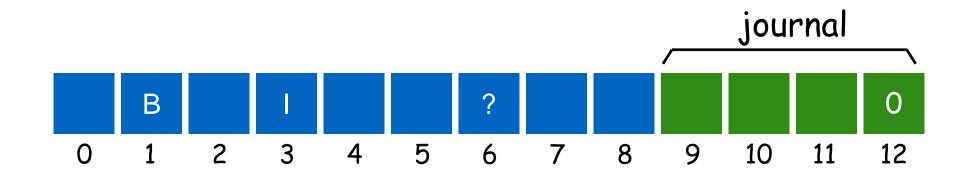
- 1) Before journal commit, ensure journal transaction entries complete
- 2) Before checkpoint, ensure journal commit complete
- 3) Before free journal, ensure in-place updates complete

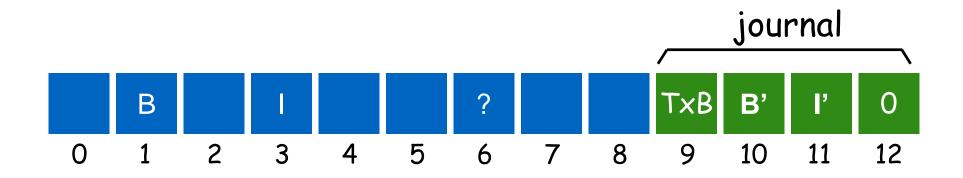
How to avoid writing all blocks twice?

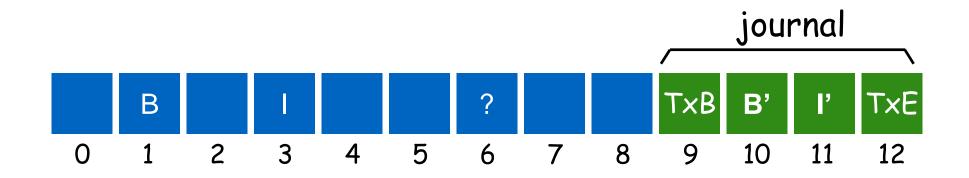
• Observation: some blocks (e.g., user data) are less important

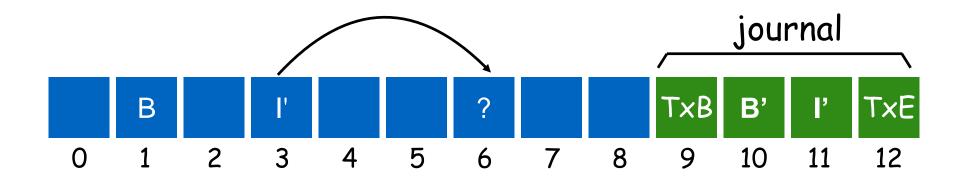
• Strategy: journal all metadata, including: superblock, bitmaps, inodes, indirects, directories

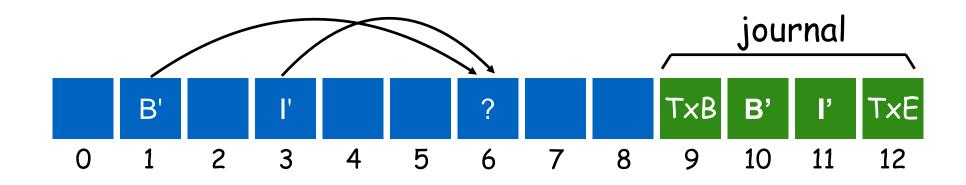
For regular data, write it back whenever convenient.
 Of course, files may contain garbage.





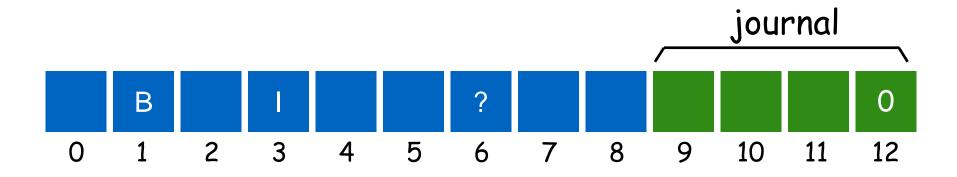


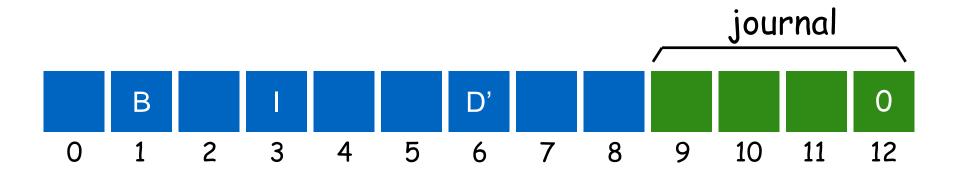




transaction: append to inode I

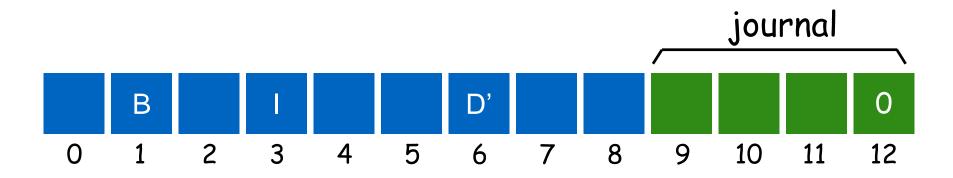
what if we crash now? Solutions?





transaction: append to inode I

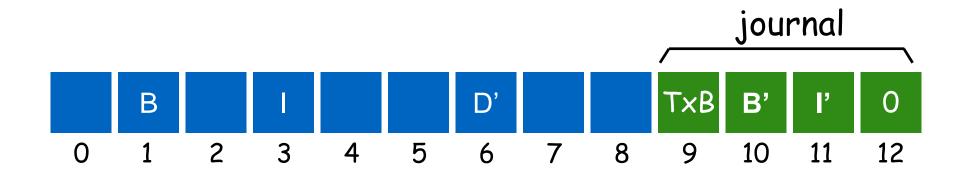
Write data before journaling

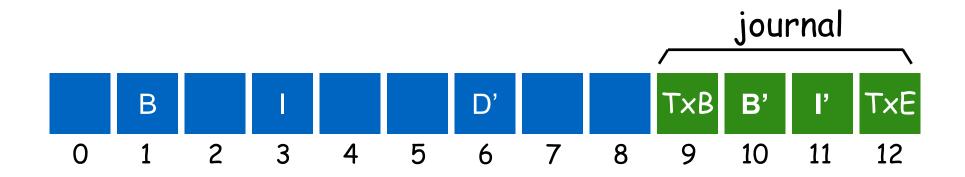


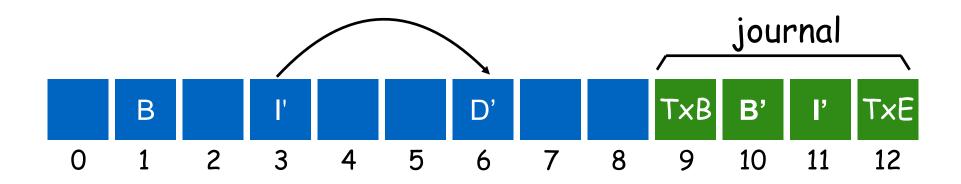
transaction: append to inode I

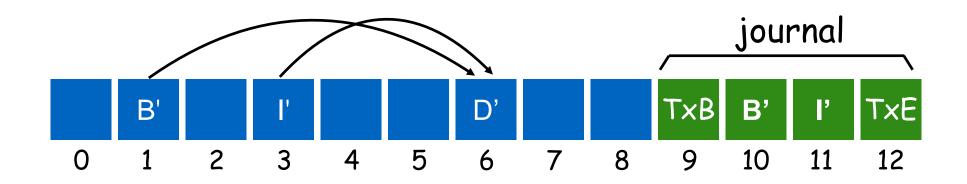
Write data before journaling

What happens if crash now? B indicates D' currently free, I does not point to D'; Lose D', but that might be acceptable









Which one to choose?

No need to choose one

- Provide all option to the end users.
 - · Let them pick it
- Most modern file systems use journals
 - ordered-mode for meta-data is popular

Disclaimer

• Some of the materials in this lecture slides are from the lecture slides by Prof. Arpaci, Prof. Youjip, and other educators. Thanks to all of them.